

Contents

August 2022	10
August 26, 2022	10
August 28, 2022	10
September 2022	11
September 06, 2022	11
September 08, 2022	11
September 09, 2022	12
September 12, 2022	12
September 13, 2022	12
September 14, 2022	12
September 15, 2022	13
September 16, 2022	13
September 19, 2022	14
September 20, 2022	14
September 21, 2022	15
September 22, 2022	15
September 23, 2022	15
September 26, 2022	16
September 27, 2022	16
September 28, 2022	16
September 29, 2022	17
September 30, 2022	18
October 2022	19
October 03, 2022	19
October 04, 2022	19
October 05, 2022	19
October 06, 2022	20
October 07, 2022	20
October 10, 2022	20
October 11, 2022	21
October 12, 2022	21
October 13, 2022	22
October 14, 2022	23
October 17, 2022	23
October 18, 2022	23
October 19, 2022	24

October 20, 2022	24
October 21, 2022	24
October 24, 2022	25
October 25, 2022	25
October 26, 2022	26
October 27, 2022	26
October 28, 2022	26
October 31, 2022	27
November 2022	28
November 01, 2022	28
November 02, 2022	28
November 03, 2022	29
November 04, 2022	29
November 07, 2022	30
November 08, 2022	30
November 09, 2022	31
November 10, 2022	31
November 11, 2022	32
November 14, 2022	32
November 15, 2022	33
November 16, 2022	33
November 17, 2022	33
November 18, 2022	34
November 21, 2022	34
November 22, 2022	35
November 23, 2022	35
November 28, 2022	37
November 29, 2022	37
November 30, 2022	37
December 2022	39
December 01, 2022	39
December 02, 2022	40
December 05, 2022	40
December 06, 2022	41
December 07, 2022	41
December 08, 2022	42
December 09, 2022	42
December 12, 2022	42

December 13, 2022	43
December 14, 2022	43
December 15, 2022	44
December 19, 2022	44
December 20, 2022	45
December 21, 2022	45
December 22, 2022	45
December 27, 2022	45
December 28, 2022	46
December 29, 2022	46
January 2023	47
January 02, 2023	47
January 03, 2023	47
January 04, 2023	47
January 05, 2023	48
January 06, 2023	48
January 09, 2023	49
January 10, 2023	49
January 11, 2023	50
January 12, 2023	50
January 13, 2023	50
January 14, 2023	51
January 17, 2023	51
January 18, 2023	51
January 19, 2023	52
January 20, 2023	52
January 23, 2023	53
January 24, 2023	54
January 25, 2023	54
January 26, 2023	55
January 27, 2023	56
January 30, 2023	56
January 31, 2023	57
February 2023	59
February 01, 2023	59
February 02, 2023	59
February 03, 2023	61
February 06, 2023	61

February 07, 2023	61
February 08, 2023	62
February 09, 2023	62
February 10, 2023	62
February 13, 2023	63
February 15, 2023	64
February 16, 2023	65
February 17, 2023	65
February 20, 2023	65
February 21, 2023	65
February 22, 2023	66
February 23, 2023	67
February 24, 2023	67
February 25, 2023	68
February 27, 2023	68
February 28, 2023	69
March 2023	70
March 01, 2023	70
March 02, 2023	70
March 03, 2023	71
March 06, 2023	71
March 07, 2023	72
March 08, 2023	72
March 09, 2023	72
March 10, 2023	73
March 12, 2023	74
March 13, 2023	74
March 14, 2023	74
March 14, 2023	75
March 15, 2023	75
March 16, 2023	76
March 17, 2023	77
March 20, 2023	77
March 21, 2023	78
March 22, 2023	78
March 23, 2023	78
March 24, 2023	79
March 29, 2023	79

March 30, 2023	80
March 31, 2023	81
April 2023	82
April 03, 2023	82
April 04, 2023	82
April 05, 2023	83
April 06, 2023	84
April 07, 2023	85
April 08, 2023	85
April 10, 2023	86
April 11, 2023	87
April 12, 2023	88
April 13, 2023	88
April 14, 2023	89
April 15, 2023	90
April 16, 2023	90
April 17, 2023	90
April 18, 2023	91
April 19, 2023	91
April 21, 2023	91
April 22, 2023	92
April 25, 2023	92
April 26, 2023	93
April 27, 2023	93
April 28, 2023	94
April 28, 2023	95
May 2023	96
May 01, 2023	96
May 02, 2023	97
May 03, 2023	98
May 04, 2023	98
May 05, 2023	99
May 08, 2023	99
May 09, 2023	100
May 10, 2023	100
May 11, 2023	101
May 12, 2023	101
May 14, 2023	101

May 15, 2023	101
May 16, 2023	102
May 17, 2023	103
May 18, 2023	103
May 19, 2023	104
May 22, 2023	104
May 23, 2023	106
May 24, 2023	106
May 25, 2023	107
May 26, 2023	107
May 30, 2023	108
May 31, 2023	108
June 2023	110
June 01, 2023	110
June 02, 2023	110
June 05, 2023	110
June 07, 2023	111
June 08, 2023	111
June 09, 2023	111
June 12, 2023	112
June 13, 2023	112
June 14, 2023	113
June 15, 2023	114
June 16, 2023	114
June 20, 2023	115
June 21, 2023	115
June 22, 2023	116
June 23, 2023	116
June 26, 2023	117
June 27, 2023	117
June 28, 2023	118
June 29, 2023	118
June 30, 2023	119
July 2023	120
July 05, 2023	120
July 06, 2023	120
July 07, 2023	121
July 10, 2023	122

July 11, 2023	122
July 12, 2023	123
July 13, 2023	123
July 14, 2023	124
July 17, 2023	125
July 18, 2023	125
July 19, 2023	126
July 20, 2023	126
July 21, 2023	127
July 24, 2023	127
July 25, 2023	128
July 26, 2023	129
July 27, 2023	130
July 28, 2023	131
July 31, 2023	132
August 2023	133
August 01, 2023	133
August 02, 2023	133
August 03, 2023	134
August 04, 2023	134
August 07, 2023	135
August 08, 2023	135
August 09, 2023	136
August 10, 2023	136
August 11, 2023	137
August 14, 2023	137
August 15, 2023	138
August 16, 2023	138
August 17, 2023	139
August 18, 2023	139
August 21, 2023	139
August 22, 2023	140
August 23, 2023	141
August 24, 2023	141
August 25, 2023	142
August 28, 2023	143
August 29, 2023	143
August 30, 2023	144

August 31, 2023	144
September 2023	147
September 01, 2023	147
September 03, 2023	147
September 05, 2023	148
September 06, 2023	148
September 07, 2023	149
September 08, 2023	150
September 11, 2023	151
September 12, 2023	151
September 13, 2023	151
September 14, 2023	152
September 15, 2023	153
September 18, 2023	154
September 19, 2023	154
September 20, 2023	154
September 21, 2023	155
September 22, 2023	156
September 25, 2023	157
September 26, 2023	159
September 27, 2023	159
September 28, 2023	160
September 29, 2023	161
October 2023	163
October 02, 2023	163
October 03, 2023	163
October 04, 2023	164
October 05, 2023	165
October 06, 2023	165
October 09, 2023	166
October 10, 2023	167
October 11, 2023	167
October 12, 2023	168
October 13, 2023	169
October 16, 2023	169
October 17, 2023	170
October 18, 2023	170
October 19, 2023	170

October 20, 2023	171
October 22, 2023	171
October 23, 2023	171
October 24, 2023	171
October 25, 2023	173
October 26, 2023	173
October 27, 2023	174
October 30, 2023	174
October 31, 2023	175
November 2023	177
November 01, 2023	177
November 02, 2023	177
November 03, 2023	178
November 07, 2023	178

August 2022

August 26, 2022



So, my work at

EarCandy starts on September 6th, that's, after the holidays. I already received a computer, so there. We will see. I have one last interview in one global company, on Monday. Not that I'm interested, just curious.

August 28, 2022

Got an email from ebay, saying that my qualifications are not enough for the position of MTS-1, Software Engineer.

That's about how views differ for different people and companies.

September 2022

September 06, 2022

Everything's fine and relaxed at EarCandy. But just that I don't have a VPN. It turned out, nobody does. I'm registering in a bunch of places; then a 1-1 with Jacob, my manager; then a 1-1 with Victor, who was helping me, showing to me what's where. Then we had a standup, everybody introducing themselves; it does not look like there's a lot of work to do (well, I found a bug in their Jira case: internationalization is trickier than English speakers imagine); then another 1-1 with Victor: he was showing again, what's where in github, and told me that William, the biggest talent there, just left to do IoT in Rust. Well. Sad. Somehow I have missed Rust. I tried it on project Euler, and it didn't click. Just didn't click.

At 9pm we chatted with Blanka, over MS Teams; I showed her my passport (but did not give it to her to hold).

Oh, and VPN is broken in the company. So what now? I have to regulate my work day. E.g., Noon to 9pm, with a short break for dinner and a walk. That would be good.

September 08, 2022

Ok, a progress: I got access to projects on github, and to AWS, and other stuff.

But yesterday Cisco VPN had eaten 60 gig of traffic, with all this Indian nonsense inside. It's Cisco, I know them. As a result I'm out of my 80 gig that I have per month. Nice!

Well, ok, now I'm saving on bytes.

With Victor, we watched all those graphs in Grafana. Funny; CDN goes for updates once every 30 minutes for some services, and for others it's 15 minutes. Hence we have peaks. Can they, instead, go with a period, in milliseconds, that is a prime number? No peaks, right? Well, I tried to sell this idea in I don't know how many companies. They don't seem to get it.

Another service is visited every morning at 10am, with a request for a huge list. Currently we are investigating, who's that visitor, and what exactly does he/she want?

The code, though, looks relatively clean. According to the logs, errors only happen in some ancient code that was not converted to Scala yet. The errors are not NPEs, they are "alternative solutions", like "HTTP 404", we are replacing non-existent data with default values, since we are in the Eilenberg-Moore category. Sure, Java people do it all the time. They have long gone, actually, but the code stays. Not a problem to fix it, but when, and who is going to fix it? Maybe later.

That's it for today. I love honeymoons in new companies.

September 09, 2022

Nothing special. Adding logins, configuring stuff.

Learned about a new cli, called gh. An interesting thing, have to investigate.

Karen is asking, why do I suffer reading docs online, they are all in github; clone it and read it. Fuck, was not I a little bit dumb all the time?

September 12, 2022

Nothing special so far.

From morning to 4pm French time I was trying to access the internet, that including attempts to get out of this “Omsk” to Agen nearby, and talk to the people from there. Eventually managed to connect.

I’m installing NIX now (an amazing thing) and aws cli (the most important part is getting an access, the rest is trivial).

The scrum meeting was pretty short; then we had a chat with Kaaren, who gave me links to docs. Victor is sick with Covid, suddenly, so we are not bothering him today.

Then I had a 1-1 with Jacob; he was showing me what’s where, and we discussed the priorities. So far so good. Working on it.

There’s another company, something related to electric cars, they want me for a Principal position. Somehow I feel too lazy for that. I don’t know, how much can they offer so that I rush again to make things beautiful and fight for the code quality.

And another company wants to talk to me, after September 20th. I don’t think I need any of this.

September 13, 2022

Nothing special so far. Nix, AWS, sbt clean package - so far none of these succeeds.

Had a scrum (fast, efficient); then Karen tried to help me with AWS, and told me that sbt won’t pick the jars from jfrog until I configure nix properly. Ok, thanks. The doc is good, but slightly disorganized.

By the end of the day wrote to Swathi about my aws access problems. It’s 22:02 already, will wait, maybe she responds, who knows.

September 14, 2022

So, the install process is going on. AWS unavailable, so I can’t install NIX, and sbt does not build. Sent them a message.

After a standup Jacob, and Swathi, and Karen started bugging me regarding setting up a zoom session; I wrote to jacob, and with Swathi and Karen we started fixing my configs. So far I have a login (Karen suggested me to fix the docs: I do it with pleasure!), and Swathi said that I “don’t have a role”, and she’ll write to Jonathan, and Jonathan will give me a role, and then everything will be fine (I’ll shine on the stage (It happened to me once, IRL, I even sang, imagine).

Regarding NIX, Karen has shown me the ropes, now I have stuff to do. Then I went for a dinner, and meanwhile, they invited me for a sprint-planning-meeting, and I wasn’t available. Oops. I had to add these meetings to my calendar.

Right now I’m (de-)bugging NIX. The script is beautiful, I’m really impressed. Never saw such humane scripts ever, it was only me who wrote such scripts, from time to time. Ok, of course, no kidding, NIX is better than sliced bread, isn’t it?

So there, and I continue my work. I’m not accustomed to accounting my work time, and working from 10 to 10 is pretty stupid too. I need to self-organize.

September 15, 2022

Not much of interesting stuff. I followed instructions, I configured, I wrote down what wasn’t working, and then, when Swathi woke up and got online, together we had walked through the defect; she told me what’s up here and there (and I wrote it down); then we discussed it with Karen; I will go ahead and fix the docs, to make it up-to-date. And I still have problems with JFrog.

Jacob too tried to help me, but I wrote to him that I won’t be bugging him for a while, while there’s a progress; as soon as it jams, I’ll ask for his help. Tomorrow, most probably. Today I’ll go fix the docs.

The major trick with github here is that nobody merges into no master at all. First you fork, then you do whatever you like in your fork, then you send a PR for a merge to master (or is it main?) Interesting. I’ll have to try it with the docs.

Aha, done configuring Nix. Karen says that we have tons of useful scripts, for github, for instance. Haha, I have some too, in my stash. We will see. I’ve already written a nice script for AWS login (as if I need it).

Strangely, I wrote this entry in Russian; had to translate it. Weird.

September 16, 2022

Well, spent the day trying to configure and build (not much success). By the end of the day it turned out that the docs that are deployed are fresher than the docs that are in github. Karen updated them first in her fork, then deployed, then eventually merged - and I was totally confused: doctor, I see not what I hear.

Anyway, that's fixed now.

Another thing is brew files; pretty good stuff, you can describe all the required configuration without having to look around, where do I take this tool, that tool, etc. `brew bundle --file=(your brew file)`. I recommend.

That's it for today; will go do something tomorrow, not much either.

September 19, 2022

Shit, AWS is causing problems with NIX, NIX is causing problems with JFrog, and as a result, nothing builds. I only sent a PR to Karen, regarding docs updates.

Chatted with Swathi, with Karen, with Victor, who managed to get cured from his Covid real fast; with Jacob who gave me a couple of ideas; but in general it all just sucks.

With Victor, we have kicked out sbt, cleaned ivy cache, with no progress. Now sbt wants to install Scala 2.10: where did it dug out this stewardess from, I wonder.

In short, a pretty shitty day today. Congratulations.

September 20, 2022

Today was a day of successes. Ok, I played with aws and eks, and it was failing, and I bugged Karen and Swathi with that. Then we had a “Jira/Scrum” training, where our “product” Renato was telling me which buttons to click. I figured the guy is a little bit deranged, and was just rading dw, while he was trying to explain to me how to add points. He ignored my hint that there’s a write-up on Linkedin that I wrote about 4 years ago, relating to this.

Then we got together with Sean and solved the funny problem with Fastly that I was having. Identification problem.

Anyway, then there was a scrum, when I told again about my troubles; then we got together with Swathi. Did I tell you I respect Indian girls a lot. Patient, smart. We fixed the configs for aws and eks, and I just copy-pasted my terminal log, so that next if something happens, I’ll see how we fixed it. Whatever.

Then we experimented with Karen regarding sbt build failures (even startup failures) in my env. Karen (wrongly) convinced me that it’s not related to Nix. It is, actually, but William left a month ago, and there’s no Nix expert in the company. I’m afraid (or rather glad) to become one. Nix is cool, totally. Tubular.

There’s also another Jira case for me; I believe I know how to fix that stuff, but I want to be able to run builds locally and test the code. BTW, what I saw in the code so far is pretty beautiful. But again, just

a first impression. A honeymoon.

September 21, 2022

Ну что, понял, что с NIX тут никто ничего. Был Джо - ушел год назад. Был Вильям - ушел месяц назад. Теперь я должен разбираться.

Ну это фигня, первый шаг сделан, разобрался, на уровне, что сейчас нужно.

Смотрю, хрен ли не билдится. Такое ощущение, что jfrog меня не пускает, и не показывает, в чем проблема. “Нету файла”. якобы. Это он шлет какую-нибудь 404. Нет чтобы sbt сначала проверил, доступ есть вообще или нету. Ну чисто по-человечески. Эх. Ну… я даже имена забыл этих ребят, что sbt пилили так интенсивно. Давно дело было. Джош, он уже давно постарел, вырастил бороду… кто там сейчас за это отвечает, хз.

Ладно. Договорились с Виктором сегодня без пятнадцати девять (нашего времени) пройтись по этой херне, по билду.

Еще в восемь был митинг, а меня туда подписали через 10 минут после начала; а я увидел через 20 минут, зашел - а там уже все закончили. Ну и ладушки.

September 22, 2022

Да ну блин, даже на митинг опоздал, который в 5 был. Пришел домой в 6. Ну и вот, початился с Джейкобом, с Карен. Она мне показала пару трюков. Дженкинсные билды конфигурируются файлами, генерируемыми из groovy. Все чистенько так. И все в гитхабе, а не как попало на странице дженкинса.

Но вот заставить интеллиджей работать с конфигом из NIX пока что хрен там. SBT норм, а интеллиджей дурит. Ну… ну посмотрим. Кейс там у меня есть, надо ж гонять все эти тесты. Ну не все, ну хотя бы юниттесты.

September 23, 2022

Well, not a very smart day. Ok, slapped together a PR (remove this from there); removed too much, restored some stuff. Also, managed to configure intellij, so it almost works on those projects. Also, walked through a crazy process of *preparing a new build for deployment*. Bored, and waiting for a pr approval.

September 26, 2022

Monday as usual. Walked through the process of deployment, me asking Karen again about this and that, and recording. The test suite failed once, randomly. FIIK why.

So, still in the process of manually running a highly automated sequence of processes. Weird.

Anyway. One “standup”, one 1-1 with my manager, Jacob. Everything’s fine. So far.

September 27, 2022

Ok, now the dumb parts.

Most of the day I was writing a doc, recording my experience in updating the web api. Then “code freeze” was closed (as if I saw when it’s on and when it’s off, there’s no signal for that), and it turned out the stuff can be deployed to prod.

And then it turned out that kaboom, swagger shows the old api. And why is it? It’s because “doc build” (which produces swagger) should be run 6 times, for each “namespace”, and while it runs “automatically”, it does it for only one “namespace”. Neat. This is not hard to fix, of course, but... but I’ll put it in jira when I get deeper into the system.

Otherwise, there’s a prod to deploy. Karen, in her Chinese style, wrote me instructions on Slack. I’ll copy it over, but tomorrow. I’ll finish documenting my shit first.

Anyway, do I care? Not much. Also, should I work too much? Seems like I spend too much time at work. I should not. There’s French to be done, and other things.

September 28, 2022

My biggest problem is that I can’t organize my work time. I’d rather spread it evenly throughout the day, and then figure out, how many hours did I work? 4? 12? Wtf...

As to the business, I created an update for docs, sent for review; Karen did a good review (will go ahead and fix it); also, Karen told me to prepare prod update (via kube config), so I did. Github definitely misbehaves when I do updates via branch in fork; I’ll need to find more reliable tricks.

Is it programming what I do now? It’s bullshit.

They gave me another jira case, do something with FB token and KMS encryption. Again, it’s not programming, it’s slapping things together.

Had a nice 1-1 with Victor; pretty cool. He does 1-1 with everyone. Makes sense. Maybe I should do the same. Introduce each other. There are a couple of shy guys, and one of them seems to be a ZIO pro. With Victor, discussed the problems of build/update automation. A lot of stuff automated, making the

developer to go ahead and update this file and this file and this file, etc, while these updates could have been automated. Not that I fucking care.

Oh, and also, had an interview, for some reason. Again ZIO was mentioned; I said that I never used it, but I have no problems with monads, monad transformers, etc. The interview consisted in running some stuff in IntelliJ; they had a project, and then I had to slap together something that eventually grew to something that needs map/reduce (which I managed not to mention... who cares).

We will see.

September 29, 2022

So, I still can't figure out my working hours. In the morning, while they were sleeping, I went ahead and updated the docs that I sent for review. They did provide a nice collection of good comments, I enjoyed them. So, the next step of pr reviews.

Then I went ahead and tried to deploy my little shit to prod. No success. Ok, let's wait until Karen wakes up.

Then I looked at the next jira case. That's about KMS encryption of a token that is currently in the code. Ok, not a big deal, except that a) encryption keys vary; and b) be careful with deployment, start with a dry run.

Later we had a "standup", no problem.

Then Karen got online, and we spent over an hour working on that "prod deployment".

Again, my change is that a change in the description of several API methods. She showed me how to actually do prod deployment (duh... but undocumented), and besides, what to do during the process. 1) open a page where you see logs. Look for errors. (What errors?) 2) open grafana to see the behavior. Now I'm in trouble. I can't log in to grafana. So I wrote to Evan, and he said he'd add me right away. The rest of the hour we waited for Evan, but ok, meanwhile, I deployed this shit for Australia, and saw that nothing broke. We agreed that I'll deploy to other continents when I get access to Grafana.

Then I discussed with Karen my plan to first do a dry run, then enable it. She said that it's not worth it, just deploy your (possibly fatal) functionality, and then see. What? No way. That's not how I do stuff. But anyway, I'll first have to write the PR.

Several years ago I would be already running in circles, screaming, swearing, smoking mj. Now I just take breaks.

So we went to another village to buy a couch and a chair. When we returned, there was supposed to be a talk regarding Mongo, but it did not happen. I was exhausted (it was 8pm), and just went to sleep. After a dinner, of course.

September 30, 2022

Ok, so I had about three questions; one of them was a strange structure of a shell command that nobody could explain, but told me to use; by the end of the day still nobody could explain.

What else. KMS, some issues with design, but I'll figure it out. Karen was wrong about just throwing in an update without having a dry run.

Also, eventually I got it how to log in to kube dashboard. Except that today is Friday, and we don't deploy on Friday.

Anyway, will get back to work tomorrow. People are mostly away today. Good company.

Also, my recruiter, with whom we discuss French series, is going to call me tomorrow. He said that it's ok now to change job. Not sure whether I want it, but I enjoy development, not devopsing.

Anyway, it's just what my "daughter in law" calls "paranoia". Maybe.

I need more time to do more stuff!

October 2022

October 03, 2022

Aha, it's Monday. Today I spent some time deploying the docs I was updating. It takes time, and it's slow - you deploy it, the version in kube is ok, but the functionality is still old. Until eventually...

Anyway.

The doc I wrote is still under review; the Karens are finding if not typos then unnecessary html tags. Ok, do I mind. Will improve it next time. Why did not they write a doc, I wonder.

Also, studying the code, looking for a reasonable solution for keeping a token encrypted without using different encryption keys for different namespaces. We will see. Talked with Jacob, he says that we actually use one key. Looked through the code. The code is 5 years old, but is pretty good, so no need to fix it.

That's it. Another meeting starts right now, oops. It's 8:30pm.

October 04, 2022

In the morning I did update the doc I was "working on" again, and Victor approved it, thanks go. Two Karens (Karen and Qiong) were finding little mistakes here and there (like an anchor that is not used), and I was updating this little stupid doc accordingly, but now Victor came first to review, and approved it. He's also Chinese, but I guess he's from Hongkong (judging by his family name).

Anyway, merged it asap, and left it for now. Will update it again, later... or not.

Then there was a long break, then there was a standup, to which I could only connect with my phone - the computer did not get zoom to that meeting - so my speech was short. "Working on my case".

Then I worked on my case, found the chunk of code that does this KMS conversion, added some tests, figured that it works for me, and eventually just asked the team, whether 'tis nobler to add a library dependency or just copy-paste. That's what I did for today.

October 05, 2022

This morning I started investigating KMS functionality; some functionality was either missing or just kept in a test. So I decided to enrich it a little bit and start using it in the module I'm working on.

By the end of the day Karen convinced me that all those tricks are already there, and done automatically - if a field of a config is listed as encrypted, it will be decrypted. The only thing is that you should properly encrypt it, because keys vary for different platforms.

But that's ok. So, is not it tricky! Is not it smart! Will do it tomorrow. Tired today.

Also, the company decided to abandon Zoom and switch to MS Teams, just to save money. Oh, whatever. Not my money.

And I'm tired.

October 06, 2022

Got my shit together, and decided to write that PR, with encrypted fb token - but carefully, since who knows what will happen IRL. So as the first step I'll be decrypting and comparing with the expected value, and carefully report the first seven failures, always returning the right one. We will see how it behaves in staging, and then in prod.

The only thing was, I forgot how to do proper encryption (with the proper key), so spent some time looking.

By the standup meeting my PR was ready, tests passed, sent it for review. 30 minutes later got a pretty good comment from Karen; fixed my shit right away, sent for review, and she approved it, and I merged it. Will check tomorrow, deploy and watch the logs.

Oh, ok, maybe will have to update kube config too. NP, tomorrow.

At 8pm there was another meeting, guys from mongo were showing their tricks (mostly about performance improvement by eliminating unnecessary indexes). They have a pretty dashboard now, just beautiful. The rest is boring.

October 07, 2022

Almost nothing today.

Tried to update a config for kube to deploy my latest code. Did not work in Nix - spent half a day investigating, eventually found out that no, nobody does it in Nix. Anymore. So I tried to do it outside of Nix. Got another failure, informed Swathi, because the bug is hers. That's it for today. Sad.

Oh, also, gave another interview to a company with which I've been talking since the end of July. And? FIIK.

October 10, 2022

Ok, I did run make in kube-config, as Swathi said, the error happens, but the script does its job, so in a sense, nobody cares (except me); sent my PR, and the only issue was, how exactly do I insert the diff link in the comment. Never mind, Charles showed it to me. Now I know.

Eventually Karen has approved it, and I'll deploy it to staging, and we'll see.

There was a standup at 5pm; since we are on ms teams, connection was problematic, so when I joined eventually, my turn was already passed, but who cares.

At 7:30 I was supposed to have a 1-1 with Jacob, but he's sick, so I sent him my "points" (what I think is worth doing), and he sent me to Karen; ok, sent it to Karen.

Ok, with the feature I'm deploying, have to figure out how to test it. Scripting is so far beyond my capabilities here, but at least manually. No misbehavior expected anyway.

Oh, and the fun continues. Deploying something that has configs is async: you either deploy configs first (and so the older service will get confused), or you deploy the code first (and it will be confused by the old configs). Will have to know it next time. Currently my code is pretty resilient, but well, wtf... At least I'm only deploying it for just one continent.

October 11, 2022

As it becomes usual, I work almost by EST zone now. Anyway, it turned out that the service in staging is not getting any requests, so there's nothing to look for in the logs.

I did some investigation, but eventually it was Karen who told me that I also have to deploy "social job", and that's what I did late in the evening; now I see that a) no trace of any my requests getting into the service, and b) a bunch of HTTP 400 in some interactions, but it's all websockets, and my requests are plain HTTP.

Still trying to fuck with kubernetes, tokens, logins, jobs, services. So far some weird stuff.

Anyway. Not boring, but weird. Can we just write code please? Something that many people are clueless about.

October 12, 2022

Ok, so I added some logging to the service I was tweaking, and then it turned out that the only consumer is some "JobRunner", so I added some logging to that runner, and then... then I tried to run it locally, as Karen suggested yesterday.

Because wtf, it's not a service I was tweaking, just a class that hardly has any api (it does have some).

The rest of the day I spent trying to figure out wtf is with the config files. It's Config from Typesafe. The logic in the code is that if we don't find a config file, we just ignore it. That's what I could not believe, and had to trace the Typesafe's monadic coding style. To me the code looks a bit ancient (ok, 2017), and I'd rather use something better, but anyway.

Eventually made it run, but well, it also wants to talk to a config server, and Cisco, in response to that, tells me that I HAVE TO install some shit from them. Wrote to our SREs and closed that topic

meanwhile.

We also had a planning meeting, so I had to create a couple of Jira cases. And that was it.

Oh, no. Hari, a guy that is leaving the company next week, was giving an amazing talk regarding Scalameta. Which is being used in our company. And Hari, I'm not 100% sure, but he was supposed to be my student at SCU, since I was teaching in COEN dept, and he graduated from there, and my course was required. I don't have to remember all of them, but this guy is pretty brilliant, I should have.

Anyway, that's it. Ok, and had a 1-1 with Victor; nice little chat. I don't know Victor's role, but he seems to be a facilitator for all issues, and at the same time he's technically, well... ept? opposite to inept. Good guy.

Anyway, it's 9pm, and normally I would have worked more, but let's hear from SRE before I proceed accessing stuff.

Oh, and I wrote a script to properly log in to kubernetes, or rather to get kube's token, depending on which server we are trying to reach.

So, this kind of job, it's pretty lazy, but still I don't have much time for anything at all. Weird.

October 13, 2022

In short, I've cleaned up logging, slapped together a new test, started trying to run it again. It has to go to Amazon, decode something... and it doesn't. It doesn't reach Amazon.

Well, I had to investigate it. Judging by the code, this token, the one that is in the config, is not being used at all. Instead, some tricky shit is being used, individual for each user. And? What now? Now I have to a) find a way to get real data, what exactly are we doing here, and b) create a test that validates that we are not producing nonsense, but that we do what we are supposed to do. Been there, when I was fixing one buggy open-source SAML implementation. It turned out that our partners (Fidelity actually) were sending us a request where a) all the data are encrypted and signed; and b) a certificate is provided, literally certifying that "I certify my own signature". Well, anybody can do that (as a veterinary once said to his doctor). Seriously. What do you certify? I certify the fact that I certify.

In short, it's just a total bullshit.

Meanwhile, my code, with logging, was approved - but what's the point/ Zero. We agreed with Jacob to discuss this today, or, rather, tonight, at 9pm (our local French time).

So, see you tomorrow.

October 14, 2022

It's Friday, and all the meetings are in the evening, company-wide ones. About a hundred of us got together to listen to the news that we could actually just google. Except for one interesting news, that we beat one large company, but was it interesting? Not for me.

In other news: starting this morning, I was failing to log in to Amazon; it started working only when our devops girl asked me what's wrong. I tried it again, and it worked now. Ok, never mind.

As you see, I could not get to Amazon, hence I created a special config file with nothing encrypted. It would be ok for a test. And then what, then I have proven my hypothesis that tokens are coming from users records; then I read something about Facebook, and it became clear to me that sure, of course. Wrote to Victor and Jacob, asking their opinions. And then silence. It's Friday. Free day. So there.

Right now, that's, at 9pm local time in France, we have another meeting, a farewell for Hari (Hari? Not Harry? - asking in 2025). He's leaving, and we'll drink to his health, via Microsoft Teams.

Oh, and besides, it turned out that the problem with (not-)accessing Amazon was that our devops people were playing with DNS, and so we just had no access to Amazon. Well, what can we do. Next time mock this Amazon, it's not the network that we are testing in our half-integration tests, right?

October 17, 2022

In the morning I looked into the calendar and panicked. They had the first ZIO training session on Sunday. So I wrote to them, asking whether it was recorded. By the end of the day it turned out that there was no session on Sunday. But I managed to download the training project and get through at least half of the first day training. Pretty cool actually, if you do it on your own, without advisors.

Anyway, and then I canceled one Jira case, updated another, sent one PR for review, sent that other Jira case for review, wrote a PR for that case, sent it for review, and now I'm waiting for it to be re-reviewed.

Had a short 1-1 with Jacob, showed him my script for kubernetes token extraction, discussed ZIO, discussed that Jira case, and basically that was it. 10 minutes from now ZIO training starts. So I have a 10-minute break, and then will be there for the next 3 hours, till 11pm. Oh my. Ok.

October 18, 2022

So, last night I was on ZIO training, till 11pm, almost fell asleep. ZIO is great. $ZIO[R, E, A] = (A+E)^R^A$, categorically speaking.

Anyway, in the morning had to finish all the exercises from that first part, and then, just in case, from the second. So I hardly had much time for work; but there were a couple of important notes from

Karen, and one good idea about logging tokens, so that's what I did, implemented what she suggested.

Now there's the next session of ZIO training, and, thank god, Jorge does not rush, but reviews the yesterday's material.

That's it for today.

October 19, 2022

Pretty busy. Did not see the video link from last night's ZIO training, so did the exercises while I could. Of course with some mistakes. Anyway, by the end of the day the video showed up, but I had to work.

Also in the morning I was pretty sleepy, so the PR to deploy the service was for prod, instead of for staging. Good that I figured it before it got approved and merged. It's really hard to roll back (which is bad).

So updated the docs meanwhile, sent the pr for review (ok, in another pr improved the logging, merge went bad).

How many meetings did I have today? This one is number four, but it's training. During this training I also did my daily French homework.

Now I'm in the class; got a jar of coffee (or else I'll fall asleep), drank half, and sitting there alert as a cat at the window.

October 20, 2022

Hardly anything happened. I was drowsy most of the day. Update some docs, adding links to the ZIO training talks. Could not find traces of the job I was supposed to deploy. Until by the end of the day Karen eventually got a minute to get through this, and we discovered that I did deploy the service, but not the job. Oops. Sent a pr for kubeconfig to deploy the job. Then the ZIO training started, part 4.

Before that I actually did spend some time watching the previous recordings.

That's it. This will end at 11pm, and now it's just 8:22. Shit.

FiberRef is a pretty powerful data kind in ZIO.

October 21, 2022

Mostly spent time rewatching the recorded ZIO training. Somewhere in the second half of the second day. Also, sent four questions to Karen, because these versions bumps were totally unclear to me. Eventually she answered, everything's fine, will implement it by Monday noon.

But feeling sleepy all day, including this evening now, when Jorge is talking about aspects, layers, loggers, metrics... boring. There are just five “students” present. Jorge promised to give us tests in the end. I’m afraid I’ll be asleep by that time, although I got a fresh jar of coffee. We will see. Just saw Jorge using contramap, pretty casually. Contramap is map for contravariant functors.

October 24, 2022

Ok, it was all pretty boring. I’m still struggling with specifying versions of service, job, config, storing this shit in a bunch of different files and making sure that they all match properly. Deploy a job to kube, it fails, because a config file is missing, because in another config file I specified an older version of a config file, while I already bumped the version of a config file in a build. What happened then to an older version, eh?

Anyway, it does not take much time.

Now I’m waiting for Karen to get free from her meeting to help me figure out how come my latest deployment, successfully finished by Jenkins, is not visible in Kubernetes dashboard. The dashboard shows my failed deployment, a couple of 3-week-old deployments, and one failed 7 months ago.

Meanwhile I’m studying French (which is a little bit boring too).

At 8pm we had a meeting regarding wtf to do with Akka. It costs money now, and the company is not used to pay money for software... generally speaking. They don’t have a budget for other people’s software. On which side am I? Generally speaking, on developers’ side, but not in this specific case of Akka. Jonas is a CEO of Lightbend, I did not know. The software was free, but now they decided they want money. Nobody wants to pay.

October 25, 2022

Looks like just a summary of work events.

While on the road to Bordeaux, was working on the docs, since planned to bug the people regarding that job deployment failure, and everybody was sleeping.

Posted the pr, and, basically, because my sugar was too low, stopped it until lunch. Lunch was good. You know, french, menu... pretty cool. But after that I did not work, and did nothing.

Then, in Bordeaux, I was having a hard time, because the meeting was supposed to start at 5, but by 4:50 there still was no internet in the apartment we were renting. So I was stressed like hell, sitting in the bedroom, using the precious gigabytes from my french phone plan (80 gig per month). But Dieu Merci, right after the meeting they found another wifi in this big building, and we connected to it. Good.

Anyway, at the meeting I shared my problems (Karen is sick), some people volunteered; one guy showed me where to find logs in kube dashboard (sequence of clicks), and I saw that, well, shit, the config file versioning may be not working at all, and it's neither mine nor Karen's area of expertise, and probably both expert left at least two months ago. Profit, I'm just rolling back my paranoid pr.

Anyway, then there was a meeting where the organizer did not show up, so it was canceled - and we went for a walk in this spectacular town. After that we kind of rushed home so I can catch the 7:30 meeting, but that meeting was meaningless, because it was just Swathi's placeholder, and I confirmed with her that I have nothing to do with the problems of accessing postgres in prod (it's a new devop in India that is supposed to help, not me). So I left, and now I'll have a cup of tea, and there will be another meeting at 9:30 PM, and then I can go to bed.

October 26, 2022

Nothing's going on, except that my rollback did not have any effect, and I wrote to the team and to Jacob, asking wtf I can do - but then I looked through the log carefully, and got it. I'm checking the file presence, but the path is actually an http url, so there. Ok then, this was easy.

Now waiting for deployment approval.

Also, had a planning meeting, and our PM insisted on me and Karen first finishing what we are currently doing. Really cool.

October 27, 2022

In short, I was thinking that I managed to track the job's bad behavior: but no. No logs in the view. The job does start; and when I inject an error, it does successfully crash.

Karen suggested to use the "old logger" (not the one that Victor suggests). I don't think it will help, but I'll try.

October 28, 2022

The whole morning, almost the whole morning, I was busy visiting museums (ok, some time in the morning I managed to prepare a pr, and it was approved during the day), but anyway, when we got home, I was exhausted, and just stayed in bed for a while.

Then reported on slack what I was doing last night and that today I'm not in a good shape yet.

But later on I deployed it on staging; yes, no logs in Kibana. So, logging problems. OTOH, no errors (no exception thrown) in the last couple of weeks, so what did we break? Nothing, right? Let's merge to prod (Monday), and create another jira case: "fix logging".

That's what I did. And slapped together a pr for deployment to prod, sent it to the people, let them sigh and get worried.

October 31, 2022

In the morning, while getting from Bordeaux back home, I created 6 jira cases (and left 2 for the future), and wrote to Renato about them. Renato is a PM.) And then took a break. Then promised to deploy to prod, and started looking for any links or documentation, how t.f. to deploy to prod. I did not document it last time, left for the future.

By the end of the day I have looked through all my records, ok, nothing in text. Then there must be some screenshots in Desktop, and these screenshots have the urls of prod dashboards etc. I don't even know where's prod jenkins.

So I updated all the documentation according to the info extracted from my scraps, and decided to look through the screenshots, before begging Safa or Karen. Not documented, so I will.

1-1 with Jacob did not happen, it's Halloween, so he, I guess, is busy walking around his neighborhood.

Also, saw a very good resume of a great guy from Twitter that just applied. Thank you Musk. Twitter was a good Scala school for many. Now the community will be getting the crop, and Musk will get half of Bangalor, cheap by dozen, to explain Twitter code to him. People on Twitter are already laughing at all this.

November 2022

November 01, 2022

Well, I have to deploy to prod today. So I did. But somehow I deployed an old version.

Bugged Safa, then Swathi, then Karen. Karen suggested a) to bump config, b) to do that build, that build... ok, now we have a plan.

Did not work, service crashes on deployment. Which is kind of good, but it does not tell wtf. Asked the people. Qiang sent me a link to the logs. No, I was clueless about this trick, kube logs whatever. Anyway, it's a config being wrong. But my config is ok. Now what.

Then I had a 1-1 with my manager, and he looked into the problem, and found it: I should not have bumped the config version, as Karen suggested. Ok, ok, now I know. I'll just update the code to point to a new config.

Then Qiang wrote to me that my code points to config 6.0, but the config service returns config 4.0. Wtf, right? This sounds reasonable; if a config service had 7.0, it should have had 6.0 as well.

So I wrote to Qiang, just to confirm that I understand it correctly (and where t.f. is the build log?!), and will do two bumps asap. Big deal, one can make 100 bumps, if previous configs are still preserved (logically they should be). But the question is, why t.f. no logs show the exact problem? Is everybody so smart?

Ok, and in another development more Twitter people are filing their resumes. Good. Like.

November 02, 2022

Not much to say today. Figured that repairing logging now is beyond my knowledge level (logs getting from kube to some aggregators, to some services, and then not visible in kibana - I just don't know that shit yet), so I decided to fix one build script, but first finish docs. Spent half a day fixing the docs, sent them for review, merged them, and also informed the colleagues that the .md we use is not the standard .md, but a flavor, and see how lists and links are supposed to be formatted.

That's it for today. I missed one meeting, late in the day, because I thought that it's Tuesday.

Oh, but at a meeting before one guy was talking about... well, something. About his code, the solution. The code looks mostly beautiful, but it was impossible to read anything, because he was jumping through files while talking. 5 seconds per screen, when was I supposed to read the code? I'm a slow reader. Others? Hard to tell. Same was with swagger files - how can I read the docs preceding the curl calls? No way. So I just decided to ignore. Whatever. I can read the code myself, with my speed. One thing is strange, the typical trick of writing

```
someBoolean match {
```

```
case true  => doSomething(some_params)
case false => doSomethingElse(other_params)
}
```

What's the point? No clue. To make sure that we always check the false branch?

November 03, 2022

Not much to say. Tried to investigate one case, where SwaggerPlugin is failing; the guy that last worked with it wrote that “it’s not fatal”. So what. But anyway, he’s not interested in cooperating on it, ok. Also working on another case, which is 1.5 years old. Some code looks a bit strange. Like, the guy divides a long into a long, implicitly converts it to Double, then takes math.floor, then converts to long. Wtf, right?

And a couple of similar things. Anyway, these were “just” compiler warnings.

Otherwise, the code is pretty good, actually.

Meetings, just one. A new guy in the team, a devops, pretty cool. This devopsology, I’m totally ignorant in that area.

Anyway, just sitting working on that small stuff. Builds are long, so I have time to listen to Dakh Daughters and write this dybr.

My buddy recruiter is telling me that crowds of developers from Twitter are bombing their company with resumes. The problem is, they all want a million each.

November 04, 2022

Worked on two cases, small problems with Swagger, and disabling some functionality for anonymous users. The reason for the latter is that even if we store the anonymous user’s choices, we won’t be able to restore them exactly due to anonymity. The frontend takes care of this, but the backend is a little bit too agnostic.

The only problem is that half of our projects don’t interop with IntelliJ. Have an sbt file, create an IntelliJ project - kaboom, it can’t find the source, it can’t build in IntelliJ... wtf, right? I’ll have to investigate it properly.

At 7pm my time we had a farewell party, Charles is leaving. Jacob, the manager, was playing mandolin. So, it happens every other week (not the mandolin, just people leaving). Interesting. I guess it’s all about money. I’d rather stay, getting money from other sources. Even if not, we have enough for a humble living. But frankly, I just enjoy working in this industry.

November 07, 2022

Spent the day a) looking for a solution: most apis require both user id and session id (why? who knows), but the api I have to fix only takes a user id. Solution 1: update the api, let the frontend people send me the user id; solution 2: find, how come we can't discover the session id(s) by the user id, or vice versa. Turned out that a user may have more than one session. Not a problem, I'll just scan all existing sessions.

Now the problem is, sbt project does not help IntelliJ to create a proper project, so IntelliJ cannot even create a new Scala class. So spent time experimenting with sbt versions (this one is ancient, and I can't update it, because of those dumb dependencies, so it's another story to file and fix). Eventually configured everything, and basically I'm ready to do some refactorings. Good.

Had a short 1-1 with Jacob; discussed Twitter, Meta, and the expected flow of candidates. So far actually the flow is in the opposite direction: Charles left last Friday.

But I'm totally fine. Nobody interferes. We are good. Even our crazy PM Renato is absolutely nice. The trick is, don't hire Indian male slave drivers, they are horrible. Ok, the company is East-Coast, so they, New Yorkers, will hardly appreciate the psychological tricks of those demigods. Which is good.

Nothing else. Also, tweaking the swagger plugin project, hoping to detect the failures I'm observing in my builds and reproduce them in tests.

So, the honeymoon continues (exactly two months today).

November 08, 2022

Had a lot of fun refactoring and streamlining the services in my specific service. Twice got stuck with tests, until eventually figured that it's just vpn dropping, and instead of running around, I'll need to check at the very beginning (of all tests) whether encryptor is available, and that's it.

No standup today, and I wrote to Karen that if she does not have questions, I don't either, and we can cancel it for today.

Also asked the guys whether it's okay to split a "feature PR" into smaller chunks. Victor gave a good example how he does it. Okey-dokey, so another common-sense factor, who would have guessed. And nobody is telling me to squash my commits, because, you know github knows how to squash.

Just logged out of an "ML presentation", where the speaker said that they are calculating how far is one artist from another. I asked him how they calculate the distance, which metric? And he said "we have a program in Scala, we run it in Spark". Ok, I see. The guy is an idiot. I bet there are tons of idiots in machine learning. No, not all of them are, but...

In 5 minutes we'll have another meeting, planning. So that's it for today. Good night...

November 09, 2022

So, all morning I was working on tweaking the test suite so that if we are off VPN they would tell us, not crash pathetically, and the same regarding aws login - report instead of crashing. Eventually I got tired, so I have the solution half-ready. Bill's Scalatest is a little bit weird in this relation. Anyway, it works, but not as beautifully.

Filed a pr, 26 files - mostly with warnings fixes and test tweaks, but also prepared the area for injecting the functionality I'm working on. That would be next.

Oh, and also, since I'm also in the process of interviewing people, and I'm added to the channel, I looked up what did they say about me. Amazing words. Like "Wow, Kirsha is joining? Can we afford?" (but actually, thank God, I don't care much about earning high salaries anymore; over a year ago I decided to find a GOOD company, whatever the money. I found one, and I feel good. They also were happy for the junior recruiter that dragged me into the company. Me too.

Now I'm looking, meanwhile, into our Swagger Plugin for SBT. Even if we can't use swagger 2.0 (and all that opendoc 2.0), we have to either start using 1.0 (ours is 0.2 or something), or fix the bugs. In any case, I have to get familiar with this shit.

Had a standup, had a planning meeting, and in the evening a long talk with Victor, a director. Funny, I have 1-1s with three levels of management.

Talking with Victor, discussed this Swagger problem (currently small defects, but who t.f. knows what comes next, we need to be in control), and about how the company almost declared bankruptcy about 5 years ago, due to a very similar case as described in Cryptonomicon (Dentist's attack). As a result, salaries here are still kind of low.

November 10, 2022

Karen found an error in my code, I took an id for sessionId. Moral: use typed identifiers. Fixed that, also investigated the code, looking for how exactly session id is being passed. It may be present in the request, but... but, I'd rather count on session service to get all sessions.

The code I was looking through was kind of over-functional. Meaning, a mix of Haskell and Java. No fixed style; parts of a service use Cats, other parts are just plain scala, reasonably functional, and more readable.

People go crazy when they decide that they are functional programmers now, and write code not expecting anybody to understand. Passing around implicits, connecting futures, returning Option[Future[Some[Value]]] and the like. Just some common sense would make it readable and testable.

Oh, whatever. Everything's fine. Waiting for Karen to re-review, working on swagger plugin, and currently I'm in a meeting regarding hiring people. What kind of candidates we want (and what kind we don't).

November 11, 2022

Did not work much. Karen found a small teeny-weeny defect in my pull request (logging can be improved); I wanted to do this in the next stage, but, as an experiment, clicked "accept suggestion" - kaboom, the code updated, tests, pass, need another review. Shit. We have 9 hours of time diff with Karen. So I left it till later in the day and instead did the "obligatory training" - how to do business honestly and legally. How to behave if you are a singer and say something from the scene. How not to give bribes. Etc. Done. While reading dreamwith and facebook.

Then there was a small all-hands, where we were told about our successes, especially in penetrating another reality (virtual one). We are good.

Currently I'm interviewing a pretty cool guy from one of the two famous companies. He's good with Scala, with ZIO, had an experience with Cats; then I gave him a hackerrank problem, and that's where he's stuck. It's been 45 minutes - not once did he run tests. Just exercising in nice Scala code. Really nice. Totally meaningless. Is it qualified as CCD?

Actually, I may be wrong, by the end I asked him to explain his solutions, and he was pretty convincing that when done, the code should work. I would not say was convinced, but... but it's okay. Wrote down my impressions.

November 14, 2022

So, short. Sent a PR for deploying the first stage to staging; it was approved; I deployed, and then could not figure out what was failing. Eventually found - it was the previous deployment, it was failing when it was getting killed. Everything's fine.

Working on the next stage, trying to combine the existing code which is written in CPS with the new piece; not sure what is the right way for composition in CPS. Probably call within call. Does not look beautiful, but ok.

At 7:30 had a meeting discussing that guy we interviewed Friday. I said I'd rather abstain. Chris was in favor of the guy, but three others were rather negative. So eventually voted No. My position was - I'd enjoy working with him, but I would not take responsibility for hiring him. Really, he only solved one problem, probably the simplest. Karen gave him "reverse a number" - given a number, think of it as a sequence of decimals, reverse the sequence, return a number. He failed. In Scala, it's something like

```
def rev(n: Int): Int = {
```

```
@tailrec def step(s: Int, r: Int) = if (r == 0) s else step(s*10 + r%10, r/10)
  step(0, n)
}
```

See, it's a curious case. The guy's Scala level was perfect. But... all symptoms of CCD. It's a pity.

November 15, 2022

Well, I did slap together a decent solution, now I had to add test cases. And shit, I spent all day trying to refactor tests, with all those mocks. Now it's 21:14, and I decided to roll back all the tests, and instead just write one more, my own test suite, to check the functionality. Mockito is unusable, in my case. Partial mocking seems to be impossible to use.

That's it for today. Found another "palindrome problem", will use it with another candidate, a Chinese guy.

Had a short 1-1 with Karen, basically, I don't have problems to discuss, so that's it for today.

November 16, 2022

Really tired of struggling with Mockito. I don't think it's even required, I could just have an actual simple mock in the test. And it's all when I add a couple of methods to the class that is mocked in the test.

Interviewed a Chinese guy. Pretty cool. Good Scala, good approach... well, he did not finish his solution, but he was on the right path. And I saw that he's an avid Scala programmer. So, 90/100.

November 17, 2022

So, I decided to get rid of Mockito in my code. Split that huge test suite into "cases that create" and "cases that don't create"; and started tweaking the suite that does create, adding one more test case when creation is not permitted (for tricky reasons). In that latter suite I kicked out `mock[SessionService]`, and just mocked it manually. Big deal, with partial mocking.

Just finishing this right now, then I'll clean up the code and send for review (my push script will also run tests, of course).

Just had a meeting where we discussed that Chinese Scala guy; I hope we'll hire him, seems like everyone's positive.

That's it, basically. Ok, one test is still failing, but must not be a big deal. The only problem is, WTF I can't run this stuff in IntelliJ? Ok, that's not a priority now; later.

November 18, 2022

So, I was busy all day. Karen found some minor defects (like, a mock service is not used anymore in my tests, so it should have been deleted), and I fixed it all, but waited for Karen to wake up (9 hours diff, she's in LA).

Meanwhile I tried to figure out what can I do with SBT Swagger Plugin - first, how come its build file does not even compile in my intelliJ, etc, etc; and what's the binary, and how do I check if my changes (in build file, in Scala version etc) have an impact on the binary; and what are the existing tests actually testing (we have at least three minor defects in using the plugin, that's why I slapped together a jira case with all that shit)... oh, and yes, we have our own Swagger Plugin, and it may make sense to upgrade to the real one, I don't know.

When Karen woke up, she approved the PR again, and I merged it; also reverted an update where a README.md was changed, and as a result the build fails, because I did not update the official version. That's our build system, с прибамбасами.

Also, interviewed John, who was nervous, his phone was ringing all the time, and he did not turn on the camera - on the other hand, he's obviously a Scala/ZIO guy, and since he's getting an emergency call to roll back some commit, he's obviously a hands-on guy.

Am I too naive? He got that emergency exactly when I showed him the problem to solve: "generate all possible Scala identifiers". If so, good luck doing it offline, I won't repeat the problem.

So I told him, let's stop right here, cancel the interview, get back to it next week, because of your work emergency; he was very thankful.

Told Samantha, and that's it, will interview him next week.

There's a problem with our company, though. Salary too low for twitterers. Maybe we should wait for three months, till the time they are out of severance. Not my problem though.

Then we had a 1-1 with Jacob, my manager. Discussed Mockito, discussed security problems with our app, discussed hiring process, etc.; he asked me whether I'm planning to take any time off. Not yet. (Also, he may have forgotten the employee manual: it says that the first 6 months I can't take any PTO.)

November 21, 2022

Ok, slapped together a PR for deployment to staging (prod is frozen till the end of Thanksgiving days); sent it for review; only in the afternoon it was approved, since Karen is away, and so Victor approved it; and I deployed it.

Now I figured that we don't have a proper framework for integration tests. Oops, fck. I'll have to slap together something similar to what I did at Tubi and Machinify. Ok, ok, tomorrow.

Spent some time with our two swagger plugins. Came to conclusion that it's us who actually maintain this plugin. Our company. Surprise, surprise. This should be me, then. The problem is, I've a problem with jfrog. I mean, intelliJ, in our installation, has problems with jfrog. Maybe I need a professional version. Let's see what Victor thinks.

Also, fck, both our swagger plugins behave strangely in tests. One just does not run any tests (they exist), another one crashes in tests, and again, it crashes thanks to jfrog. Ok, it's my responsibility now too.

Tomorrow interviewing some guy, will give him this "generate Scala identifiers" "problem". And the next one will be "generate an html table of all those" Greek words". It's not a puzzle, it's a test for writing clean code. CCD or not.

That's it for today. I love my job.

November 22, 2022

Well, since deployment to prod is blocked until Monday, I decided to do other things. First, do we have integration tests that talk to staging server? No, zilch. Oops. Ok, there is a project that has some kind of such tests, written by a qa, a long time ago. That's it. Bad.

Then I tried to ping people regarding swagger plugin. Ok, I see that it is probably my responsibility to take care of it and maintain it, it's opensource, and it was being maintained by our company. Interesting. But so far I'm having problems configuring jfrog access properly. Karen pointed me the files that may need tweaks. Ok.

Then we had a "standup", then we had a "planning", where I signed up for a couple of new cases.

Now I have more meetings ahead: 1) retrospective, 2) interview a guy whose resume looks like brilliant paula bean (will fry him on his variety of claims, from perl to scala), 3) a large demo.

So currently I'm just having tea and taking a break.

November 23, 2022

In the morning I started with trying to slap together a script that checks whether we are under vpn or not. After several attempt wrote something like curl -i (some url); it returns 6 if I cannot reach the url - and threw it to my build script(s). Big deal, right? Whatever.

Then I played with the build for the project I basically finished, to see if it will work with the modified jfrog config in .sbt - strangely, it did. That's weird. Because sbt play swagger build still fails. Weird. I'll get back to it Monday.

Then I looked through the two new jira cases, but not much.

Because we were in Emmaüs around the time I had to come home for a standup, I almost ran home. I was not late, but the meeting was converted into a slack report. So I did, reported.

Then we had a meeting discussing John, a Hispanic guy, and Bancroft, a British guy, the two candidates. Turned out that John's interviews were not finished, after I interrupted it, after he got a work emergency - so we have to get back to it when done with John. But regarding Bancroft, funny. Davis was strong positive, since Bancroft knew everything in Scala. Interesting. I was strong negative: the guy pretends to know regexes, but he did not; and he did not write anything that was producing anything like "Java identifiers": he just did not get the essence of the problem. But he used familiar tricks from a ZIO course. Which does not help if you can't code.

Then Karen was also mildly positive regarding Bancroft, except that Bancroft was rather vague about SQL.

Then Victor, who was in the passenger seat of a car, heading South (judging by the shadow on his face), with a bunch of kids in the back seat, said that he's mildly negative, because, well, Bancroft does not understand SQL, and in general, he talks much but does not understand much.

Then Jacob was also mildly positive: a nice guy (I agree, partying with him would be fun), but well, he was a bit weird. Jacob asked him about his "negative sides", and he said that he rarely shows up at the office on time. Interesting. It's not a problem per se (my experience tells me that the best programmers come late and stay late), but we are online now. If he's in the Bay Area, he's supposed to join a meeting at 8AMPST. Anyway, that part is not my business. But also he lives in a trailer, and his internet connection is just pathetic. How is he supposed to e.g. do git clone or sbt clean test? Some projects are large, and dependencies are just huge.

Anyway, we decided to get back to him after we review other candidates. Obviously that smart Chinese guy from Twitter decided not to join. Money, of course.

Anyway, I'd wait until February, when all these Twitterers, Amazoners, Facebookers and maybe lots of others discover that nobody needs them anymore, and probably nobody will be paying them \$500k anymore. Who knows, though. But startups are on the freeze until 2024, right?

And also, since Victor wrote to everyone, on slack, that since 1PMEST the working day is over, and since our code hound is down (for the obvious reasons that SREs want to do their stuff), I closed my work computer till Monday.

See you Monday then. I will have to deploy my shit to prod Monday morning... first will have to check if deployment freeze is over by then.

November 28, 2022

Not much. Created a PR for deployment. Planned to do it Wednesday. Seems like I'll be spending Wednesday in airports. So... so, nothing. Too much nerves, just can't focus on work.

Told the colleagues about my troubles with the flight back home.

Basically, Lufthansa canceled our flight and moved our departure to 30th, which is beyond our 90-day stay limit. They notified expedia; neither expedia nor lufthansa notified us. Spending the whole evening on the phone(s) with these two entities, trying to avoid a big fine and/or travel ban. Fuck.

November 29, 2022

Yesterday, in the list of candidates, found the name of Prasad Mokkapati. This is the person whose C code (for Visibroker) still impresses me by its quality and impeccable style. I was amazed by the opportunity, wrote to them that we have to do all we can to grab him, while the stars are falling from the sky into our hands.

Well, and what, Karen responds: "well, he has a lot of experience as an architect and a founder, maybe we should hold him back, until we find an appropriate position..." Yeah, right, as if he's going to wait for our next move.

But it's an interesting signal. People are afraid of someone who's too good.

I observed this from time to time, in various companies.

November 30, 2022

Today I did not do much, thanks to all kinds of troubles. Wrote to them on slack.

Still, got a decent size time slot, later in the evening, when everything settled, and deployed "collection service" to production (5 different namespaces, each for another country. Watched the logs in Kibana; my stuff went OK, but saw a lot of errors reported for a bunch of other services. Maybe I should create a jira case for all that (not assigning it all to myself, just to keep the information).

That was it. Oh, no, also, while in the pharmacy, did run a project to see what's wrong with joda time transformations in a json parser - mine reports some missing classes, but that project has everything fine. I guess it's project.sbt that is incorrect in the module I was working on, but without working samples I just could not figure out wtf is going on. Will investigate.

But anyway, I'll have to start another "production jira case", and do that swagger problem in the back burner.

Also, there was a pretty funny discussion with Chris regarding the problems I was asking interviewers. Obviously he could not solve a problem like "can you write a regex to check parentheses parity in an

expression”, and also interpreted “you can’t use regex for parsing XML” as “XML has a lot of entities and specifics of XML, and regexes can’t take into account all the rich features of XML”. I tried to explain the levels of language hierarchy (btw, invented by Chomsky, afaik), but no success. I have to be careful with this guy. Finding a new job these days is not as easy as last year.

December 2022

December 01, 2022

Spend the morning looking through a bunch of project, trying to find which one gets the “streams” request. Since what I thought was doing it was not. Eventually wrote to a channel on slack, asking if somebody knows. Christian turned out to know: it’s not even in our repository, it’s a service provided by another team, and the code is another repo, so it’s not my fucking business at all. Cool! Updated the jira case, assigned it back to whoever created it, with comments.

Meanwhile another project, where “Joda DateTime” jsonification code compiles properly, I’m going to find what’s so good with it. Remember, all the people that knew are gone now. It’s a little bit like Twitter in 2023. They could have left a backdoor, nobody would notice.

Had a standup, where I could not connect my headset, so I was just waving and smiling and showing the thumb, like teletubbies or like lapins-cretins. Only later connected this new headset with the silencer and made sure that it all works.

Then there was a demo of Snyk. Is it a Ukrainian word? Could not find it via Google. Anyway, at least a couple of Ukrainians were presenting. But... but, when I asked them if they can demo how it works with Scala code (it analizes the code for security holes), they said they’d do it later - but they did not. They showed some cases in other languages; it’s beautiful, but what kind of problem is it when a file is left unclosed - so what? We may be a little bit short of file handles, but it’s not Rust anyway, right?

Anyway, a little bit impressed, but I’d see more before judging.

Then there was an interview with John. Although he’s John, he’s Hispanic (ok, probably Mexican, judging by his accent). We already tried to interview him a couple of weeks ago, but he was busy. Now I think he was not. Also, he turned on the camera, but never looked into it. I guess I had to do the same, look to the side, etc.

Anyway, we just started, I was trying to reformulate the problem I was giving to him last time, when his phone rang, and he started talking to his brother, about their father. I did not have a chance to tell him that yo entiendo español bastante bien, when he hung up, and told me that his brother only calls him once a year, when it’s his birthday; I tried to congratulate him, but he said this is not the case.

In short. First I gave him “generate all Java identifiers”. He understood the problem, but said that it’s too hard for him. Ok, I said, given a 4-letter alphabet, “abcd”, generate all possible strings out of this alphabet. No, too hard. Ok, I said, generate all possible 7-letter words out of this alphabet. He wrote “abcd”.toSet.combinations.foreach(println) I noticed that what he printed is 4-letter words, I need 7-letter words. He added “efg” to the string. I asked him, if English language has 26 characters, should all English words be 26-character combinations of them? (The fact that a programmer does not know what an alphabet is, for me, is a signal that the guy is an idiot.)

He eventually understood the problem, but said he can't. Too much for him.

Ok, then. Do you have anything in github, any project that you can show? He did not show anything, except for a link. I asked him if he has any question? None.

So I went to his project. It was the standard task from ZIO workshop. And obviously too much for this kind of guy; he just copied it from what John de Goes showed the audience. 8 months ago. That's it.

These people, what do they do, in general? They talk, condescendingly.

Anyway, I wrote to our recruiter, Samantha, all that I wanted to say. But politely.

Well, and I missed the security design meeting that was happening at the same time today! Such a pity.

December 02, 2022

I'm not feeling well, brought some European virus (or Washingtonian one) home, and so there. Working from my bed. Took me some time to figure out wtf I actually should do. The case that I reassigned to Mike, it's actually my case. The thing is, there's an API, and jobs use that API, but they are not aware about updates in the records they receive. So I found the service that produces that shit, and added the field, and, eventually, wrote a test. This thing will be "deployed" (actually, go to a library), and then it can be used. The client side code ("jobs") that pull the data does not care about the structure, so there's nothing to test there or even add.

Also, fixed a bunch of comp warnings etc.

At 3 we have an emergency meeting. We got a new "Senior VP". With a community college education. Кухарка будет управлять. We will see. Who fucking cares, actually. Unless he starts introducing mutual hatred, like in California. But he's in Georgia, so we will see.

He also confessed that there was a small layoff, in "data science" team. Nothing that has, is. Christian Data Science. Also he said that last year there was a 30% layoff. Which explains how come the majority of the pros are gone.

Anyway, by the end of the day I slapped together a PR, 13 files, one with the functionality, one with a new test case, and others with warnings fixes. Mostly the fixes are in arithmetics, like `val n = Math.floor(a / b)`, where `a` and `b` are longs. Not a big deal, of course, but the compiler WARNED.

And again, I'm not feeling well.

December 05, 2022

Now this pathetic Covid patient is trying to pretend to be working. Since my PR was ready by Friday, sent it for review, and that was all I could do today. During the meeting we were all getting introduced

to a new guy from Columbia, Wilfred, or Wilson, something like this. I have a brain fog, that's an excuse. As to the work itself, Renato is asking me whether the stuff will be ready by Wednesday. It is ready, but we need to deploy it. I was too dumb in the morning to split this task in two parts: one being "make it", another being "deploy it". And I had also to explain to him that we don't have tools to test whether it works, so we should be optimistic and keep hoping.

Wrote to them on Slack that so-and-so, I have this popular disease. And closed my work computer.

But later we had a 1-1 with Jacob; it was pretty heart-warming. He had Covid twice already; now his parents are sick. Isn't it horrible. So, we are kind of lucky, actually. Discussed with him the issues of swagger maintenance. And the issues of interviewing. He wants to create a list of interview questions. That's a good idea. Sort by topics and by complexity.

Then my brain fog got thicker, and I opened my work computer again around 6:30. Nobody reviewed my code, so I wrote to them again, what's up etc.

Victor writes that yes, brain fog is think during this disease. Two Indian girls (probably one of them is a Tamil, actually), they both wished me well. You know, I adore Indian female programmers, they are the best. I have a bunch of them in connections on Facebook and on LinkedIn. They are the smartest (much smarter than males).

That's it for today. Still a dense brain fog, I can't think.

December 06, 2022

Using my excuse as an excuse not to do much of anything. So, yes, I did some stuff, deployed something; Victor found that another project should be fixed as well (agreed to meet tomorrow), and... and we had three meetings.

Nothing important, nothing serious, I'm still too dumb and slow. Got an impression that I'm just lazy. No drive.

December 07, 2022

I have an excuse for being lazy, I got covid.

So, had three meetings today; one was about sprint grooming, and included looking into the Jira cases that SNYK automatically creates. What it creates, I think, is just a bs. Complaints that our sbt build refers external libraries which may be dangerous. Anyway, worth looking, but even worthier asking SNYK people, are they sure it's a bug? Maybe it is. But if it is, they should formulate it properly.

My "case" was split into three, since Victor said there's one more project to be updated (if I know what exactly).

So I planned to talk about it with Victor, but instead we dedicated the whole 1-1 hour to the company progress, to interviewing process, to using hackerrank for creating custom questions (that I will do), about Will that left in August (that's a big pity for both of us).

And I decided to investigate the project myself, but oops, jenkins on staging is just down. So I wrote to the guys (or how they say it in New York? Dudes?), and will now play with HackerRank questions. Will just go get a cup of tea and proceed.

See, I've been lazy lately.

December 08, 2022

It's still a bit hard to work - lazy, foggy, not much, but.

Deployed to staging that service that I committed yesterday; then looked into wtf am I supposed to do next to make it all work. There are a couple of stages. Turned out updating "records" service means just referring the latest deployment, so I did; sent a pr, it was approved, and I've just merged it. With long long breaks between. Will deploy it tomorrow. There's one more step, will look into it later on. And there's HackerRank problem I've been trying to create; so far I'm confused.

December 09, 2022

Ok, dream hardly working today; I was also in a brain fog, tried to deploy to prod instead of staging. Eventually deployed two services ok, but there's a job to deploy; will do it Monday, it's better.

That's it for today. Hope my brains will be Ok by Monday.

December 12, 2022

So, my plan was to deploy a couple of services to prod, and see what's up with the job on staging that only runs the Australian version - and how even this job deployment happens if there's no config for that.

Ok, prod deployment, I decided not to do it until I check what's wrong with the jobs.

Turned out all staging deployment for that job fails, except the Australian one. And why is it? Because some resources used in deployment show 404. Missing. WTF? Asked people. It turned out, our colleagues had deployed something. Something that serves the resource perfectly when requested from a browser or curl, but denies its existence when called from Jenkins.

So, I rested my case, pinged as many people as I could; also, updated a README and slightly tweaked the error handling in that job code.

That's it for today. Karen is on the West coast, so let's see if she can comment or find our brave colleagues that broke the code.

Oh, and also I've slapped together a hackerrank problem that I love. I mean, I added a problem to our list of problems. Cool.

We were supposed to have a 1-1 with Jacob, but neither had anything to discuss, so we are cool.

December 13, 2022

I can only explain it by my Covid. So, what I deployed is not enough. Jobs are deployed, but there's more, something to be deployed to Jenkins (and it will deploy the jobs). Ok, documented it, but not well enough, so still working with Karen on this.

Also, quit my today's attempts to deploy recs service to prod, because, strangely, it's not there... was not it there? Seems like it was.

So, spent tons of time on all that undocumented deployment nonsense.

Also, there's a script that you are supposed to run before pushing configs. The problem is, the script always fails. Looking into the script, I see that it should. But the first part (before crash) does something, which I would not say I know. And I can remove that crash (it checks a condition which is always true, and probably the condition is wrong). But then there are weird steps that run for about 10 minutes, and I have no clue what it is. And nobody knows either. So, should I investigate? I don't know yet. Even the first part is a mystery as a script - but the results it produces (the ones we need) are obvious. What else is there? Who knows. I think I should work on those scripts.

December 14, 2022

Ok, the services are deployed to prod. Now I have to deploy a couple of jobs, for all the locations. Turned out one version was wrongly specified. Had to roll back. So I did. Github just hangs on that pr build. So? So sent to Karen a PR whereby I just update a README. We've been discussing this README for a couple of days.

Anyway, bored. Renato asked me if I'll be able to deploy all to prod before Dec 21 (freeze date). I said I hope I will. But our process is crazy. And almost nobody knows wtf is there. Karen knows. I'm documenting what Karen knows.

Last night I was brain-foggy, asked the people how I can deploy a service to prod. Karen wrote: "it's in the document you wrote". OMFG. Sure it is. Just brain fog. But anyway, that's how we are. Bad.

December 15, 2022

Half of the day I was kind of working on stuff. Deploying the remaining jobs to staging, figuring out what's up with versions mismatch, and the like. Also, discussing with Wilmar (a new guy) wtf is going on with Nix, which he's trying to set up on his machine. Seems like the problem is that his machine is an M1, and kaboom, Nix throws some known errors ("null is not a string") on that machine. Meaning, I guess, that some property is not getting retrieved (it's good that null is not a string, in Nix language).

Then I got suddenly sick again, closed the discourse, and only watched how Wilmar and Qiang discuss these issues.

So, that's it for today. Tomorrow will ask Karen how can I check jobs on staging (she sent me one URL with the data that we are supposed to get through, but it's not obvious to me yet, especially with Covid).

December 19, 2022

So, today is basically the end of the year. Deployment freeze tomorrow.

I had two prod deployments left. One PR was already posted yesterday; it was approved today, I deployed it. Worked Ok.

Reassigned my jira case regarding Nix to our Colombian guy, who is a Nix expert, relatively. To be more precise, reassigned, and asked him if he can handle it, otherwise he's free to bounce it back to me.

Slapped together another PR, for the last deployment, sent for review, it was approved, and I deployed it. Mission accomplished.

Then we had a "standup", I did not have much to say.

Then we had a 1-1 with Jacob. First, we agreed that we may not even need Nix. It's not being used in Jenkins, and besides, it's crashing on the Colombian guys' computer which is M1. Nice! Neat! In short, we may need to reformat the case for the Colombian guy: investigate whether we fucking need Nix.

Also, Jacob said that he got connected with my favorite recruiter, Josh, so Josh will be supplying us with good engineers. Cool!

The rest of the day I spent documenting this "deployment process" for these "import jobs", the info that only Karen knew by heart. I don't know anything by heart, I just got rid of Covid and my brain is foggy. I record everything.

So, the remaining days I'll be trying to figure out how to handle our swagger plugin (and fix bugs there).

December 20, 2022

Just a list of actions - wrote to Wilmar regarding using Nix, asking his opinion, he said: who cares - wrote to the team asking their opinion regarding not using Nix; two people were happy - wrote to Wilmar that we should just close the case and move Nix to a freezer - closed all my cases in the current sprint - fixed a build failure ("a build with this version number already exists") - started looking into Swagger; still can't figure out where's everything

December 21, 2022

So, what was I doing? Ok, formally, I was supposed to work with swagger, but I did not. Instead, closed the jira case that I assigned to Wilmar, since we don't fucking care about Nix anymore. Then updated the docs (pulling out all mentions of Nix to a separate file). Still working on this PR. Also, added a build.sh to docs project, since building it is not by sbt build or something, and errors were reported in jenkins but not locally. Fixed that.

Had a 1-1 with Victor. In short, he's leaving. He got an offer, and our company failed to make a counter-offer, so there. Sad. It was great working with him. He'll still be around for a month. Talked about this and that and Scala, etc.

1-1 was about "what are we going to do next two weeks". Nothing!

December 22, 2022

Not much happened. Karen reviewed my docs update, gave pretty cool advices, I updated the docs, will deploy a bit later today. Then we had a "party", where, by Karen's suggestion, we played some online game; funny, but probably not for my age, so I played just once.

Also they discussed the price of visiting Disneyland, where you waste as much money in a day as you could spend staying for a week in Costa Rica, all-inclusive.

I have to force myself to open my work computer again and do something meaningful today.

December 27, 2022

Today is a fake work day. People just pop up in slack, "reporting the progress", which consists mostly of "reading docs".

Ok, and I did what? Updated some docs, exchanged opinions with the Columbian guy, regarding Nix, and started looking into wtf it is with swagger plugin. The freshest version in the opensource is 6 years old, depending on sbt 0.13.9, which I'm not sure if it still exists. Anyway, wrote to the guys that

maintained that plugin 6 years ago, and wrote to rssh, who forked it 6 years ago. If I don't find a fresher version, I'll have to work on what we have. Etc.

Also, the HR wrote that K.L. filed his resume. Nice. I wrote to him several days ago, advertising our company and suggested him to apply, but I asked him to send his resume to me. Anyway, we will see. People are looking up at him. Why so, our codebase is pretty decent. The architecture can be improved, but I'm not ready yet to interfere.

Btw, one loquacious participant is now out of this blog, for personal reasons (really, he compared me to Hitler for not disclosing my sources, I suspect he may have mental problems, but is it my problem?) - anyway, I know some people read it silently. It's okay.

December 28, 2022

Pretty much everyone is kind of on vacation. I half-worked too. Got some info from Wilmar, regarding Nix. Will send it to the Nix guy. Also, figured wtf is with our sbt versions. A mixture of different versions.

So I created a jira case regarding unifying the versions (e.g. to 1.4.4), and started experimenting with swagger sbt plugin. It's a horrible mess. Basically, I could not find, on github, any sbt plugin project that I could clone, build and test. They all fail. Including those from tutorials. Mostly because the Scala methods they use are deprecated.

So, not sure yet what I can do. Probably try to grow the shit from scratch. From one of the simple samples, but from scratch.

That's it for today.

December 29, 2022

Not much was happening these days. I was struggling with sbt swagger plugin. My last decision was to start a new project from scratch, and gradually modify it to make it a swagger plugin project. So far so good. The things that IntelliJ does not recognize in the old project get accepted in the new one.

January 2023

January 02, 2023

Ha, I was sure it's a working day today. So, did some stuff with Sbt Swagger Plugin, then, at 11, went for a stand-up. There were two of us, me and another guy. Nobody else. So we just said "Happy New Year" and closed the session.

Then I was supposed to have 1-1 with Jacob; wrote him that well, do we need it - and he wrote back that it's a holiday today. Nice!

January 03, 2023

Feeling weird today. Tried to work on that plugin, tweaking dependencies, turning managed one to unmanaged, etc - but well, feeling weird, dizzy, etc. Probably it's something like covid again.

At the meeting we just exchanged opinions, and all other meetings were canceled. The Columbian guy asked if it's okay to upgrade sbt to 1.4.9, he was, seems like, only interested in my opinion. I said "sure, go ahead, but we mostly have 1.4.4 in our projects (with one exception for 1.4.3)". So, in a sense, I approved the upgrade. Feeling weird.

But I did not open my work computer after that. Just feeling too dumb now. Watching movies, doing French, sleeping, looking at carpenters raising high the roofbeams.

January 04, 2023

Well, a bit of brain fog, so I even missed the second meeting (and the third did not happen, Victor is too busy).

Agreed to pick up a case, or, rather, an epic, as they say in Jira. It's about changing the way we are handling a stream of updates (they seem to be coming in as a stream, then stored in a db, and only then get visible in the app - wrong!). It was Victor's epic, but he's leaving in February.

The Columbian guy found a bypass for the weird case of Nix not working on his M1. Cool!

We have a new guy, Mahesh. I'll have to check out his resume, it's in Slack.

Also, a bunch of new resumes: former twitter and facebook engineers are flooding us with their applications. Some are pretty cool. One is Russian, from SPb; it took me some effort to block my, you know, attitude, and look at his resume professionally. Looks decent.

By the end of the day I was pretty desperate in my attempts to just build sbt-javagent (I only need its doc jar). Sbt version, JVM version, Scala version, what else? Who knows. I was tweaking the project,

no good. So eventually I've cloned it on my own computer - and everything goes fine here. Omfg. So I'll have to find an env var that is wrong. Tomorrow, tomorrow.

Wrote on slack about it, maybe Wilmar knows, or, I don't know, Karen, Victor...

January 05, 2023

Had three meetings, discussing things, problems and solutions. Basically, currently a bunch of data are getting to us via an integrator, then an api dumps it to kafka, then a job scans kafka and runs something in spark to stuff it to mongo, then another job eventually prepares the statuses to be sent over to clients on demand. Takes half an hour at least. No, we need to make it a stream, of course.

So I have to study the whole architecture, so far I was not much aware of how the things are connected.

Today Victor gave a demo, drawing pics and explaining the flow. Not documented, of course.

Seems like my job would be to figure out and find a streaming solution. Ok, ok.

But then, after lunch, our docs site went down, for some reason. Ok, I already raised a question, why t.f. do we need a docs site, docs build etc, if we already have a wiki site where other teams store their knowledge. Let's be simpler. As simple as we can.

Anyway, since that site was dead, I kind of asked Safa, wtf (in nice words), and just closed my work computer. Opened it later in the evening - silence there. Ok, ok. Tomorrow. Mañana. Manhã, tão bonita manhã de um dia feliz que chegou.

January 06, 2023

I would not say I did much today, and mostly because I was looking through our docs, and could not extract any meaning out of them. They were written about 7 years ago, mostly in future tense, and they describe some plans that were abandoned probably five years ago. Anyway, by the end of the day I found a collection of videos; it's too late to watch them now (switching my headset is a pain, so... tomorrow).

Victor was giving another "knowledge sharing session", mostly about how to build and deploy. I reminded that I have it all documented, and the noobs are wholeheartedly welcome to update the doc as soon as they find something missing or wrong. (So far Wilmar did not bother to do it, though.)

Then Jacob scheduled a 1-1 regarding how do we document what we do. Right. No picture, who knows where the information is coming from and where it goes. No big picture. Maybe Jacob and Karen know. And Victor does, but Victor leaves in February. So, I got a feeling it must be my responsibility to keep and maintain the big picture, kick out unmanageable stuff like Nix (unless Wilmar bothers to maintain it), and keep things simple and safe. Because if something crashes, we are like Twitter now. Like Salesforce. Like all those MBAs left without developers.

At 2pm there was an all-hands, where there were like a couple hundred participants, everybody was excited, everybody was looking forward to a big progress, but I don't know who is going to work on all this? That's why Victor left: his negotiations with the management did not work out. How do I know? He told me.

So, interesting, what next. I'd say, we have to divide the work. Karen would not, Chinese people, I never saw a Chinese Engineer willing to divide the work. Different mentality, nothing wrong with it, just a different mentality. And of course everybody can ask me a reasonable question: who the fuck are you? And my answer is: nobody. Just curious. Just someone looking for the big picture.

January 09, 2023

So, not much work done. Tried to watch Victor's videos, but hardly had any chance. Then missed the standup. Then had a 1-1 with Jacob; agreed to have a "leadership meeting" some time within a month. Then Renato told me that swagger's page with APIs was not updated with the new code. Oops. I was not changing the API, but the structure was updated, I had to update the swagger docs. And I forgot how to do it, and my docs were not good enough to remind me.

So I asked Karen; she told me; I'll update the docs. But the problem is, swagger docs were not updated on the server. Now I know why. I had to update the version in the config file. And update the docs that say I have to do that. Fck. Anyway. I was busy and tired. Will fix it all tomorrow.

January 10, 2023

A little bit of Covid, so there. Watched a presentation online.

Renato asked me how I'm doing with this API update in Swagger, he did not see anything. Well, I had to read docs again, because I was not sure. But I just deployed the docs yesterday, to staging. But kaboom, the API looks fresh. Mission accomplished, in a mysterious way.

Then we had a 1-1 with Karen. Basically, just walked through some difficulties of our code base, and exchanged smiles. Everything's fine.

Then I worked a little bit - watching videos, etc. Watched Joe, the legendary genius that created all this complicated paradise. Not a master of speech, but ok.

Then Victor was giving a show, knowledge transfer. Kaboom, he showed the "main architecture picture". It is there! So, Victor's talk was very good. And I learned some abbreviations, too. E.g. NLB.

Right after that show, a 1-1 with Victor. I had two questions: 1. wtf with our swagger plugin, and how do I configure the contents of our jfrog. Got answers. 2. What was that about switching from packet processing to stream processing of stations updates? Oh, it turns out that the providers of that information want to use push, while currently it's pull. Currently we pull it 4 times an hour. Better get

it sooner and faster. Only problem is reindexing elastic search. But that's a side issue. So, I'll create cases in Jira for all that spike that I was working on, and then let Renato do planning or whatever.

Then I had to do some, well, social stuff on wikipedia (fix it after some idiots). Then my Covid turned me into an idiot or something. Done for today, good night.

January 11, 2023

Frankly, just out of luck today. Renato sent me an example of a curl call that does not return the field it's supposed to return. But I could not do anything, because our wifi is broken, and my work computer hardly connects through the phone, FIIK why.

Even slack sucks, but not on the phone.

In short, I participated in meetings, but could not do anything meaningful. They'll fix the internet tomorrow.

January 12, 2023

In the morning was not doing anything, because the quality of my phone wifi...

Then the internet was fixed, and I got back to work. Basically found the cause why that field was not passed around to mongo and back. Ha, it was not mentioned in the transformer. Now, why t.f. did all the tests pass? And I discovered that the appropriate tests are just fake. I tweak that field, making it totally corrupt. Tests pass.

So, I'm looking for the tests to update. And discovered that intellij, with a project containing several modules, has been misbehaving since about 2012. There is some Sergey Baranov in Jetbrains, who is answering all error reports with some stupid shit. I wonder if this is Sergey Sergeevich Baranov. His father is totally lunatic. Was one already in 1991. Oh, whatever. I'm getting more and more disappointed in IntelliJ.

So, but I'll slap together a proper test, and file the pr, hopefully later tonight. Customers are nervous.

January 13, 2023

I wanted to start designing the new API, but Mike, instead of agreeing or objecting to what I wrote to him, suggested to write more APIs, but first chat on MS Teams. Ok. But I missed the meeting I scheduled - I was in traffic, so we agreed to meet Tuesday.

Also, merged and built the bugfix, but there were some issues in the build system... in short, the stg deployment is probably delayed till Tuesday too. Sent the request for review, though. Will check a bit later today, it's still daytime in California.

Also, almost missed Victor's talk, because, fuck, as usual, it was not in the calendar, and suddenly he started, and kaboom, I was not there.

Anyway, will watch the shows later on.

That's it for today.

January 14, 2023

It's a weekend, generally speaking, but since last night Karen approved my PR, I just deployed the service to staging.

An our later I figured that it was not enough, and I also have to deploy jobs. Oops. Tomorrow. It's a weekend, after all.

January 17, 2023

Had a bunch of meetings and talks. Kind of agreed on the APIs that I should specify; I'd have written the doc, but there's an urgent failure which I was investigating. Yes, with the code I've been working with.

It turned out that there are two processes, one is the service, another is the job; both do the same thing, but the data structure they use, which is supposed to be the same, is declared independently, and is different. And yet another, under the same name, is declared in another service, but that one is different, so I just renamed it.

So, instead of writing specs, I'm bugging our partners (from a different dept) regarding what exactly are they sending us (or we are sucking from them) - that structure is supposed to be the source of truth, but we don't even know where it is documented.

That's how I ended this day.

Also, thanks to Karen, I now know that what that team calls "playlists" is called "collections" in our codebase. Fucking weird. Why not have a glossary? Or a thesaurus. Will have to talk with Jacob about it. Microsoft had, a while ago, a book of terminology. So should we. Just a page.

January 18, 2023

I don't know, a special day today, a meeting after a meeting. Standup - explained again that the cases that I'm doing are somehow not listed in kanban, and explained again what exactly I'm working on: fixing the data structure in import job, and preparing a document to specify the actions to implement the switch from package pull to stream push.

Then Victor was giving a talk, “knowledge transfer” - good talks. I’m downloading all these videos.

Then we had a sprint planning meeting, and I again explained what I’ve been doing, and how many sprints the second part (pull to push) will take, agreed on sprint and a half; it also turned out that Renato’s understanding of what “point” is (all these people are idiots, I believe; the scale is logarithmic, but it’s too complicated for them) much smaller than what I saw everywhere else. NP then. I’ll give it all more points. Point inflation, probably. Who cares, right?

Then I had a 1-1 with Jacob (my manager), and we talked about ensuring that we are in control of swagger plugin, and figuring out what refactorings can we do, and can we do more ZIO in the code, and maybe even switch to Scala 3 (dreams, dreams, dreams, what Scala 3 if we have Spark).

Then I had a 1-1 with Victor, it’s his last week here. He explained to me that shit with the history of incompatible data structure definitions, so I decided, fuck, I’ll fix this one. Then we chatted for about an hour, on this and that topic, and he said that if we have questions, we can always contact him. Weird. Nice.

Then I started refactoring, kicking out the record definition duplication, and injecting explicit dependency, and then trying to make sure the tests pass (they don’t yet).

And then it got dark, and, well, time to do my French, right? À demain!

January 19, 2023

It’s kind of hard to find time to work.

Was figuring out how the data transformations in Jobs fails on the data that are ok in Service. Found eventually, there are some nulls inside a map inside a json. Well, a regular shit. So I have to update the service first, then use it in jobs as a dependency. Oh, whatever.

Also, wrote a design doc for “pull -> push” refactoring.

Had 3 or 4 meetings; at least one was good, when Victor was telling about, well... about this and that. Mostly about legal issues, which are really funny in the US.

Also, had a “roadmap” meeting, where we voted on this and that feature/fixture/improvement. People are enthusiastic about switching to Scala 2.13 and about starting using Service Mesh, of which I only heard today, on Facebook, from brilliant Oleg Tsarev (and then went ahead and read the docs).

So, otherwise, everything’s fine. It’s 7pm now, and I’ve closed my work computer.

January 20, 2023

Ok, my first PR was approved, and I merged it - it was about getting the missing functionality from one project and stuffing into another. Pushed it. Then I worked on that other project, making sure it

has dependencies, not copy-paste, fixing the tests and getting rid of warnings too. Tests are not very good, but well, fixing them would take more time.

Also, had a “slack standup” - you just write what you have.

Mike was happy about me having written a design doc. Seems like chatGPT did not penetrate our realm yet, so anybody who can use English is kind of respected.

Then Victor was giving another talk on user authentication etc. Pretty cool.

Then we talked with Karen on how, now that Victor is gone, will we test the candidates’ level of Scala. No big deal, in my opinion, but it is not exactly relevant, anybody can learn Scala. We should only check if the person claiming good Scala does have it. No big deal, again. Variance, loops, monads, implicits, why not vars, Options, nulls, parsers-combinators, Nothing/Object/AnyRef/Any, and the like.

And then we had Victor farewell party. Mostly the veterans (Victor was there for 7 years) remembered the old days, showed pictures of the parties in their NY office, Odersky coming over with a talk, Joe as the brightest star in the sky, etc. Now the office is mostly empty, hardly anybody stops by. If Jacob organizes our meeting there, several “senior” engineers coming, that would be cool. Some time in a not so remote future.

I actually finished that other PR around 6 pm. There’s another PR, but it’s not urgent, it’s just for the clarity of classes naming (getting rid of duplication).

So, feeling strange that I’m not slow, but I’m rather fast, relatively. Ok, until they hire some Hindi asshole, we are fine. There are good Hindi people, but they mostly stay in India. Well, I know several good ones in the Bay Area, and one in NY. But the Tamils and the Malayalis, who are nice and cool, I never met a bad one. Call me racist. I don’t know much about races.

January 23, 2023

What a day, seems like I’m doing a typical devops job, deploy this, deploy that, bump version in this config, bump it in that config.

And since I’m not a friend of nix, it just got me. There’s a script (in ammonite) that bumps into Scala 2.13 that I have, and Scala 2.13 does not like the old style, so I had to fix that script. Now there’s an ‘import’ that is not working in that script. Anyway, I bypass the script, but have to fix it eventually. For that I’ll have to investigate how amm files get their dependencies. From an ivy config? Oh, whatever.

Oh, and looked through an interesting resume, with the introduction probably written by chatGPT. How he is socially responsible, how he wants to work with a company that shares his concern about the future of tech and ethical programming practices, and how he, while being a biology student learned programming because he was a Student Government’s IT director.

And the languages he knows are Scala, Zio, Java, Javascript (React, Node, TypeScript), SQL, NoSQL.

And how he worked in Scala 3 as a senior engineer. And how he “Generated REGEX scripts to validate input fields and fulfill security requirements”.

I looked him up on LinkedIn. A totally different story, and a totally different resume. No “senior”, no “backend”, just some frontend, pretty cheap.

Anyway, wrote to our HR person that this guy is obviously just fake.

But the idea! One resume on LinkedIn, for the people who know him, and another in private applications.

January 24, 2023

Pretty dense day. I was supposed to deploy all those jobs that use the updated record shape and import the missing data. Deployed eventually. Then looked around and got a suspicion that that was not enough. Talked with Karen. She said no, it’s fine. Then I figured the code I deployed has nothing to do with that import. Then Karen managed to convince me that it does. Ok, it does.

So, today I’ve learned how to search for texts in Kibana, and how to query MongoDB online, and sucked in all the data that we are pulling every hour, and started wondering. They give us 2700+ records. We somehow store to MongoDB 2814 records. And the field that is supposed to be there is still missing.

So, do we really pull from the same url that I used to pull? Do we store in the same Mongo that I used to browse? Or wtf. The code I work with is pretty good, and covered with tests. And it does not work.

Also had to explain to Renato and one more guy... Rabindranat? Some long name - that I’m working on it, and it is not working yet, and deploying to prod makes no sense (that was Jacob’s point, don’t deploy untested code, good point).

Also, did not have time to struggle with Nix/Ammonite config, but that’s second priority.

Oh, and found time to edit my doc regarding switching from pull to push.

And I’m wondering why t.f. should we pipe 2700 records through Kafka if I can e.g. read it all via curl or on my browser (that’s how I sucked it in, just copy-pasted a page from Chrome, formatted via IntelliJ, and used jq to do search etc). Anyway, this Kafkaesque design style is coming to an end soon.

January 25, 2023

Spent most of the time working on that Import Job; it turned out that this one, the true one, does not give a shit about Kafka, streams, etc - it just reads the whole dataset as a Json array, and is supposed to behave, but there are no tests. So working on adding a test for something like real data. Mongo behaves in a weird way; and there are all those cps-style futures, which make no sense, because we

have a Sequence, and nobody parallelizes anything (but we could, but there are just 2714 records in total, so why bother).

Anyway, told them I'm working on it and, hopefully, will solve the problem tomorrow. The problem is mostly about record ids being handled by Mongo... and the code structure where futures are being passed around for no reason whatsoever.

Also interviewed one bullshitter, who forgot immediately what I was asking him to do, although he, in the beginning, did rephrase it his way, and I showed him examples why that was not the right idea. Anyway, he happily copy-pasted permutations implementation from somewhere, I don't even remember, was it Java? Anyway. The guy who sits silently looking into his screen that he does not share, how, he thinks, we are going to work together? But when you ask "general" questions, he gives "general" answers. He even started looking up parsers-combinators in the first edition of Scala book, after I had to spell it in chat. I was not going to ask him about parsers-combinators, I was only looking for the boundaries of the topics he can comfortably talk about.

Anyway. Worked until 8pm.

January 26, 2023

Slapped together a PR for this import job, including a big test case and some checks for missing data that were supposedly imported. Sent for review several minutes before standup.

At 1 gave a talk regarding the new API, pull instead of push. Our service is more or less a data tier, like Kafka, but specialized, so it should not be pulling. Anyway, had a good discussion, including replacing "deletion" with "disabling" (not that it's important, but why not). Everybody was happy.

Later in the evening got a review from Karen for my PR.

The standard question was: I replaced trait SomeFunctionality {one method declared here} class SomeFunctionalityImpl extends SomeFunctionality with just class SomeFuntionality meaning that "Impl" is meaningless. It's not an abstraction if it has just one implementation. Karen did not mind much, just asked. And I really hate the idea of having "Impl" accompanying traits. She also had a couple of good remarks, so I merged this one and slapped together another pr, to fix these two teeny-weeny things.

Also, Renato wrote 4 or 5 Jira cases, for me to fill in, but it was all there, I just added points. Since in our company 2 points is 2 days, and nothing, with our process, can be finished in 2 days (sucks!), I put 3 points everywhere. I have not met a PM yet that would understand logarithmic scale (which is what points are in), so why argue. Just inflate.

That's it for today. Low sugar, tired; it's 7:30PM now.

January 27, 2023

The day was pretty intensive. Or intense. I had two PRs, but Karen is in San Diego, so had to wait. So I started investigating the new feature to work on.

One jira case, I estimated it as “4 points”; Renato told me that there’s no such thing as 4 points, either 3 or 5. Some PMs have problems with natural numbers, I guess. Fuck if I care. Let it be 5.

Then I had an interview with Chris W. The guy graduated from MIT, worked with hardware half of his life, then kaboom, discovered Scala, and started doing Scala. He did not exactly solve the problem, but was pretty close, and his idea was ok, would work when implemented. Seems like other people liked him too.

Now the problem at work that I’m trying to solve. The tier where I’m at is basically a data tier. A server that serves data, with some specifics. The data, in my specific case, are coming from Texas. Some Json. What the code does: get that Json, converts to a Scala class, then converts that scala class to Bson and stores in Mongo. When requested, that Bson from Mongo is read, converted to a Scala class, then converted to Json, and sent over on GET request.

Why Scala? The vague idea in the heads of some, hmm, murky geniuses was to keep data “type safe”. Of course between the time that code was written and now the sender added a bunch of fields, but we are clueless, we are strongly-typed. Nothing penetrates Mongo. And nothing is returned. So they asked me to add one specific field, and since the functionality is spread evenly in the codebase, I’ve been spending a bunch of time tracing what the fuck is going on with the values that I receive and store in Mongo. Well, also the class in the receiver and the class in the server were totally distinct pieces of code.

So, I figured, fuck it all, just take Json, transform to Bson, store it; on request gake Bson, transform to Json, send it back. Actually, we would not even need a Mongo, there are just 2714 records, who cares.

So I spent time trying to figure out how to convert Json to Bson and back, and how to test it properly. Last time I worked with Mongo it was year 2011.

The code is there now, and that’s when I just closed my work computer for this week. Till Monday.

January 30, 2023

Meetings most of the day. stanup, grooming, a short 1-1 with Karen, a 1-1 with Jacob, an interview with Vicrant.

Seems like when I finished the design doc, they want me, instead of just writing the code implementing all this shit, to write a document for those 4 jira “spikes”. Ok, I don’t mind, will include the things I’m working on.

Vikrant is a nice Sikh, but well, his coding skills are kind of kid-level. But he wrote some Scala, no errors. Just five times longer than needed. Instead of regexes, searching indexes in strings, etc. So, I don't know. Will talk with the guys tomorrow.

For our HR Samantha reviewed a bunch of resumes. One was really impressive:

"Lead data migration of static analysis data from Jfrog/Artifactory to Databricks using Spark streaming, Delta Tables and Glue Crawlers. During this project mentored four teammates to build the platform."

Here I just wrote to Samantha - "no way, she has no clue about the words she's using".

The rest of the day I was struggling with implementing what chatGPT advised me; not so simple, because the project I was in was using Scala 2.13, but the library I want to use is better in Scala 2.11, so I was jumping over a version, and had to figure out which version of each library to use now. The code is not a big deal, but still. And had to "migrate in Jfrog/Artifactory", meanwhile.

That's it for today; it's 7pm now, and since I started early, it's time to close the work computer.

January 31, 2023

So, meeting after meeting.

The Sikh guy, we discussed him, opinions vary. I eventually came to a conclusion that he's okay, and he is able to learn, and he will. So I'm rather positive. Some people were pretty positive; Chris was rather negative. Oh, and talking of another Chris, whom we interviewed. He chose another company. Why? Money.

Also, deployed to prod all the parts that I had to. This morning I figured that the job was deployed, but the service was not; so I rushed to deploy the service too. Works now. But... but we need good integration tests! So I wrote to our slack channel my complaint about integration tests missing. Karen and Jacob totally agreed. So I'll slap together some prototype, and we will see.

Oh, yes, and pinged the guys in Texas, telling them that the stuff deployed to prod is supposed to work (what a shame not to have tests to run with prod and tell me everything's fine); they confirmed, and the issue can be closed.

At 12 we had a 1-1 with Renato, regarding how do we deal with Jira, who creates which cases, what's the meaning of points, and how one can do anything if one's time is 60%-70% occupied with meetings. His solution? Add points. Well, okay; he's not the brightest star in the sky, but he's pretty smart and nice to deal with.

Had a planning meeting. Not much for me. Also, discussed with Mike the simplification of API. Instead of post/put/delete just write one post, and not bother with "editing a json with another json"; been there, even if I implement it (been there, starting with Google and our brilliant Brazilian colleagues),

users will be totally pissed off trying to understand how the fuck it works. Better make it as simple as possible.

The remaining time I was experimenting with reactivemongo and reshuffling the modules in the service project. Unfortunately, can't run the tests in intellij, because on every build sbt creates yet another copy of a couple of generated files, and complains. Fck. Anyway, tomorrow.

What else? Nothing else. Working on it.

February 2023

February 01, 2023

Kind of crazy, not of the best kind.

While most of my day was occupied by meetings, I was spending my backend time trying to get a decent library to convert JSON to BSON and back. Outcome: I failed. I think I could write a decent one myself, but is it worth it?

There are libraries in maven. But our jfrog does not list them, so I had to add the ref to my build.sbt. Eventually it worked. And it turned out that some smart guys somewhere had called the object that has the functionality json. There are tons of json things around, so the compilation (for some reason) behaves differently in intellij and in sbt. When I eventually pointed to the right one (not documented, I just looked up the source), it turned out they won't be able to convert *our* json.

Of course there's a cheap solution: JSON -> String -> BSON (and back). That's what I will do when I get back to it. Anyway, had a nice day.

Meanwhile, Renato set up a big meeting dedicated to one "very important" feature, or, rather, big piece of functionality, that we need to deliver asap. And I got a strong feeling that the company has just one programmer. Guess his name.

So, when they wanted me to wrap up what I've been doing and switch to that "emergency programming", I started asking them - how much time do I have to finish the case that was supposed to take a week? Two hours? Eventually they came up with a solution that I remember from my first job. Create a program but don't write it. Just document it. Well, I'm not 23 now, so I just smile and say "ok". The rest of the day they were pinging me, asking when I will finish the thing that was supposed to be done some time next week. I promised to document it tomorrow.

Actually, to document it, I'll need about 20 minutes. But I was busy! Context switch, anybody heard? And I'm from St.Petersburg (a spellchecker suggest Parkersburg), if you know what I mean.

Anyway, since my experiment failed, I'm free to write a doc tomorrow morning, before "standup". For the next "big feature" they also actually want just a design doc from me. This is really no big deal; I'm a writer, ain't I? Since nobody else documents anything, I'm feeling like I'm an internal chatGPT in flesh. Well, you see. I've been writing blogs for ages. Ok, for decades. Like Maria Wassiltschikowa.

February 02, 2023

I was kind of pissed off in the morning, they want me to drop the shit I was doing and do some other shit. I don't like it, I want to finalize and then get distracted. The only thing was that I strongly suspected I could make it exactly the same piece of code (abstraction!).

But since that big team of managers wanted me to slap together something fast (MVP?), they wanted me first to finish that case, which was not even a case, but a spike. A spike means I should write a doc. But I'm a writer. Writing that doc will take about 20 minutes. So the first part of the morning I was walking around the house, answering my mails, selling META and other stuff, looking at the papers that I'll need to be done with, and eventually I sat down and wrote that doc. Sent it to Mike, as they are our customers, for review. Renato wanted me to discuss things with the team first, and then with the customers. This is not my approach. Customers first. Mike was happy, he said he'll send it to his team. Fine!

The solution was: FUCK REST. Have just POST and GET. POST will send the whole record. No separate call for "unpublish", no separate call for "update". Have you ever implemented an update of a Json from a Json, or a PB from a PB, or anything like that. It's horrible. I implemented it three times already (Google, Salesforce, HealthExpense), and eventually came to a conclusion that this is the way to make customers' lives a hell. No. Just replace one record with another.

So that's what I sent them, and I hope they'll be happy.

At the "standup" meeting I just reported that I'm working on it. And what I saw in their "kanband"? No other case on me. The next one is for the next sprint or something. So I'm done with one (it's WIP, and that other team will be skiing, as Mike informed me). Cool. I have time for having fun/sex/fight with JSON/BSON transformations. Exactly what chatGPT told me last week, that it's easy. Fuck it's easy. There's a bunch of immature hackers writing some insane APIs for that. They are Java people forced to use Scala.

So I decided to kick it all out, and either write the stuff myself (since none of what I found works with our sample(s) - I have the full dataset that I pulled from our customers' service url, and no library is able to parse these jsons. Fine.

But then, miracle, it seems like there's a simpler solution: hide Json inside Bson, and that would be it; the authors of Mongo even have a code example for that. This is all we need.

Except that I have to slap together a test for that (almost there). And I have Friday for this! Glory to our company of lazy bums!

Also, talked with Karen about yesterday's interviewee; we have almost opposite impressions. For me, he just started the solution with writing a regex; Karen was sure that he's incompetent with regexes. What. Anyway, opinions vary.

Also, I slapped together a github project with interview questions and code samples. And btw, Davis tried my other question, and chatGPT solved it easily. With all the modifications. So... have to work on it.

What else? I have to wrap it up for today.

February 03, 2023

Not much to say. Finished slapping together a piece of code, in a test, that isomorphically stores jsons in Mongo (pretty inefficient, but so far it's ok).

Just spent a day working on it. That would be just one day, if I did not listen to that idiot chatGPT. Nunca jamas.

Posted my PR, and I'm done with this spike.

February 06, 2023

People did not like my pr with a sample file of 300k lines of json. Ok, do I care, I just made it 24k lines. They also suggested that, based on this test, I slap together an integration test and pull it from the service. Good idea. Next time.

While working on that test, found that our "mongo extension" library behaves a bit badly on missing resources (Scala 2.13 has it fixed, but this one uses 2.11). So I took that library project, fixed that small piece of code, sent a pr. But also, I'd say, I discovered this library. It was not touched since the last layoff, 2 years ago. Ok, and that's where I should put the stuff I was going to add.

The library itself is a bit crazy. I would not call it an ancient style, because it was in vogue about 3-4 years ago, but to me it's just crazy and stupid, with macros, and what is it all for? For a type-safe parsing of Json and Bson. Come on, there are libraries for that. More, you don't even need to parse anything, because you don't need the data in those jsons and bsns. It's just a storage. Ok, will work on it.

Looked at the calendar - wow, where are all those scheduled interviews? Gone. Why?

At a standup Jacob explained why. Hiring freeze. Also, the market shows: -7%. Ok, ok. I did not buy this garbage anyway. Maybe next week. Nobody told me not to.

Meanwhile, lazily, started investigating the new shit case they gave me. Karen suggested we get together and discuss. Instead, I sent her six questions, and she gave me six perfect answers. I'll document it and will talk to the guy who *manually* transfers some date from their team's server to ours, via our ui. No, I'm not rushing to automate other people's job, I'm thinking about having a reasonable solution for keeping TMS256 keys somewhere safely and casually. Will document, and will talk to Karen regarding our options.

February 07, 2023

Ok, not feeling well today. Brain fog, etc.

So, with on PR, still not pushing - but it's just an experiment. Will move it to "mongo extension" later on.

Also, studying this stuff, tokens, HS256 encryption keys. (Why symmetric? Beats me. Where's hash?)

Anyway, reading, documenting. Found a piece of code (thank you, hound app) that already does some stuff with triton tokens. Written by Jorge. Jorge is still with us, so why not give it to him? Weird, but I don't mind. I'll document it and make it work for a large set of providers. No big deal.

But again, some little covid returned, so there.

February 08, 2023

Basically spent all the time trying to figure out where to store the secret keys (encrypted), and how these secret keys are used for generating triton tokens. Still experimenting.

Exchanged a couple of messages with the people that are anxious about these token being available any time soon.

Figure that this app uses postgres, so I'll be storing the stuff I do there too, big deal. Alternatives were: mongo or aws files or dynamo.

But still can't find a trace of those secret keys. Oh, whatever. I have a couple of days ahead.

February 09, 2023

So, spent most of the day studying the oauth service. It has integration tests, but they were not run since before Covid. Tweaked them a little bit, but have to work more.

Anyway, I know all the tricks now, saw all the ropes, and just finished writing a design doc. Not 100% sure, but it seems like that's how spikes should end, right? Design docs.

And thanks to Karen, she reminded me the trick with encrypted values in configs. com.typesafe.config does funny tricks.

That's it for today.

Interesting how most of the work these days is done by me walking around the neighborhood and pondering. Not crazy typing or desperately googling.

February 10, 2023

Today is a silent day, "standup is by slack".

Added the table schema for the keys in the doc I wrote. Sent it again to Karen and to Kerry, let them review. Kerry seems to be pretty satisfied. Will have to schedule a team presentation for the next week, I guess.

Since all my cases except one are closed, started investigating the last one. It's about updating elastic search index by the data we get from "editors". Ok then, looked around in the codebase. Cloned a huge huge project, with Java beans inside. Horrible. 2014 is not that far away, but the technology looks like it was 2004. Horrible. Thank God it turned out it only mentions elastic search, but does not do anything there. So asked the people around, where do we update it? They pointed me to a rather new project (2022), which looks ok. Wrong design though; but I'll fix it, no problem. Not much to fix, and it's independent from everything else. Will just write another design doc, big deal. But first have to study the code, to make sure I know exactly how it works.

Then Jacob wrote about SNYK, do we use it? Jorge writes - "will it really find anything bad"? I wrote - "hardly, our code is professional Scala". Anyway, installed the plugin. But shit, the tool is undocumented. Of course we can document it, but what are we paying for? Posted the question, and decided, ok, let's discuss it next week maybe. Jacob has a similar opinion.

In spite of the fact that we are in hiring freeze, Samantha asked me to share the list of technical questions. It's in github, so I sent her the link, asking if she's able to access it. I guess she's afraid she can lose her job. She's really nice, good to work with. Ok, New York, New York. Cutthroat assholes left in the Bay Area.

February 13, 2023

Fuck, it's already half past eight... So, studied our app for elastic search. Found that we will need one more dataset from "edit" tier, from Texas. Ok. It's just 73 json records, but somehow our "job" reads it twice an hour. Anyway, that will go. This additional data it pulls is related to the other project, the token service, which is in oauth, and uses postgres. So what, we need to keep some data in mongo and some in postgres?

By the end of the day, after talking to Karen, who also suggested to use Domino, I figured that wtf, we have a piece of code that works and is configured properly; why bother, ok?

Planned to write a doc regarding managing that elastic search, but then asked Mike, the guy in Texas, what are our options. He sent me back a list of data sets which they want us to be receiving. Cool. So we can reshuffle the list a little bit and proceed. I'll use mongo.

Before the 11AM standup, asked Renato regarding how should a spike be closed, some code deployed, or? He said, a wise guy, that a spike ends with a document that is discussed - so (I have a document) I suggested to discuss it tomorrow (and Renato sets up a meeting, neat!)

At 1:30 had a 1-1 with Jacob. Showed him our woods (I was sitting outside), and he said (he's in Wisconsin) that they are also interested in moving over to South Carolina. Welcome to the South! (Wisconsin is not bad too, but too cold.) Then we talked about problems (turned out there are none, zero, zilch), then about SNYK, and he showed me how to do it in intelliJ (via help->tools); then I asked him how

am I supposed to request a vacation (just a couple of days in April). He said: nobody does; people just label their absence as “OOO” (out of office), and that’s it. Of course if you are absent one week every months... well, that’s not my habit. I kind of just work.

So that’s it. Then I ran SNYK, and it complained about a missing sbt plugin, and Qing wrote that I have to add the plugin is sbt.plugins (ok, thanks). Also asked if anybody is interested in attending the conference in Lisbon (nobody), asked Jorge specifically, and he says that from his country he would need a visa and covid proofs etc), so he’ll do it online. He’ll probably give a talk on ZIO HTTP (one new hot dish these days).

And also some small discussions, and I’m ready to document what I’ll do with elastic search, and Renato scheduled a presentation tomorrow. Cool! I love it so far. I’ve been here for almost half a year and still nobody seems to hate me. Weird. It’s like HealthExpense, the first couple of years.

February 15, 2023

Well, I like our work organization. Did some stuff in the morning... actually, investigating elastic search indexing. Then we had a standup; Renato asked again if I think my spikes are all done, I said yes, and he asked me to create cases for development defined by those spikes.

It means we write plans and then follow them. I love it.

So, I got distracted, but 30 minutes before the planning meeting I slapped together 5 cases, and Renato wrote one for me, too.

During the meeting we just covered all those cases, and that’s it. Now I can go ahead and follow my plan. Last time I worked like this it was in the USSR, I wrote plans for the whole team, and then we followed the plan.

But I did not have much time for anything; after grabbing my lunch, joined another meeting. That was something like a conference between Texas and us, new-yorkers... I am actually a yorker. Whatever.

Those guys made a presentation regarding their architecture, development and deployment (I liked the “deploy” button); then our guys, Jacob and Chris, talked about ours. Neither set of slides did make too much sense, but at least I’m familiar with ours - and asked and got a pdf with their slides.

Later I tried to work on case 1, which is a db schema update in postgres - but the appropriate test was failing because we had a VPN problem, and when I decided to fix the config library code (to make sure it reports this kind of failures properly), I could not, because, guess what, we had a VPN problem. Ok, I just closed my work computer. Tomorrow, tomorrow.

February 16, 2023

Spent a day testing my db update in integration tests. The worst part was adding a trigger to postgres; eventually decided to ignore it.

But added a test to validate the db update. Ok. And when I posted it to Jenkins, that integration tests started failing. Last time it ran it was April 2022. So, notified Swathi, wtf is with aws access from Jenkins, and now trying to block that test for now. Jenkins shows pretty weird messages.

That's how I'm spending my day.

February 17, 2023

So, my db update pr was approved, now I tried to update the db, in staging. Turned out people do it manually, via psql (ok, the tests with that db did pass, but updating manually... ok, I'll slap together a jira case for that automation).

So, asked Swathi how does she do it, and whether it's fucking documented anywhere. Of course not. So I'll torture her Tuesday, to slap together the document.

Besides, at last I fixed my “gitpush” script, to open the PR page on every commit/push. Convenient.

What else? Nothing so far. WIP.

There was a meeting at 2pm, all-hands, but I was just napping that time. Oh, whatever. Nobody cares.

The company had a small layoff. Our team is not touched, because, well, we already lost a bunch of people. Also, the bonus looks ridiculous this year, but again, who cares, right? I lost more on the market today than that bonus, but I'll get my money back next week (I hope, I hope, I hope, I hope, I hope.)

February 20, 2023

Had problems with VPN today. Just no connection. Asked my colleagues, how do they work? Strangely, they are all okay. Hmm. So I was just sitting thinking how to finish two Jira cases by next Wednesday. Of course, MVP, but not exactly so. I need, if not cut the corners, then just cut extra corners. E.g. Write, in the code, a db update, if it's not updated. No big deal, right? And stuff the data into a config file. No manual job, like logging in to postGres on production. Who am I, a devops? I'm a programmer. I follow all five principles of Amazon's good architecture: Automate Everything!

February 21, 2023

Ok, mostly meetings, and a long talk with Karen, since I'm shadowing the pager duty 1st tier, starting tomorrow. Also, the pr with the test got some decent comments, I started fixing, and kaboom. Ac-

tually, I also updated one method signature, and it seems like it's being mocked in play via guice, a stupid "di" product from Bob and Kevin. Bob and Kevin are very very nice guys, but they spent too much time in Stanford writing code, not learning how to write decent code. So they slapped together that piece of shit, and Play framework people used it. And our engineers took it, and they also took Mockito, and as a result any innocent refactoring can cause weird "class not found", "dependency not found" etc errors. I don't think Guice and Mockito should be ever used anywhere. Neither makes any sense... mathematically speaking. Oh, and neither Bob nor Kevin have been related to Guice for the last probably 10 years or more.

Anyway, the right approach is not to touch that old shit. So, redoing my changes.

Basically, my idea is to stop manually updating the db schema, and have files with updates listed in a directory, and pointing at them at start, making sure that they are called only on need-to-know basis, and never manually. An app starts and first makes sure that it has the right tables available, or else it creates them.

February 22, 2023

I'm so fucking tired today.

First, tried to get through my PR. Failures are random. E.g. git finds a bad formatting in the code, it does not report it, just updated the file; then, in the end of the build, jenkins says that oops, the file was changed during the build; failure.

WTF. Ok, I've eventually found how to take care of it, by creating build.sh for each project (commands vary, defined in some jenkins config files), and so let it check before I commit. Fck.

Then, I'm on pager duty today, ok, shadowing. Nothing happens, but there was some spike in the graph, in graphana, I was clueless, should I be sitting, like a guard, in front of 10 tvs watching the graphs, or? Whatever. Let's see how it will be happening in March.

Karen was doing Karen stuff. My test sql file was adding production tokens to a test db; she noticed that we should separate the tokens, one for test, one for production. Wow, at least somebody fucking cares what's the statistics on a test instance. But whatever.

Also, asked if we have a project with a common library. Or else tons of code is copy-pasted from project to project. No we don't. So I asked Swathi how we create a new project ("a new standard"). She told me, and I'll do it.

Also, there's a doc, written by Jacob, about "engineering initiatives for 2023" - my "common lib" initiative is there, but I'll have to read the whole doc. Just too fucking tired today.

Oh, whatever. I have a week to finish this API. Tons of time, in normal circumstances, right? But in these circumstances, everything's crazy slow.

But at least, while walking today, at least I came to a conclusion how I'll provide the API backend. It remains to implement it. If I'm not distracted, I'll do it fast. But it's all about meetings and context switches.

February 23, 2023

Today we had just two meetings, and one I had almost missed. There was some upper-level manager telling us that everything will be okay (quoting "Mickey Blue Eyes") and that we should not worry. Meanwhile since Karen approved my long pr, I had merged it and started doing the next case, API for put/get. Sounds ridiculous, but this codebase is a bit weird. Some data go to Domino, some to Mongo, some to Postgres. I love Postgres.

The funny part was, my integration tests did not want to run. Just did not start. After a bunch of stupid attempts (like, making a copy of the test that did start), eventually they all started running (except that my new one was failing because my sloppy url transformer had a bug. Ok, fixed that, the tests run, will see how they are tomorrow. My plan is to send the first PR tomorrow, so that I can move on. Yes, it's like a stupid hobby, programming, and I love it. Especially refactoring, converting weird code written by weird people into something OBVIOUS.

Also, Jacob notified us that we can postpone filling our капиталистические обязательства (OCR?) till April 1st. Good timing.

Ok, good night for today.

February 24, 2023

Lucky day, no meetings. Oh. Maybe I missed one. Whatever, who cares.

Spent a day playing with slick, postgres, etc. By the end of the day cleaned up the tables and wrote the code that at least fills the db and reads the records (which are in the db split into two tables).

Who cares. Just learned how to deal with pg via slick. Not bad, not bad. No json involved. Pretty much type-safe.

Sent a PR by the end of the day.

Oh, I'm also on this pager duty - but that was easy.

Ok, around 10pm figure out that I let in little Bobby Tables in my code. interpolation of a sql statement is not the brightest idea, right? So there. Added some checks; later will figure out what can I do to properly parameterize the statements. But now I just check what I'm given, and throw little Bobby Tables under the bus.

February 25, 2023

I know it's Saturday. Just for fun, updated some code, implemented a getter and a setter, wrote some tests including Bobby Tables test. Have to figure out how to prepare parameterized queries. Oops, just found. Will fix.

That's it; it's all fun.

February 27, 2023

Adventure day. While I was waiting for a review, I was throwing in more updates, including creating all the tables in the db, if missing, and populating them, if data are missing. Turned out, I want some data to come from a config, where it is encrypted, and keep it encrypted. Does not seem like doable with the current config library from typesafe. Whatever. Will find something.

Had a 1-1 with Jacob, pretty short, everything's fine? everything's fine.

While I was working on my shit, tests stopped running, because there was not enough disk space in the postgres machine. Omfg. Wrote to our devops people or something. Karen answered that we can ssh in there and check.

Karen approved my stuff, except for a typo in an exception message, which she suggested to fix; I did. And kaboom, the tests in Jenkins all failed too. I mean, the integration tests.

I was in panic. Ok, I looked through the code. Problems: - these integration tests were not running for over a year; - every test case creates a new db; - dropping that db was in a wrong branch of the code, so they all stayed.

Fck. Ok, so... what I did. Wait, it was not easy.

I could not even ping the pg db. Why? No vpn. Why? FIIK, my AT&T was blocking it or something. So I collected via my phone.

Then I had to extract the db password, which was hiding encrypted in a config file. For that, I knew how to do it, run the code in debugger, and voilà. But it did not run, IntelliJ was just hanging, taking 370% of my cores. Had to kill it, restart it, run the debugger, found the password, remembered it forever, and voilà.

So, what I did.

- installed psql (some ancient knowledge)
- found the trick how to do bulk DROP DATABASE
- dropped all 308 of those test databases.

Profit. But github did not recognize a rerun in jenkins, so I had to add another empty line in a .sql file, commit it, and profit.

Ok, I have just two days left to finish this functionality.

Oh, and by the way, I'm also on pager duty, answering and confirming all those "limits exceeded" failures.

What I will do. Check that the db deletion works. Maybe merge some test cases into bigger ones (not nice). And tomorrow I'll have to finish this db preparation part (all tests are already there), because when it comes to prod, I really don't want to log in and start manipulating the prod db manually. No.

So, what do you think. I find it all funny and entertaining.

February 28, 2023

It was complicated. See, it's 21:22, and I've just closed my work computer. First, I'm on pager duty, and a bunch of alarms were coming. Like, Kafka unable to send data to some service. Eventually Chris asked me if I'm going to take care of it.

Ok, ok. With Chris's help, eventually I went to hadoop, killed a couple of tasks, then in Jenkins launched them, etc. It works now. Also, documented my actions.

See, it was supposed to be Karen, but Karen is sick, so there.

The rest of the day I was struggling with integration tests that were written god knows when, and never dropped the databases they were creating. It did take me some time to properly code it. To drop a database, it's better to do it via command line. So I did. Oh, whatever. The biggest problem was, some shit was keeping connection open. More than one. So first I had to make sure all those connections are closed. Oh, whatever.

What else? Wrote to Renato that I won't be able to finish it by the end of the day tomorrow, and I'll need a couple more days.

That's it. I would not say it's not fun. Curious. The things are getting simpler and simpler. Oh, whatever.

Bye...

March 2023

March 01, 2023

Oh, I was working all day and just finished it (pushing my shit to github). What was going on: Postgres, handled via Slick, was hanging on certain operations. Giving it 15 seconds or something did not help. More, psql was hanging on connection. Meaning that connections were interfering. Really not my area of expertise, yet. Safa was supposed to be helpful, but well.

Anyway, now I know what happened.

Slick, when you call it, returns you a future. If you don't want to wait (why should you), you "chain it" with the next future. But the fucking truth is that it's a Scala Future. It launches right away. Which is not what I want to have, I want it start after the first one is finished, there's a dependency.

So, one solution would be to concatenate all those sqls in one big chunk, and execute; another would be to remember what I knew about Scala Futures and chain them properly. The third solution would be to use ZIO.

But wait, I can't seriously expect something different from Slick, it just returns you a launched future. I'll have to chain them via "onComplete".

Vaguely I remember dealing with this shit eons ago, no I have to fix it. Why would a promise help, I don't see. Oh, whatever. Tomorrow.

An alternative would be just use a regular CPS, since that's how it works. Not creating a future before the previous one started. That would be the sanest solution.

March 02, 2023

It's a bit crazy, to work from 10 to 9:30. But it's ok. Subtract 1h for a walk, half an hour for lunch, and it will be just 10 productive hours.

So, I fixed all those integration tests - meaning, of course, I fixed the code. I decided not to get into troubles with CPS, since we are in an environment... etc. Could do it, but instead I just pinged Jorge, asking if he has suggestions to convert all that shit with Scala/Java futures into ZIO monad. Let's see.

Meanwhile I just separated the DB initialization into sequential steps with Await. Because, well, we can't do anything until the stuff is done anyway, and it's not multithreaded (is not it? I guess it is not). DB initialization is idempotent, so no big deal.

Then, when I tweaked this and that and created a class called SQL etc, and created a PR, I discovered that shit, jenkins does not have psql installed, so I have to do it properly. So I did. added some admin operations in SQL, including dropDB. All works. Kindergarten, how many times should one write this kind of shit, eh.

Anyway.

Meanwhile, Karen updated the doc, adding a pretty nice solution for feature branch deployment to staging. No need to do stupid merge to master poisoning it with unchecked stuff. Ok, maybe integration tests could help. And of course, Karen is Chinese, so she does not automate, she prescribes. Ok, documenting is also not a Chinese feature, but I kind of started this culture, so she follows now; cool!

Well, see, this programming profession, it's a drug. I had to work on it intensely, the whole week, why? Because I could not stop until (almost) all problems are solved and everything's illuminated.

Now I can relax.

I mean, there's the next step, slap together some rest api, but it's not a big deal. Routine.

March 03, 2023

So, done with my PR, except that ordering of the operations requires some cleaning, so I asked Jorge, and Jorge sent me a pretty neat ZIO-based solution. Cool.

The rest of the day I was working on formulating the new REST API in the framework we have (that's swagger doing all that generation); not much success yet, but at least formulated.

Anyway, the weather is good, so I was sitting outside.

March 06, 2023



So, after a bunch of questions from Karen, which I answered by providing more documentation in the code (which is good, since it's not trivial, how come "populateDB()" is called exactly once, without "variables". Scala tricks), she approved it and I committed it.

Meanwhile I was working on the next step, writing some swagger config to generate proper API endpoints, and it was not trivial, since if your endpoint accepts a json, you don't write a function with a param, but use a cps style. Not a problem, but not trivial, had to read through tons of code looking for examples.

Anyway, the stuff is still not compiling properly, but it's ok. Working on it.

Meanwhile one service stopped, temporarily, responding, providing user's location, and since it's about triton service, the suspicion was on me. I just told them I did not deploy anything for the last three weeks, so not me. I believe our logging is not good enough. They had to run around and guess what's up, saving precious disk space (which is totally worthless).

So, It's 7:40 now, I'll go ahead, experiment with swagger a little bit, and close for today. Nichts neues. The only stupid thing was, I managed to slap together new updates in the main branch of my fork. Tried to roll back the branch, but it's really not easy. More, the manual reversals did not go well, due to an empty line in some files that were automatically injected. Eventually it was fixed when I merged my other pr into the main master, and then I was able to sync my master. Stupid...

March 07, 2023

Ok, I'm kind of crazy. When the things work, it entices me to work more and more and more.

I even agree now to use Mockito (which I've been hating for a long time). Why? Because for simple cases it's pretty much usable.

So, the functionality is almost there, and it's ok. Only one test case left to be written, and that's it.

All those swaggers, mockitos, etc, all solved.

But let me sit an watch a movie today.

There will be a bunch of meetings tomorrow, during which I'll create two or three test cases, make sure they run, and I'm done. Like a week ahead of schedule. Ok...

March 08, 2023

So, the only remaining test case was failing, and why? It was a setter. I spend all day trying to figure out a compatible transforms, json to instance and back. Turned out that joda datetime which I converted to international format, when serializing, was not always serialized this way, and was represented as millis. Why?

Only when I drove to the airport, later in the day, that I figured I should not even have it in the record I receive. It's our internal issue, and that's it. Receive an update without timestamps. Shit, I had to figure it out earlier. Anyway, will have to spend half a day tomorrow to redesign this thing.

Anyway.

Also, the contractors broke my internet again (cut the cable); they though they fixed it, but they did not. Will have to call AT&T again. They did "connect" the cable, but it still does not work. Not a problem, I have my phone.

March 09, 2023

Working like crazy, because, well, I knew how to refactor the shit, and so it was just full speed.

Some time after lunch that PR was ready, sent it for review. Simple, easy to understand.

Also, wrote a method, for testing results from futures, that the usage looks like

```
from(callSomethingReturningFuture(params)) expect theExpectedResult
```

And rewrote all eligible test checks like this.

Which definitely reminded me Shakespeare, <i style="color: rgb(62, 62, 62); font-family: 'Trebuchet MS', Arial, sans-serif; font-size: 20px;">“From fairest creatures we desire increase”.

Sent the PR, then reviewed Qiang’s PR, which was pretty neat, and ZIO-ish, but I don’t appreciate injecting 10-line cases inside a for-loop, so that’s what I suggested him to do: refactor. Extract method.

Then what, then the next trick I’m doing: decrypt encrypted keys that I take from a db. Decrypt them the same way that typesafe’s config does, and I’d rather use exactly the same code. And cache the results, of course. That’s WIP.

There was a meeting regarding some attempts to hack our site, coming from Portugal; I was attending the meeting, but I kept quiet, because, well, really not my area of expertise.

That’s it for today. I may work too much, but that’s my hobby, so there.

March 10, 2023

Ok, it’s almost 8pm, why am I working?

By the end of the day Karen approved my pr, but with more good questions (so I’ll have to fix these little things; I really love working with her).

Oh, and I reviewed Qiang’s pr, wrote one suggestion, he implemented it; cool.

Spent this day working on decryption solution. Ok, now I do decrypt the data coming from the db, via the credentials coming from config. Had to cannibalize some code from William Ho, written a couple of years ago (before the layoff).

I had two parallel PRs, and when Karen approved one, and the other one was ready, I spent this evening (ok, 30 minutes) merging them. Is not it fun? Now integration tests are running (swagger reporting ancient warnings for William and Joseph (who left two years ago).

Karen found a funny defect in the code I sent her for a review. Thank you Karen! The defect is fixed.

Now what. Now I add a test case. Oops. The compiler fails.

Oops. It hangs in a loop. Well, ok. Just commented out most of that test case.

Oops. Other tests fail too. What?! Omg. Had to add a comma. What was it, really?

The time will probably come when I stop it with Scala. These people, they are crazy or something. Now we have another scandal. Some idiot, “the manager of Scala community”, with an education

in “politology”, from Belgrade University, decided not to allow John de Goes as a speaker at a Scala conference. (The only beef with John is that he’s bold and muscular: obviously a racist.) Another Jon (without an ‘h’) was kicked out of everywhere, because the girlfriend of a left extremist from Berlin had publicly complained that a while ago Jon had sex with her but refused to marry.

My God, what a bullshit. I have to return to Forth, the geezers there are all peaceful and smiling (and I almost know them personally; at least by their source code, in the old days I read that code as a poetry). That’s stupid too, of course.

March 12, 2023

Ok, it’s almost 8pm, why am I working?

By the end of the day Karen approved my pr, but with more good questions (so I’ll have to fix these little things; I really love working with her).

Oh, and I reviewed Qiang’s pr, wrote one suggestion, he implemented it; cool.

Spent this day working on decryption solution. Ok, now I do decrypt the data coming from the db, via the credentials coming from config. Had to cannibalize some code from William Ho, written a couple of years ago (before the layoff).

I had two parallel PRs, and when Karen approved one, and the other one was ready, I spent this evening (ok, 30 minutes) merging them. Is not it fun? Now integration tests are running (swagger reporting ancient warnings for William and Joseph (who left two years ago).

Anyway, let me finish, and I’ll go do my French (pardon).

March 13, 2023

Not much to say; fixed all the merge troubles, decryption works, pushed one pr with all that, and another pr for deployment of the previous pr to staging.

Some of the work was done in the airport, CLT, where I am now.

March 14, 2023

Ok, I naively thought that my deployment will be ok, but late tonight (it was 21:30) I discovered that ha, constraints names in postgres must be unique, so my functionality is not working there. Fck.

So, updated it, and now running tests on that PR.

Also worked on managing the signing keys in the oauth code.

Also Karen asked if I'll be ready to deploy that shit to prod. Now obviously not - and it does not make any sense actually. Prod or staging, it's just an experiment.

Yes, on the planes I was only opening my work computer.

March 14, 2023

Exchanged opinions with Kerry and Karen and Davis, updated the doc regarding data handling for tokens. Also, wrote a doc for search engines indexing solutions, seems like Karen is ok with me moving imported data to MongoDB.

Presented my tokens doc to the team; got a couple of nice ideas and questions, and that's it. Closed this spike. I guess the search data import doc can also be closed.

What a life, no need to write code, just write docs and talk. Weird. It does not require much concentration at all, I can slap together docs pretty fast. Ok, yes, I did have to read docs and code, a lot.

Ok, we will see.

March 15, 2023

Ok, it would have been fine, just sitting in a train, working. Except that, fuck. First, I'm on pager duty. And there was an even, at about 3am EST, and I tried to investigate, because, well, I was on a French train, and the internet was disconnecting once every 5-10 minutes, and I need go get a confirmation from my phone, which wants to recognize my face - after about 36 hours of casual travel.

Also, today is the last day of a “sprint”, and it ends at 4pm eastern, which is 19:00 French time (today).

And there were meetings.

But the fun part was that I did try what I deployed on staging, and found that the code was getting an old config, and wtf? I posted a question on slack, but well, everybody's asleep. So there.

The good side was that when I took a bus from Agen, I was just looking into the window, seeing all those green field, not parched by the sun (yet).

When I got the opportunity to get back to the internets (meaning, my computers are ok, but my phone was discharged, and was unable to see my face (and report that me is me), so I had to go to a gif store or something, and wait in line, while one madame bought a flower pot without a bar code, and the next monsieur bought a frying pan, but he did not have the store card, so he could not get the discount... and it was 12 minutes before the “standup” meeting. So I ran back.

Ok, ok. Then it turned out that while I was implementing, nobody on the other side figured that, well, to get a key from my server one needs a key id... ok, at least the provider id, can do - but they just

introduced a flag: “need a token”. As if my code is a GPT and could make any wild guess regarding what t.f. that token is, and who they are, those clients.

Ok, ok. At least Davis and Karen told me why my config is not available in my code (although they are in the same repo).

Configs are tricky in this application, they are posted to a config server... In short, it's all smart, and correct, and undocumented. I had sworn to document everything I see. The 1:1 map of China.

In any case, the deadline for the “sprint” is 4pm eastern, which is 21:00 french time, but I managed to deploy it to prod and test it before 16:00 french time. So, I’m good.

Did I tell you I’m on pager duty? The good part is, the application was written by geniuses (genii?) and only cracks (not crashes) Monday morning, not even every morning, and the failure is pathetic, and one can just manually restart that service (this generation is not familiar with scripting).

Ok, anyway, and we are all good people now. Just too many meetings.

At 1pm we had a post-mortem meeting. Although nobody died. And Jorge wrote that we should immediately adapt the newest version of ZIO. There was silence, but I started giving my comments, how important it is to use the latest version of ZIO, with Layers, and how this framework is totally superior to any other stuff/shit (like Cats) - yes, it is. In short, Jorge (btw, a contractor) got his idea pushed.

I also have an idea - my server does not check the encryption, the usability of the signing keys, the non-emptiness of the keys of the signing keys. But I’ll do it anyway, big deal. We are on the right path.

At 2pm eastern we had another meeting, planning.

So, by the time I deployed to prod, I a) was totally exhausted, b) could not find how t.f. I can get errors in the logs from Kibana (I just forgot how to use kibana... I really forget too much, this morning, on the bus, I forgot both Dunning and Krueger’s names, wtf, right? How long do I have before Alzheimer hits? Fight! Struggle! Never surrender!

But it’s 20:37 now, and I’m totally ok, everything’s fine, and tomorrow will be a great day (I’ll need to implement the token generation based on keys in db, cached).

March 16, 2023

So, no alerts the whole day, except some minor complaints about some performance, low priority. But Karen is doing Karen, and she’s trying to convince me that since I’m on pd, I’m supposed to not do my usual work, but go ahead and fix minor problems here and there. Ok, sure. Working on it.

Meanwhile, Jacob notified us that copilot is available. Cool! Installed it immediately. It works! Not always guesses my thoughts, but pretty close. So typing fast is not a required skill these days. Most of the work can be done while walking in the park. Then, when ideas and solutions settle, I can come

over and type some letters, and the copilot does the rest. Thinking is (still) (temporarily) a human occupation. Cool!

Well, but before lunch I was partially distracted by some shit here, so there.

But around 4pm I went for a walk in the park (photos later), and figure out the whole solution (except the part that requires refactoring some crazy CPS code. I'll make it simpler.

When I got back, I did slap together some code, then I talked with Karen, when she told me I should not be working, but should be fixing (ok, sure, I can go ahead and fix my code and hers: it was her who introduced this crazy CPS style, instead of splitting token generation into three easy steps, overriding one step for specifics. Oh.

Anyway, I've a plan to go do some more work right after the shower. The problem is, I have not adapted yet to this new schedule. Trying to figure out.

March 17, 2023

Ok, I spent half of my sleep trying to figure out how to send more dependencies into a callback. Eventually figured out (S-combinator!). So in the morning worked on that stuff, prepared the first commit for the PR. Just to have a picture clear (or clean?)

Meanwhile, in the free time, I was reading docs regarding “pager duty”, and you know, I don’t like it. I’m supposed to work on cases, issues, old reports, anyway, but not on the thing I’m working on. The problem is, if it rotates in my head during my sleep, how can I just forget it and start doing some minor things? Well, I understand the psychological aspect, but I never actually worked like this. What I can do is just go take a walk.

So, as a result, some time after lunch and some work I did I pushed it to github, and went for a walk, and walked until it got dark.

BTW, I use copilot now. And it really makes googling obsolete.

March 20, 2023

Ok, the post I wrote just this moment, disappeared.

Anyway, everything’s fine, and I slapped together a class called Crypto that does encryption/decryption via KMS and AWS; using it in my oauth module.

What else? Everything’s fine. Pity the dybr was lost.

March 21, 2023

Worked in different places. Since Karen was asking this and that about my PR, told here that I'm actually working on the next integration test, to check token generation. But I have not finished it yet.

Also, since I'm on pager duty, I was getting some shitty notifications (kafka overflow); and started asking our devops, wtf. They are clueless.

Then I had a 1-1 with Karen, and she, being Karen, patiently showed me tons of undocumented steps to reach the logs of that server, and to kill it and restart.

I signed up on fixing this failure (I mean, not restarting, but actually investigating and fixing).

Meanwhile Chris suddenly wrote a pretty neat doc explaining all these secret steps that Karen told me.

Meanwhile, my PR was still not approved.

That's it.

I have to dosage my work time.

March 22, 2023

Why the fuck do I work now that is already 21:30 CET? Ok, I have another PR, because, well, signing worked, but I did not check validation, and for validation to work, I had to refactor a lot. Because, well, key ids are being passed around.

Besides, had to add more "standard providers".

And renamed some classes, just because their names don't fit well in JBuilder window.

Done with the pager duty, and can do some more useful stuff. Ok, at 7:00 was woken by all those pager alerts, and "confirmed" them.

Nothing else. Missed one meeting. Will have to install outlook on my phone. That's it.

March 23, 2023

So, what's up. Fixed some comments to the code comments in the PR. But then Karen found another defect, and not that I trust her findings, but it will be more reliable if I fix it. It's about the idea that encryption engine for prod is different from the engine for staging. How do I know if this is the fact? I don't. But better no rely on the keys which may be not decryptable on prod.

So I spent some time making this shit independent.

Also, started a project for "shared common library". Sbt is crazy, IntelliJ is crazy. Just creating a project from scratch seems to be not working with default settings. Oh, whatever. We, programmers, copy and paste.

And right now there's a meeting, they restart mongo and look into how come it overflows with some maybe non-trivial queries. I'd love to participate, but everybody has a large screen, so viewing it inside my small mac screen... I hardly discern the letters. I just see spikes in the curves, which titles I can't read.

BTW, it's really not my area of expertise. I'd rather write code.

Oh, whatever.

And btw, Jacob suggested me to install outlook on my phone - I spent months looking for ways to link google calendar with outlook calendar, could not make it work. Installed outlook, now I'm totally fine.

March 24, 2023

This enthusiasm does not look like sane. Ok, I was cleaning up the next stage, to make sure that we don't store in the db the keys not properly encrypted. Which means, I have to decrypt one before storing, and to test, I have to encrypt something.

While doing this, I discovered that a) the encryptor wants BASE64 encoded data.

Now imagine. We take some bytes. Encoded them to BASE64. Then encrypt. Then encode the result as BASE64. The documentation a) insists on BASE64 url, and b) says nothing about the input having to be in BASE64. Well, whatever. Also found that my "BASE64 checker" had a bug. For some reason the data encoded don't always end with at least one = - why is it, really? Well, whatever. So far it works.

Also, prepared a pr for deployment to staging; it was approved, but I did not have time (or desire) to go ahead and deploy. Monday.

But I worked till after 22:00 CET. Maybe it's too much. Did not do much else.

But at least the code is getting neater. Except that, just for readability, I wrote this function: def uncertainFutureT: Future[T] = Future.fromTry(Try(fun))

It's a natural transformation, and with this we basically hardly need any "EitherT" or any other shit from Travis and other half-educated dummies.

Ok, IMHO.

March 29, 2023

In the morning deployed to prod, it seems to be working; notified Renato that he can tell the customers ("client-side people") that they can do. And started slapping together scripts, to encrypt/decrypt, to reencrypt the same keys from staging to prod, etc. Spent the whole day scripting. Ok, it worked by the end of the day (by 22:40), but I found that what was deployed to prod does not accept my keys, for

two reasons. 1: wrong regex (my fault, I did not cover it with tests), 2: seems like prod won't take my updates from France. But that's not a problem, I can ask someone in America. Anyway, just running tests on my new pr that fixes all that.

Meanwhile, we had three meetings! 1. Standup. That was easy. 2. Postmortem. That was nice, we are all good people and communicate (no kidding). 3. Planning. Oh shit, that was hell. I don't know how many times I explained that my new spike has nothing to do with encrypted signing keys, and that covers, more or less, all future demands from our "editor people"; and then Renato found some "back burner" jira case that was postponed because they wanted me to do the keys.

I was exhausted. And then, when I discovered that my updates, sent via "legal API", don't make it (ok, I could of course just write a couple of sql statements... four actually, or rather generate them), then it was a really nice outcome.

Just sent a pr for review. Oh, whatever.

March 30, 2023

In the morning I thought again that it was not the brightest idea to take over the jira case with "add more fields" to the api that should be agnostic - and instead, I can as well leave it to Qiang. Also I remembered that the necessary code was mostly already there, so it required a pretty small tweak to start passing around intact records. For the last two months the whole database was repopulated with the right data, so there. 70% of the code should be kicked out. Don't do "type checking" of the data, the jsons, that are just being tunneled through.

And I knew that in San Diego, where Karen is, it was still 2AM, so I have time to think about it.

And since it was a spike, not a dev case, I could just write a document - so I slapped together a document, sent it out to the people, specifically mentioning Karen and Renato. The document is a little bit like revolutionary manifesto, mentioning typeless lambdas and the obsession of FP people with strong types where typing is not their business at all.

Meanwhile, my other two PRs were not being reviewed (ok, morning over there), and also my power block stopped working, so I decided to go to a neighboring city and buy another power supply for macs, leaving a slack message to the team that oops, I'm out of luck. But then it started working again; weird!

So I chatted on slack, and discuss my scripts with Qiang, and that was it for today.

Everything seems to be fine today. Tomorrow morning will check all those things that I left for reviewing... ok, ok, one other thing. So far I can't not just log in to the production oauth database, but even find it in DNS. Obviously in production the DNS is different. So I have to deploy the update first, and then populate the db with new keys.

March 31, 2023

Almost nothing happened. I was busy.

Sent a PR to deploy my “script update” to staging; it was approved, so I deployed. But testing it was too late. So I’ll do it tomorrow, when I’m less sleepy.

Basically, I was just busy all day.

April 2023

April 03, 2023

In principle, everything was ready today, and production deployment was approved. And another small perf fix was approved too. So I merged both, and deployed my oauth to prod. Btw, it's a new quarter, and a lot of management is somewhere away.

Deployed to prod, yes, and then decided, let me see if my fixed api populates the db with new signing keys. Nope. It was not responding.

Ok, then I went to swagger and tried, first, the apis that I wrote - oops, it hangs, then the other apis on that server - it all hangs. Fuck. Without thinking much, notified my colleagues, and asked them how t.f. can I roll back? Oh, they say, write a pr, run it on Jenkins - with changing the build version back to what you want.

So I did, rolled it back from 1.0.113 to 1.0.112 - the version that was on prod since last week.

Deployed it. Does not work. Fuck.

So I rolled it two versions back, to 1.0.110 - and this one works.

Grabbed, in Kibana, the logs of those failures of my code. It's all Futures timing out after 5 seconds waiting for postgres, in the beginning of the server lifecycle, trying to populate the db, just in case. WTF is going on I don't know, but I believe there are two things (or three) I should put in before deploying.

First, crash if the db is not healthy on start. Second, make not more than N attempts before crashing. Third, log properly.

I know that the db is ok; I can actually populate it with psql, first generating a proper sql statement... and maybe running it from a script. Good idea, no? Maybe I won't need to check the sanity of the db.

I don't know. But it's 22 heures et demie, and I'll think about it tomorrow morning, and deploy a proper solution.

April 04, 2023

Ok, so, relatively early in the morning (9:00 CET, 3AM EST) I decided to fix my too-smart code, remove all these data base checkup that starts pummeling poor postgres in 20 threads in 5 nodes (not a big deal, I'd say, but there are probably configuration problems), and just accept the assumption that the db is ok.

It was not a big deal to kick out the code that I don't like - and run the tests. And kaboom. 4 tests crashed mysteriously - "postgres password missing". Wtf, what happened to it? Ok, git stash, try again. No good. And I spent the next hour and a half debugging into typesafe.config.

It turned out that the component of config that was coming from the config server was just empty. Empty json. Meaning, just an empty file. 0 bytes. Why t.f. will typesafe accept such a config and... ok, I know what to do with it. This config should be <i style="color: rgb(62, 62, 62); font-family: 'Trebuchet MS', Arial, sans-serif; font-size: 20px;">required. Thanks for listening.

Now ok, why was it so? Why empty file? Traced it again. Ha, it was not even HTTP 200. It was HTTP 301. And typesafe config did not bother to check and follow the link. Fucking Swiss or whatever politically-correct hackers. They got rid of good people and now what.

So I updated urls, s/http:/https:/ and all worked now. Tests passed. Then I fixed the warnings, then I ran snyk (NP, mate), and sent the PR for review.

Later that day Chris told me that I should not do code>s/http:/https:/, and the bug is in the config server, and they are working on it, and he will let me know when done.

Perfect.

Also, we tried to have a 1-1 with Karen, but first, my clock did not ring the bell, then she was busy.

Then we had a 1-1 with Jacob, discussed this and that, and future challenges, and Jacob was holding on his lap his 3-year old adopted kid; they have three natural ones, and now they adopted one. I think these people are heroes. The boy is also very beautiful. Vital_sol, if you are reading this, he is like your younger one, but with dark hair.

Then I was distracted again, we walked through all those stupid American sites that don't let you log in from France and check your bill; while we were doing this, I looked at slack again, Chris wrote to me that they fixed the config server, and I reran the tests in github/jenkins (poste /retest in the PR comments), and that was it for today.

Ok, I'm also working on our lib project, because I don't have enough nerves to work on two serious threads at the same time (probably just age, I don't know - I could juggle 4 threads at Google, but it was 17 years ago).

April 05, 2023

It's quiet in the morning. So, I slapped together a "deploy to staging" pull request, and, not much else to do, slapped together also a small "library" project, and sent a PR to review. This one, it seems, nobody was interested. But I can as well just merge it myself, since there's no Jenkins project for it yet. Later.

Oh, and in that library copilot, when I was adding scaladoc to a test file, decided that it was

```
/**  
 * Created by Rajiv
```

```
*/  
class CryptoTest ...
```

It made my day.

When my deployment PR was approved, I merged it, and tried to deploy. Kaboom fails, “timeout”, in less than a second. WTF? Mahesh had the same problems earlier today.

It turned out some keys were expired, so that... so that we can just relax and wait until Swathi fixes all this.

Also, had four meetings. It’s the start of a quarter, so everything’s sleepy and dreamy (of future achievements).

One was the “standup”, and it was ok. Next was a meeting where... Rajiv? Someone like that, I rarely see this guy, was talking about the problems he’s solving, something with ads serving etc. He was talking fast and not very clearly, so I turned off the microphone and did my French.

Then a 1-1 with Karen; it’s mostly about life, since all technical questions are being addressed immediately (but it’s up to Swathi to fix).

Then something like a “planning meeting”, but Renato is away, so Jacob was leading it. And ok, ok, what can we accomplish next sprint, or something. As much as we can. The “velocity”, as they call the speed, really does not depend on development speed. We do have impediments, and we can’t do anything about it. So I was doing my French.

Then, after that, Karen sent me links to two of my Jira cases, finished and closed, asking me to add more text (that’s something like documentation). That I did with pleasure.

April 06, 2023

So, kind of exhausted again. In the morning, there was no trace of good deployment to staging. As a result, I finished my library PR, sent it for review; later Karen asked to update the README, which I did. Approved, merged. Now I’ll have to create a Jenkins job for that project. That’s tomorrow.

Since there was nothing much to do, I’ve also updated the project that contains our scripts, adding several of my own scripts, sent it for review.

Later today it was disclosed that now jenkins deploys somewhere else. To another pod, or whatever it’s called. Another k8s. And I’ll have to update this and that in my scripts, to go to something like “server23”.

I was clueless; and Chris was also telling me that my service’s calls to config server were failing.

Eventually, by the end of the day, I got all the urls, configs, server names, etc, to put into my script.

Discovered that actually the deployment succeeded yesterday, but Swathi only disclosed the deployment pod name, or whatever, today. Ok, it's morning over there.

So, everything works for me, and since it's already Thursday evening, rushing with deployment to prod makes no sense (we don't deploy on Fridays), and I have a good slot of time to do performance tests, stress tests, etc. Will just write some code for that, and see how it works. Have to pummel it to get tokens generated for given signing keys, from the real service on staging. We do have some performance tests, but not many.

Tomorrow; I'll have to sleep on it.

Tonight I'd rather go watch "Les Trois Mousquetaires" in our cinéma.

April 07, 2023

Today is something like almost a holiday, I'm clueless what it's about. Also, Renato is away, and Jacob is busy with the court regarding their adopted kid; in short, there's nobody, almost.

I spent some time figuring out how to slap together a stress test on staging. Ok, started with writing some code for statistics - funny, but it's not obviously available.

Oh, even before that, I decided, what the heck, let me write it in ZIO 2. Ok. ZIO 2. Nothing works when you add it to our codebase. Pinged Jorge, but he's busy with other jobs, I guess.

So I decided, next time, I'll just fill the database with hundreds of imaginary users, and will pummel the staging system on their behalf.

Let the guys (and gals) know, and they told me not to dare do the same in prod (no worries, I'm paranoid too), and then... and then we had a "team's virtual offsite". There were five, then seven of us; they wanted to play some games; I'm really not good at games, so I said goodbye, and that's how my day ended today.

But tomorrow I'll create hundreds of clones in our staging database. Just too much armagnac today. And too much nerve regarding the things going on with some of my relatives.

April 08, 2023

What did I do today. Getting a clearer understanding that with the tools we have, I'll slap together the necessary shell scripts way faster than I write it in Scala. Probably I need to fix this problem. If we had `curl` in Scala, I could write it in several lines (plus a couple of json files). But now... I'd say, it's a defect, it's just a defect. Why stuff together all those futures and actors? No shit, `play` uses actors to send an http request and get a response, and also requires a "reader" for that response. In shell, `jq` reads json professionally.

Oh, whatever. I've just created 500 users via curl in shell, when I figured that Scala, with all that nonsense that we have here, should be thrown out as a tool for making tools.

April 10, 2023

In the morning I slapped together scripts that create tokens for those hundreds of users that I created earlier, with distinct providers. Ran them in parallel, in 5 processes (or 6?), got the output of the time it took, and then in Scala got stats from those script runs. Ok, the range is 0.7 to 1.9 seconds; once a day per user, no big deal.

During the standup notified the guys about my readiness to deploy to prod, and ok, it was approved, as well as my other pr, with new test code for experiments.

I planned to deploy earlier, but Mahesh and Serkan organized a meeting regarding how t.f. to make sbt work. Turned out that if we don't use Nix, our documentation sucks, and Nix also sucks without good maintenance.

There were about six of us, helping the guys, and you know, it worked! Serkan never dealt with Java or Scala, so for him it's all a weird mystery. For me sbt is a bit of a mystery too, but well, we together know definitely more.

Then I deployed to prod, and looked at the behavior. It was ok. Then I decided to check the behavior with the new keys. And oops, no trace. Then I looked in the database, and kaboom, new keys were missing. Ok, np, I have a script for that. Got the new keys again. Now what. Now I had to test. For that I had to find profileId in prod, and sessionId. Normally, they are in X-something headers in requests: I open our company's page, started some music (Ok, Prelude and Fugue in c minor), and looked at network exchange. No such headers. Karen had such headers. So I copied her id, and she told me how to grab session id.

So, I wrote my scripts, and found that one new key works ok, but another (contains space char in it) does not work. Fck.

It's almost 23:00, and I've just sent my pr for review. No luck, actually: Karen is off by now.

Anyway, I have stuff to do, and some keys work ok, and if it gets urgent, I'll just rename the provider, to get rid of spaces (stupid, right?)

One revolutionary thing was, I added logging, and limited logging by a date, May 1st, 00:00 - by that time everything will be clear. If Karen asks for keeping it forever, I'd love to. But so far I heard strange ideas about saving logs space. Heard it from Csaba a while ago. As if space costs anything, and our time costs nothing.

April 11, 2023

Ok, I was pretty nervous, why t.f. do not the keys I post get into the system on prod? That's why I added logging and tried to get an approval to deploy this stuff to staging.

Of course my innovative way of logging (setting time limit) was not understood, and they, traditionally, asked me to use logger.debug(), and update the logback.xml to enable this logging. See, the operation happens several times a year maybe, but we are saving on disk space. And of course if you enable debugging in the config file, who is going to disable it? I see comments in those files, like "This is here very temporarily, will remove it soon" - from over 2 years ago.

Oh, whatever. It's really not in my interests to push the boundaries of vision in the conservative people that are basically responsible for the functioning of the system. Should I replace them? No way.

But how will I figure out wtf is going on?

The problem is, we don't have any tests that contact staging or even prod and check the behavior. So I'll need to have some. Meanwhile, I was just running scripts.

And I found that, hmm, the data gets into the database. (I have psql open.) But the server was not showing it. And then, after I took a break, it did.

OMFG. It's cache. I have cache. I intentionally used guava cache (no shit, it's something that I, together with Bob Lee, aka crazybob wrote when I was in Google... then, of course, there were lots of refactorings.

So, everything works, and I was just paranoid. Pinging cache to refresh would not help, because there's more than one instance anyway. And it was intentional: do not rush, it'll settle, we'll have <i style="color: rgb(62, 62, 62); font-family: "Trebuchet MS", Arial, sans-serif; font-size: 20px;">eventual consistency.

Fck. Of course, the mood has changed.

Then I went for a walk, then we had a meeting with Jorge, to figure out wtf is happening with using ZIO 2.0 in this project. Ha. It was Scala 2.11 - who, in year 2023, uses Scala 2.11? Conservatives only. So I decided to switch to Scala 2.12 (why not 2.13?) - and that's where I'm currently stuck.

Then we had a 1-1 with Karen, and talked about logs, and, since I got on Kibana through all the logs produced by this app, oauth, I know what this code logs - mostly spam. Karen suggested me to clean the logging, I suggested to add memory for logs. And went ahead with looking into where all this shit is coming from. It's not from our code. It's from k8s. Basically, none of our business. Just filter shit out. Ok, found some weird error messages, and found that I could improve the meaning, so I updated just one little line, to make the log message more meaningful.

That was all for fun, because I did not plan to take the next project today. Tomorrow, tomorrow. Mañana. Morgen, morgen, nur nicht heute. Утро начинается с рассвета. Manhã a bonita manhã.

At this moment I had almost closed my work computer.

But not exactly. There was one more meeting, where Davis suggested to discuss a new endpoint. What's there to discuss, the endpoint exists already, and he wanted to give it an external access, and was asking whether it's secure. It sure is.

April 12, 2023

In the morning I found that a bunch of people want from me examples and explanations, how can they validate the tokens that the server generates. Ok, so I wrote a README.md, explaining in details all those manipulations with keys, tokens, base64, and why the last decoding fails after showing a decent json (the reason is: the command-line base64 hates non-ascii chars, which the signature contains).

They were happy, but in the end of the day they wanted a token for another provider, and ran the script again and send it to them. How to make people happy by about 200 characters of base64 that they can't read. Interesting, right?

Then I experimented with ideas regarding storing json inside bson (alternatives exist); then we had a standup, then we had a postmortem for sprint (it was a good one, we seem to be thinking in parallel. The team cooperates - I never saw it on the West Coast.

Then we had a planning meeting. Pretty boring, but at least agreed on the plans. The topics of code quality and tools quality... Customers did not ask for it.

When I returned from a walk, I reread what Chris wrote to me about restarding swagger docs server on prod; did it, and the docs now look perfect. Also I generated another base64 that people asked in the beginning of this chapter.

April 13, 2023

A special day. We had today a “Permanent Production Meeting”, like they had at the Soviet factory I worked, 49 years ago. This all-day meeting was scheduled for, basically, the whole company (not exactly, but a lot). While they were talking, I opened a can with ratatouille, heated it... omg, omfg! It's so good! You saw the movie, right? I'm like that connoisseur in the movie. It's just beautiful. Even from the can.

Before the meeting started (at 16:00 CET), I had a chance to work, to do some docs updates, to figure out how to hold my experiment with JSON and Mongo - basically, I decided to go the short way, split two spikes, and do them separately. One, switching from pull to push (accepting deltas), and another to make it free-form. There's the third problem, dealing with Lucene search, but since I have no plans to kick out any data (later, later), it won't hurt to postpone it. Wrote to Renato about all this, but it was early morning for him.

Then I decided that unit tests with mocks don't solve the problem of testing whether the real Mongo will handle all of this properly (I already have tests that use mocks, and, well, what does it prove? Nothing.)

Started writing all this (kind of copy-pasting from my code in another module: making the library module usable will take some time).

Then I could not find a way to charge my work computer. Or my home computer. Oh, WTF.

Spent most of the rest of the day trying to solve this problem.

When I was in Agen, I at least checked slack (and sent my standup report); found that our partners (consumers of the stuff I wrote) ask strange questions about the contents of the tokens. Could not do anything while I was in my car, sitting in a parking lot of Boulanger, under a pouring rain, so shared my doubts about their request. Will answer in details later.

It all works now, and I'm in a meeting; I like what the guys talk about, and the plans (ok, I'm a naïve optimist), but I can't go into any details. Only when a guy was talking about Triton tokens, he was saying strange things, so I (impolitely) interfered, noticing that tokens are not encrypted, only base64-encoded, and if we want them encrypted, we can do it, but we have to communicate with Triton regarding that.

Oh. Anyway, the meeting has ended earlier than I expected (it's 21:56 now, it was supposed to last till 23:00). So, let me write my blogs and go to bed.

Good night.

April 14, 2023

Had a "slack standup" today, and a "token meeting", and also a slack discussion.

And what did I do? Kerry and Karen were worried about the tokens I generate not having a payload. So I spent half a day trying to figure out wtf, and where's the payload. Turned out that JWT creates a token like this: - header map - payload map - signature

All three base64url-d separately, and concatenated with ‘.’ as a separator. Karen even pointed, after I discovered all this and posted a two-liner script to show the data, to an amazon site that does exactly this. Well, amazon site is hard to copypaste into a shell script, right?

Anyway, problem solved.

The "token meeting" was about using more tokens, in more places (since we have a success with the triton token, and since we now have a token pro (which is yours truly)).

Then I took a break and after that turned to my experiments with "station data". What I need is a good integration test that actually talks to Mongo, not to a mock that can be faked without any practical

validation.

So I had to make the tests to get information from the proper properties, etc. Doable, but needs to be done.

That's it for today. The tests are running now, but I'm too sleepy.

April 15, 2023

Strangely, I decided to work today. On that project. Strangely, it was not even compiling in intelliJ - "duplicate class" (generated by swagger) So I rolled it all back, pulled from master - same problem... almost. Eventually I figured. There's NO NEED to have a separate directory for integration tests. Just put them together. Call the tests that are IT "ITest". And that's it. Will make all those dependencies much easier.

So, I did run the tests this way and that way, and wrote to my colleagues, and that's it.

Also updated my work calendar, notifying that I'll be away half a day on Monday, and also Wednesday-Friday, when I'll be in Lisbon (inshallah).

And closed my work computer.

April 16, 2023

Strangely, I decided to work today. On that project. Strangely, it was not even compiling in intelliJ - "duplicate class" (generated by swagger) So I rolled it all back, pulled from master - same problem... almost. Eventually I figured. There's NO NEED to have a separate directory for integration tests. Just put them together. Call the tests that are IT "ITest". And that's it. Will make all those dependencies much easier.

So, I did run the tests this way and that way, and wrote to my colleagues, and that's it.

Also updated my work calendar, notifying that I'll be away half a day on Monday, and also Wednesday-Friday, when I'll be in Lisbon (inshallah).

And closed my work computer.

April 17, 2023

Starting this morning, worked on my mongo stuff - it turned out that writing was not even declared. Only reading. Writing is in another project.

Mongo is weird of course. So spent time, finishing it right now. Hopefully.

Had a 1-1 with Jacob. I got a feeling he wants me to do bigger things. Well, with pleasure. They definitely need an architect, but do they view me as one? See, I'm not a big fan of careers, never made any, just enjoyed doing what I was doing. I sure did design a lot of stuff. But do I care? What do I care about? Fun.

April 18, 2023

Ok, so my home computer is dead, posting it here for now.

Spent the morning trying to figure out how come I write records to Mongo, but when I read them, some are missing. Abandoned Futures, waited for everything. Etc.

And eventually found that the mongo library for dealing with these specific records has a data filter; it ignores the “genres” listed in the config. Not coming from user requests, but from a config. Of course IRL no config contains anything like that (and nobody knows), but in the test config it lists two “genres” to be ignored. OMFG.

Anyway, tests pass, sent a PR for review, and, generally, I'll need to do two things:

1. Run the same test (data compatibility, new format and old format) on the whole import dataset
2. Write a design doc, or rather two design docs. That part, I'll do it tomorrow morning.

And regarding import, spent some time looking for the module that does that import. Found it, everything's fine. Tomorrow.

As you see, I'm going to post my dybrs right here, for a while. Writing from my work computer.

April 19, 2023

I'm at a Scala conference, in Lisbon.

Not much to write about work, everything's just fine. Wrote a design doc in the morning, with a table specifying what has to be done, which items depend on which, to which spike each item is related, and how many points each. Sent it to them. The management seems to be happy.

April 21, 2023

I'm currently in Lisbon, on a Scala conference. So, I didn't spend much time working.

In the evening found a question from Karen, how can we update the keys in the database. Wrote to her. There are examples in the script, but it's not documented; bad, I'll have to update the docs.

April 22, 2023

It's Saturday, and we are walking around Lisbon, especially Alfama, where we are staying.

But I suddenly got a precious idea of refactoring for the project I'm working on. It's hard to stop the brain, right?

April 25, 2023

Another crazy day. In the morning I updated the design doc, slightly changing the implementation plan, basically because of those two lines that make things easier.¡

And everything was fine, except that I really did not want to attend a standup meeting while sitting in the airport waiting for boarding. So I just posted my progress on slack.

Oh, and I was investigating the tricks to do to add a project to Jenkins. Basically, it's just another pr, but there's a puzzle. One thing is to add a groovy script with description, and another to add three totally obscure xml files - as I understand, these files were being copy-pasted with just project names being replaced. I think we could make it all much simpler, but nobody cares. So I copy pasted some shit, in a minimal volume; let's see how and whether it will work for me.

Then our "customers", I mean, people in Houston, reported problems. First, a "provider name" was wrong. Ok, it was wrong in Jira. Ok, let me do a db update. Which meant, I added another provider, with a correct name. Then they asked me to remove the old one. That would be a problem, because tokens were already generated, and they point to that "wrong" provider. Then it turned out that our provider ids are wrong. Because we should take ids from them (and why did not they give us those ids?) We were just generating ids (in a shell script, so replacing them with parameters is really not a problem).

And when I went ahead to do production db update, the airport internet decided that I had enough of it, and disconnected me. Why so? Because, well, our flight was delayed for 4 hours, and I overused the airport's internet. Ok, so I pulled my French phone - but it decided not to share the internet on wifi, so, it was the first time, I connected to it via bluetooth. The internet was ok, but port 22 - no way. So I turned off something (FIIK what) on my French phone, and the internet became available. But not port 22. So I told my colleagues that oops, later, later, probably 5 hours from now. Because they changed our departure time again, now it's not 20:45, but 18:30.

That's almost it for today. I think.

I also could not pull some code from some repo, because, again, of port 22. Will handle it later tonight, or tomorrow. Tomorrow I'll have to present my design doc, btw.

Having come home, logged in to prod db and updated the records I promised to update. The guys are in Texas and San Diego, so it's still daytime for them.

April 26, 2023

A little bit crazy day today.

First, I'm on pager duty. Some services are out of resources, etc. Started a discussion dedicated to this - how do we end our "investigation"? Jira case? Or just "got some experience"?

Anyway; and also studied two jira cases assigned to me, where our Texas partners decided that they now know what they need after I implemented what they asked for. They wanted to pass us a providerName, and I did it all. Now they want instead to pass a providerId. Instead? I'd rather make both optional. Who knows what they decide to choose in the end.

After a long process eventually Karen disclosed that they (the partners) have one jira case to do, after which we will have the data, and I won't need to take 4 points to implement all that shit. Nice. They never told me, they told Karen.

Had 4 meetings in total, now sitting in the fifth. No energy left to work on code.

Ok, tomorrow a big meeting where I'll be doing this "show and tell" again. Nobody prepares slides these days, just switching files in IntelliJ seems to be fine for the people. Nice.

April 27, 2023

I think I work too much.

Ok, I implemented that stuff with providerName and providerId. And then got stuck with tests - just added some tests, and found out that neither the new ones, nor the old ones behave properly. It's almost 22:00 now, and I've just finished fixing my tests. Swagger docs is next to work on.

Anyway, I'm also on pager duty, except that the pager is silent (magic! Thank you Karen!)

At 5 we had a "standup", and at 5:30 I was supposed to present my stuff - refactorings from pull to push and from "strict types" to "free-form" json. At 5:15 the "standup" ended, and I figured that having all the files ready in IntelliJ is not enough, and I need to slap together slides. Ok, I have 15 minutes. Meaning, no time for fancy pictures, just well-organized text on 5-7 slides. Exactly what I did. Another funny thing was figuring out how I can show about 100 lines of Json in a slide. The solution was: get out of "presentation mode" and show it right in Google Slides, scrolling way down the window. Impressive.

So at 5:30, when a crowd of managers gathered from around half the country, I spent about 10 minutes with these slides, then showed them the work schedule in my doc, then the code in IntelliJ. Then there was a discussion. One idiot-architect started telling us how important it is to have strong types (being passed around the internets, right?), and how our client software should abandon everything it was doing and start using GraphQL to request the jsons that they want from our server. Also, our server

is not required, because they already have a server that has all that data (I asked if they are ready to serve those millions of clients requests, which we are currently serving from our cache).

So gradually we kind of, magically, forced that idiot-savant to shut up and listen to more reasonable people. The only question was whether I'm going to pass a "json blob" as a separate field in the regular json response. Of course not. The response will have the same structure it had before, but enriched with whatever they want us to cache.

So, everyone was happy, and we closed the meeting. Nice! East Coast culture!

The only sad thing is that Karen could not attend it, because she's in Tokyo, and it was over midnight there.

Then I took a break and continued to work.

April 28, 2023

By the time the day ends, I'm totally exhausted, and hardly remember wtf was going on today. So, just a list of actions.

1. Read their instructions regarding OKR, and filled some forms. Writing this kind o stuff... did it for the last 40 or something years.
2. Spent too much time trying to push a pretty obvious PR to github - but some of the tests were failing in a weird way.
3. Causes:
 - a) An exception thrown inside future is not so trivial to propagate to the top, especially when the top is not on top, but inside a closure;
 - b) Found 4 NPEs in AWS libraries; one their own (a variable not initialized - Java!), and another in java library, in which, to do base64 encoding, the author calls some JAXB, which uses some "ModelBuilder", written by Kohsuke Kawaguchi, which pulls some XML, which is not present, so that the "ModelBuilder" just throws an NPE.
4. Cleaned up the code and the Swagger docs - but Swagger is crazy, it counts leading spaces, so there.
5. I'm on Pager Duty, so there was a huge latency (5 sec instead of 1 sec, and later it was even 20 sec); I found the piece of code that holds this back, but there's no logging there, so go guess.
6. Had a meeting at 20:30, "what should we do to fix stuff"; I posted a poster "do more logging", and everybody was against it.
7. Instead they just guessed that one service (our partners) was too slow to return the data (as if there's a lot of data and we can't cache). They don't need logs if they can just guess.

8. Also was told that during PD week I must read emails, because that's where customer support sends us their complaints. Oh shit. I have almost 500 emails in that microsoft box.

That's it for today.

April 28, 2023

Ok, again was busy all day. First, Karen wrote that encryption for staging and prod is done on different credentials, so I had to fix my Crypto class accordingly. Fixed, sent the PR, and then I was waiting for Karen to rereview my PR update. Also, I was supposed to test the key management on staging, and for that I needed a couple of real credentials, which I forgot where to get (and it's not documented). So I was working on the lib project. Not 100% trivial, but WIP.

Oh, also Karen asked why I kicked out EitherT[E,X] in one place, replacing it with Future[Either[E,X]], and I told her that EitherT is against any reasonable science, it was just a hack slapped together by not very educated Cats people, based on the Haskell fashion of the old days (now deprecated). Yes, physicists turned computer scientists are creative, and they don't have time to read math books, so they like to invent all kind of illogical shit. At times it's an advance, but please, read the books first.

And then we had a "standup", it was already 17:00 CET, and I've already returned from my daily walk. Seems like I will be able to deploy to prod.

Except that my script to push more signing keys to staging (or prod), it was pathetic, basically because I did not know how to concoct a command to execute; it's definitely not just one string. (sigh).

Eventually wound up writing a function to send new keys, with params etc. I'm not a big shell programmer, and copilot does not always understand what I mean. But well, meanwhile I was also updating some docs, and copilot was really helpful.

Then we had a 1-1 with Karen, and she showed me the trick to get session id etc, no need to log in to staging website and look through cookies. Just an API call, curl etc. I just scripted it immediately.

Then I updated the records, removed all test records in staging, and prepared the commands to update records in prod. But first have to deploy. So I sent a pr for prod deployment. Tomorrow morning will do it, while everybody except Mahesh are asleep. Mahesh is in IST time zone.

Oh, and also, one guy was asking for help (and I'm on the second tier of pager duty), so I investigated, talked to our guys if we can let him use staging, and it's ok, so we told him so. The so-called "qa" cloud does not get maintained regularly; last deployment was done by, I guess, Joe, over a year ago. No trace in Jenkins, it's kube that provided me with this. Ok, ok.

That's it for today. I know, I know. It's boring and may sound stupid. But I'm having fun, and people are happy. Remember Algernon. I do.

May 2023

May 01, 2023

Another day.

I'm still on PD. But everything was quiet until about 19h our time (10AM in California). When the latency started getting up-up-up. But this later.

I thought I wrote that 2-liner, to get the data from JSON together with JSON. Turned out, no way. Some tests failed, and see how. When you refer to a field in JSON, you specify a path, like `(__ \ "customer"\ "name")` etc. What I need is the whole json. I wrote `(__)`. It compiles, and some tests pass. But then, the reader sees “an empty path” - and crashes. I told typesafe people, years ago, that they need category theory specialists in their library. They were clueless. They were saying they needed a compiler guy (I'm not one). So, they have a bunch of BS inside their play library, including a pretty incompetent approach to monadicity. Like in HBASE, same thing. When (speaking in SQL), `select * from users where age > 200;` will crash (since there are no such users (yet)).

Oh, whatever. Posted a question on stackoverflow (I know, chatgpt will be clueless), and had a nice discussion with another clueless guy, but eventually came to a conclusion: I'll handle it manually. It's not hard. Just takes much more than 2 lines. Maybe 20 or 30.

Then started looking through email, because Davis said that while on PD, one has to go check the emails from customer support and from users.

We had a 1-1, I told them that I a) have one PR from last week, closing one case, and another from this week, closing the case I opened regarding statistics, and plan to maybe take a look at this case that I'm supposed to do this sprint.

After the meeting, discovered a comment from Davis, on our team's channel in slack, telling me what I should do and what I should not do. I should not do what's in my sprint, because I'm on PD, but instead I should look for cases that belong to support, for example, he recently created a case, sent a pr, it did not work, so he rolled it back. Now, he says, I should take this case and implement proper caching. (I know, I know, we use guava cache, which comes from Google, from the guava cache that I wrote eons... 15 years ago - then it was rewritten by some guy called Fry.)

Sure can do. Except that he could have asked me. With respect. Maybe casually. Like, “Kirsha, could you please take a look at this case, I need help”, or something like that. I'd have done it with pleasure.

Oh, and who is Davis? Is he my manager? He's nobody's manager. Just an engineer. So I just thanked him, saying that I appreciate his attention to my work.

But sorry, now I'm just busy on alerts. Coming with an increasing frequency. The same latency thing. That “should not be logged”, because we have too many logs. Yes, I did deploy to staging the small fix that disambiguates between two operations - but I did not have time to even look into staging's logs.

Maybe they are ok. Then I can deploy to prod (thank you Safa and Swathi for approving fast). Then I can at least partially disambiguate, what kind of requests are causing troubles. The same as Friday. The same time, around 10AM PST. When Bay Area engineers and students are driving to work and turn on our app to listen.

BTW, why not log the location, for instance? Weird. We could.

Anyway, this was a good reason to relax for the day. Have to get up early tomorrow. And it's 20:41 already here.

Bonne nuit ! A demain !

May 02, 2023

I had some fun today. Yes, I was working on that “two-liner”, trying to bypass all those weird behaviors of play json, see <https://stackoverflow.com/questions/76147861/scala-play-json-how-to-capture-the-whole-json-source> (doable, yes).

And waited for people to approve my prod deployment of that stats update.

Meanwhile people over there were getting pissed off because I was not fixing production problems.

How the fuck can I fix them? “Processor is taking 90% of its capacity” - is not it good? “HAproxy cannot access this sever” - fine, let me look into that server’s logs (DynamoDB was the culprit).

But by the end of the day (Karen is on vacation), after Jacob was pissed off that I did not spend (my) night trying to fix the dead Jenkins (0 docs, and nobody was answering my calls), and Davis was kind of confused regarding why I’m not fixing the stuff, I probably convinced Davis to walk through this shit together.

Turned out that HAProxy is running in the same kubernetes whatever (cloud), and we can find its name and ip, and if it decides that it’s too busy or cannot receive responses from the servers, we just kill it. And it will resurrect again, this time doing good. Since it’s HAProxy, not our code, nobody fucking cares what exactly happened. Just shit happened, kill it.

Then it turned out that it’s all written in the workbooks. Except that links in the workbook are too old and don’t lead anywhere. And that the aws/kubectl scripts that I had since last year don’t work anymore. Because they changed urls. And Swathi even wrote about it in slack. In devops channel. Should I be reading devops channels, figuring out what in my scripts I must change? What was changed, mostly, is eks23 was replaced with eks21. Except that it was not done anywhere. So I could not run my scripts, and instead I had to just run this command and that command, as suggested by Davis. But the result was the same: switch a context, find a pod, kill it. They do it all the time. Manually.

It’s 22:25 now, and I’m eventually done with it.

Ok, also, investigated another “user case”, and eventually wrote a “spike” regarding some obvious things: “let’s log users actions when a profile is being changed”, “let’s keep the user profile history”. Because somebody changed a user’s email, and there’s no trace. And it was a hack.

Oh, whatever.

Yes, I could retire any moment. But first, what I’ll be doing? Walking around the town? And second, stocks just dropped today. Except for GLD that I was thinking about buying, yesterday. Oh, whatever.

May 03, 2023

Ok, I’m totally exhausted after last nights late shift with Davis. Davis, btw, is totally sick and did not show up today.

I just can’t work. My PD ended today, and I’m just taking some rest. Reminded that I have three PRs to review. Late in the evening Qian did leave some good comments, so I fixed it. But I could not run tests locally, because I can’t log in to aws, because they did change something, as they always do, without telling us what we should do now. Anyway, ran the tests in github/jenkins.

What else? Nothing else, I really could not do anything today. I know what to do, but I can’t. So I was doing my French.

Oh, and since we have a slack channel “pg handoff”, Renato asked me if I did the handoff. I said that it would be really cool to have a ceremony of handoff, something like what they do at Buckingham palace all the time. We agreed that it would be cool. Then Qian also commented that there are things worth writing during the handoff. Ok, it makes sense too. Like, ”I was working on this product problem till 10am, my shift ended, so here we are no, feel free to proceed.

Does it make sense? Not to me.

May 04, 2023

Ok, did I work today? Kind of. I have one problem to cover properly, but it requires focusing, and I did not have time for foThat’cuses. So I just slacked.

Well, managed to deploy to prod that service where I wanted better stats being collected - but since the guys changed the prod services urls (why?!), I don’t know now where to look. Well, whatever. Next time I’ll be on PD it’s June (and then July... oh.)

So I created a spike in Jira, specifying the problem with our HAProxy, and possible solutions. We’ll talk.

Of course while I was getting an approval for deployment, someone deployed something, and I had to bump the repo, and get an approval again. But do I care? I don’t.

Also, my “oauth” deployment got through, and I’ve just deployed it, checked the API manually (why?! Need a script or something); discovered that swagger doc needs to be updated, and updated it. Thank me, I documented all this shit, so I don’t have to remember.

That’s it for today. Tomorrow I’ll focus on the perfect solution for handling play json, the one posted by Michael Zajac on SO.

May 05, 2023

A nice day today; some pieces in those “json transformations” were missing, so I used the time to make them all work. It’s 22:01 now, the script runs to test and push. I guess it all works now.

The problem is, I lost my working glasses, somewhere at home, while looking for some books. Weird. So I’m using glasses from the store, and it’s not very comfortable. But anyway.

By the end of the day we had a happy hour (virtual). Since I’m skipping a week regarding alcohol (just testing my params), I was drinking tea. Also, seems like we’ve connected with Renato on Strava. Others are not on Strava (why?!). Talked mostly about ZIO 2.0, Scala 3.3, and the necessity to stay away from weird stuff like Nix - which is nice, but not self-maintainable (needs a dedicated person, trust me), and is subject to bit decay.

That’s it for Friday.

Well, also our partners unblocked the other case that I was supposed to do for them (they did add a param to the request they send to my oauth service). So, I wrote to them: will do next week, with pleasure.

May 08, 2023

It might have been stupid, but I was too tired after yesterday’s long hike, so I just sat on my couch in the morning and spent all day working. It’s already 22:14, I’m still working.

Just added a field to an oauth token, right? The trick is, the value I receive to add, from a request, may be incorrect (it’s being used for some stuff).

And dealing with passing an error properly from function to function, in a Future, was a bit of a challenge. Ok, the code is not beautifully written, it was written Tolstoy’s style, long sentences with complicated interconnections. Ok, Platonov. I prefer Marshak style. Primitive. Clear. With no side meanings.

So, that’s why. Currently just one suite is failing, and it’s good: it means it does check the behavior. I’ll fix it tomorrow, no rush, this jira case is for the next sprint.

Another PR that I posted last Wednesday for review, is still just hanging there. Chris reviews Davis’ code, Davis reviews Chris’ code. Mahesh is new, and he’s too shy. Karen and Qing, two nice Chinese

people, are on vacation.

No, it's a nice company. But as soon as you give developers "too much freedom", they start building ethnic groups. What can one do? I don't know.

Ok, I reviewed Mahesh's PR (had a question regarding some mismatches, but he said it's ok, so I approved).

I think I have to talk to Jacob about this. It's not very healthy, and we need a stable solution to this.

FIIC much.

May 09, 2023

By the time I'm free and can write my dybr, I'm usually too exhausted.

So, spent all day, again, working. Most of the day it was that PR that Qian is now reviewing, and he does it so well, I appreciate it a lot. The idea of making MongoRepo not a superclass but a delegate, it was right - now we don't depend on the implementation. Almost. Anyway, an adapter instead of a wrapper. Cool. And he found many little details, some several years old. Good, I can fix it all now. Also, the errors management in those methods that return Future[Either[E,T]] was really pathetic. Up to throwing exceptions in certain circumstances.

Another thing, I had another PR which I did not finish, there were a couple of tests I had to make run.

So that's how I wasted so much time working. Ok, whatever. That's it for today.

May 10, 2023

Ok, my computer overheats all the time. That's my work computer. I kind of guess who warms it, but well, is it not my business.

I had two prs, in two projects. One is still waiting for a review, although the case was supposed to be finished today. Not my fault. Another is deployed to staging, and I checked it manually, everything's fine. This ticket is for the next sprint. Well, whatever.

Had a standup, then sprint planning, then the previous sprint's postmortem. Nothing to talk about, everything's fine, total mutual agreement.

That's it for today. Friday, they say, we are supposed to work on those "infrastructure tickets" that we created. Technical debt, it's for Fridays.

May 11, 2023

Spent most of the day working on two prs in the same project. One developing a feature, another just building a minimum delta that would be a part of that feature. Also, my “oauth” pr still was not reviewed.

Also, discussed with Renato when a spike should be closed. I thought (and I believe I heard it from Karen) that a spike is closed when documented and provided with a list of tickets. Now Renato says it’s when all those tickets are closed. Ok, I don’t mind.

But there was also another spike, which I closed earlier, mostly because its solution is covered by another spike.

But now I had to reopen it, and since there was a lot of progress, maybe this spike should consist of one 2-point ticket. Ok, I don’t mind. But first I’ll have to write a design doc and give a talk to the team regarding this.

Maybe I’ll write it tomorrow. It must be one or just 1/2 page. “How to get data for search engine with no pull involved”.

Ok, never mind. I’ll write it tomorrow.

May 12, 2023

I’m so tired that all I can say is that the day is over. Deployed oauth to staging; still talking about one small PR with Chris and Davis.

The rest is ok. No, never had a chance to write a small design doc. I was driving.

May 14, 2023

Spent some time working: investigated an old code (that’s a kind of archeology, the code dates from 2014, and nobody know not just what it does, but even whether it exists, and that it works). Here’s how they built it in the old days!

In short, I figured that I have nothing to do in this spike. Everything’s done, eons ago, and new updates in the other code have no impact on it. Everything’s designed and built the way I wanted to do it. Except that they use Kafka and RabbitMq, but since it works, it’s fine. It’s just fine.

May 15, 2023

By the end of the day I don’t fucking remember what was it.

Ok, in short. Deployed (I thought) oauth to prod, but kaboom, my stuff does not work.

Turned out, my merge to master failed, so I deployed the old stuff. Ok, merged again, deployed, checked, notified our partners. Cool.

Also, my other pr was approved (it was all about parameterized id trait); ok, so I merged it, and slapped together another one, threw it for review. Karen is still on vacation.

Also, Jacob wrote some notes for my OKR; big deal, I updated them. Who fucking cares about their bonus. They could raise my salary, but do I care. It's all peanuts. I'm just having fun. My work is just a hobby.

That's it for today.

There was a couple of tech talks, but I can't talk about it. The second one was pretty stupid: an Indian guy was pushing some rather stupid philosophy about microservices and about "circuit breakers design pattern". I don't believe in design patterns, and he was just showing some pictures. Pretty stupid.

But do I care? No, I don't.

May 16, 2023

Again spent too much time working.

The service I'm working on, I'm splitting my pr into smaller chunks; one was approved, merged it, and slapped together the next one. While they are reading it, took another project to use my new service' library.

And kaboom, that new service, it consists of three modules; two of them made it to jfrog, but the one I needed did not. So, what should I do? Spent time investigating wtf is with jfrog. Eventually (also bugged Swathi and Safa) came to a conclusion that it's all in Nix, and Nix config is hiding somewhere but not in github. So, while Karen is not here, nobody knows.

But my approach is simple: this module will be gone by the end of the year, so working on its perfection makes no sense. I'll just copy-paste what I need. Will get rid of Nix later.

So I did, and slapped together the PR and sent it for review.

Then we had three meetings, including a 1-1 with Jacob. Jacob is one of those four decent managers. No, I'd call him a decent human being, first of all. With an adopted child. These people are almost heroes.

Well, our PM is also really good. And smart. And a cyclist. We are now friends on Strava.

Jacob was talking about the importance of OKR for the bonus. I'm not telling him I don't fucking care about bonus, it takes much less time to pay more attention and win this amount on stocks, without risking your integrity.

But anyway. I'm just having fun. That's why I'm still doing something, after 23:00.

May 17, 2023

Since I'm starting this Pager Duty today, I decided to throw as much as possible for the guys to review, and then take a break (until something breaks).

Ok, I have the code to review, but nobody reviews it. So what I did? I decided to investigate, how come I can't make my new "library" project get through Jenkins. That's archaeology. But as a side effect, I found how to make a module get to jfrog. Just remove one line in sbt file. "disablePublish".

OTOH, I don't need that anymore.

Meanwhile, one node in AWS got broken (AWS is investigating), and somehow no other events.

So I relaxed, went for a walk.

I skipped one meeting: my calendar did not ring the bell, FIIK why.

At 9:30 we had a spike meeting; I can't give the details, but eventually decided that Davis may need to investigate the volume of the data we will store to decide between keeping data in Mongo or in S3 and refer from Mongo. Qian says that Mongo can store a lot. To me Mongo is a stupid toy made by people that love adventure and don't care about data. Anyway, the discussion was good.

May 18, 2023

Ok, so I'm on pager duty, and also in the process of getting two PRs approved. Qian is good in reviewing, I appreciate his comments even if I kind of disagree.

Anyway, and there was a failure in prod, the system called me, I opened pd app, followed the runbook, killed one job and restarted it, all by the book. Now the question is, how come we don't have scripts for that?

Also, investigated, how t.f. can I properly create a new project in Jenkins. It's a mystery. While studying this, discovered tons of petrified Java shit since 2014. The good part is that it still works!

Then I got tired and took a break.

Then Chris was strangely sending me pointers to where else I could find failures. Devops channel in Slack. I don't even understand what is all that. I mean, that AWS management. Besides that it's what our devops are working on. How can I help them? Maybe I should come over to Oregon's Amazon data center and help with some wiring? I've no clue.

BTW, our QA, Priyanka, works from West Africa. Cool!

Then Chris found that I did not read an email from customer support. Just another case of a user whose email address was magically changed to some nonexistent domain. I created a jira spike for that, and asked a question on slack regarding how come we don't log such important actions - or,

rather, all incoming requests (like Google does). Wrote that opinions on logging vary (Google has a different opinion). Chris wrote that opinions actually don't vary (his one is correct).

So I thanked him for letting me know that there's only one correct opinion. And separately asked Jacob whether sarcasm is acceptable in our company (it was not in the previous one).

Let's see.

Stocks are doing good now. Suddenly. Strangely. Almost unexpectedly. But not all. GLD is going down - but that's as expected.

May 19, 2023

By the end of the day I hardly remember what was going on.

Ok, since I'm on pager duty, I did fix one customer's email - we have an API for that, no clue what's inside, but at least I don't have to go do psql in prod. Good. Wrote to customer support about that. Got no confirmation from them. Weird.

Yesterday... do you remember what was happening yesterday? We have some weird hacks happening, and I suggested to log all requests as Google does, and wrote that opinions vary about logging - and Chris responded that our opinions don't vary. This is cool; I wrote something sarcastic, but not for the dumb ones to figure out. Then asked Jacob if it's okay to use sarcasm, and he said - in small dosage. With which I had wholeheartedly agreed. That was yesterday.

I would not use sarcasm with the Chinese, but the Chinese don't say stupid things, so we can talk seriously. What would Qian say to me? He would say that yes, that would be nice to have, but we are saving on S3 space. Ok. We do save some gigabytes, I know.

Then, since nothing was breaking, I checked that Qian approved my pr, and I merged it, then wrote a pr to deploy it to staging, and deployed it, and then discovered that my new API is invisible. Weird. So I investigated the code generated by Swagger. Ok, I can write a test. I wrote a test. The test passes. Weird. Now what. Anyway, probably need some logging (right).

That was the end of Friday anyway, so I've just closed my work tabs. Will be back on Monday.

May 22, 2023

Still on pager duty, but nothing happens. Cool.

So I was working.

Last week I could not figure out what happened to the API I deployed to staging. Asked Qian. He had no idea.

So today I investigated. And eventually I figured out: it was not deployed.

Started looking into that. Kaboom, a weird error from kubernetes. And that's why? Found the logs eventually. Wow, the app just does not start. NPE.

And then I figured. It was a pretty bad macros: if you specify a type of a value, `s/val x = something/val x: SomeType = something/`, the app does not compile in IntelliJ, but it does compile in sbt (and in sbt in jenkins), but throws an NPE. Але, Женя! Oh, Eugene is now busy with other weird things at Google.

So, I wrote a one-liner PR to fix it, and bumped into Chris' resistance. He started teaching me to do software development. Like, you have to run the service locally first. Oh, really? How? (So far nobody can run this specific service locally.) Or you should run it in a separate container. Oh, really? Do we have any documentation? Karen wrote some, but it's unsearchable, and Karen is in Japan, will be back Friday.

So I'm still negotiating with this guy. OMG how pissed off I'd have been just several years ago. No I don't fucking care, absolutely. The guy also admitted that he would not mind reviewing a 20-file pr with all the functionality included. Nice. Good to hear. I'll remember that.

What else? Nothing. No, the service does not run locally, but it's a good challenge. I'll rather write an integration test that starts the service locally and pings it with different calls.

Fuck, how come I have to listen to all these people all the time? I don't know. Maybe I should retire? Seriously.

Also, had this discussion:

chris 9:58 PM to be honest i think your approach tends to be too academic and not very practical. at the end of the day this isn't an academic exercise, we're trying to build software that meets a business need, so focus more in that direction. if the endpoint you built doesn't do exactly the right thing the first time, we can always change it later, especially for internal endpoints. 9:59 if you don't feel the code you wrote is even worth running before you submit a PR, don't submit a PR yet.

Juan

10:03 PM

An interesting opinion. Thank you for this. I probably should have shared with you my experience in the industry, but anyway, when I hear about academic approach, I really feel good.

†

chris** 10:04 PM**

but the point is *not* to engage in academic exercises here. save that for the university. (edited)

Juan 10:05PM

:)

ī

P.S. Update. Just unlinked from Chris Friend, May 22, 2024. Fuck them, these people.ī

May 23, 2023

So, what happened today.

- Discussed with our partners the API I wrote for them. They want to be more conservative, and instead of “upsert” have a “PUT” and “POST”. Customers’ desires are (when feasible) the law, so I promised them to make it available next week.

Meanwhile, Chris continues his critique. Like, “since we’ve been calling radio stations”streams” like 8 years ago, they should still be called streams, and when you write “getStation(stationId)”, you should instead... etc. I don’t fucking care, I’m just trying to solve the puzzle, what bugs the guy. Will talk to Karen, I guess. She will find a solution, she’s been here for a while.

Anyway, what else did I do? Sounds like nothing.

Oh, no. I wrote an integration test that launches an app, internally. And you know what? The app fails, doing some weird stuff with Mongo, the stuff that it probably should not even be doing. Running index checks on four collections, sequentially, with Await, giving it 10 seconds.

That code was written by Victor in 2019 - before Covid! In the office! He was probably too busy to think.

Anyway, asked Jacob and Karen.

May 24, 2023

Yesterday Renato came up with two jira cases where we replace “upsert” with “update” and “insert” endpoints, the first one via PUT, and the second via POST.

Cool, so I started with the first one, and decided, ok, let me first make the integration test really run the app (and then, in parallel, ping it with requests).

It turned out that I need a sophisticated configuration to run this shit locally, not in kube. But, thank God, the Ancestors wrote a very good exception message, just telling me what I do wrong when running it locally. Oh, I miss those people that left. Anyway, so I did, and the service started running locally, but I still need to tweak some stuff to catch my curl's.

Then we had a meeting, and then another and then another. The last one was just an entertainment, a trivia where we were supposed to answer questions, a mix of smart and stupid ones. Like, how many

languages (scripts actually) have LTR (12). How many countries have shilling as a currency (4). When was Berlin wall destroyed. Chris said: 1964. You wish. Oh, how I remember November of 1989. Winds of change. We were really really lucky to have lived through those times.

BTW, I won that trivia. Well, I'm considered some kind of academic over there. I know too much. Hence, I'm not practical. Sweet summer child, as some wise people write. IRL, I just don't fucking care. Because I know what real science is, and all this shit that we are dealing with is so low, so much below the shining snows of Math mountains, that it's not worth serious discussion, just an amusement. Btw, I guess John de Goes has a similar view of things, he's a Mathematician by education too. As Дуся once said, we are arrogant people from SPb. Well...

That's it for today.

May 25, 2023

Well, now Karen is back online. She's still in Tokio, so the time diff is huge. I guess she works from the bed. But is it my business? At least she approves my code without all that shit. Except that she gave me an idea to add an "experimental" tag, and (as a result?) the service running in staging does not see my requests (while when it runs locally, it does).

Oh, whatever.

We had a 1-1 with Karen; I was really happy to have her back online, with no need to convince Chris about the usefulness of refactorings etc. That PR, by the way, I had just killed it, since it makes no sense now.

Yes, I mentioned Chris' comments about my "academy" etc, but Karen is Chinese, so we can only smile discussing this kind of shit. You may wonder what's my position here? My position is positive. I do appreciate the Chinese approach.

And whatever, it's already 22:04, and I'm still working? No way. What I was watching actually is a show regarding swagger and its tags.

It's okay.

Also, Renato asked me, am I sure all those strangely added jira cases to my spike are going to be a part of that spike. Eventually I looked at them again, and said no, they are pretty much separate spikes, to be split into cases, and they will, well, take something like several weeks, until probably mid-July. Let's see what Renato says. Definitely these new features were never mentioned, it was Karen who added these cases to this spike. Or was it Renato? Hmm.

May 26, 2023

Today is half-day, due to Memorial Day. But anyway, I worked too much.

Experimented with that “url tag”, which Karen told me where to insert; strangely, it works on staging, but not in the local run: it means what? Who is that proxy that makes a decision? Weird.

Anyway, added a little bit more functionality, sent for the next review.

Can do the main functionality now that the interface is implemented. Can even do TDD, because it’s clear now how it should interact.

May 30, 2023

Not much. - found that the PR I posted yesterday is from my fork master - how?! No clue. Have to add some checks in my scripts. Anyway, removed that, and created a meaningful branch, and sent the PR for review again. Since reviews are done by Karen, and Karen is in Tokyo, and our 1-1 with her was when it was 1AM over there, I can’t expect her to review it right away, but do I care? I decided to work on my “lib” project. Some problems there too; discussed it with Karen, we decided I have to minimize it. Definitely.

- had a scrum meeting, bravely reported my progress, made Renato happy. Also discussed with him his run this weekend, and my ride this weekend (we are friends on Strava), and changing bicycle chain, etc.
- Then there was another meeting, where I did not have to say anything, so I was just playing sudoku and listening.
- Then there was a presentation by Jorge regarding a weird bug he investigated, and wrote a plan to fix it exactly by July 5th. Right after fourth of July. Ok. The problem is mostly in naming, because the code calls some id as “entityId”, and what is that entity, opinions vary. It had to be ‘stationId’, but Chris objects, because in his view “stationId” is almost the same as “streamId”, and it is “entityId” if viewed from the AWS p.o.v. I suggested to document all this. (Or wait until Chris leaves the company and unify the terminology, fuck.)
- Then we were supposed to have a 1-1 with Jacob, but he moved it to tomorrow. Ok, tomorrow. Do I mind.

May 31, 2023

So, I have two PRs and nobody is much interested in reviewing them.

We have a “sync-up” meeting today, and Renato suggested to throw in topics to discuss; one of two my topics was changing our policy of posting PRs for review. I tried to put it as gently as possible. Seems like it worked. Sent one to Jorge. Karen looked at another, there’s a pretty interesting issue of rate limiting that I have to check before merging.

Anyway, cool. Meanwhile I was working on making my “library project” to build on Jenkins. So far no success. I stripped it of all the fluff - no way.

Then we had a “QA for API” meeting: our QA from Africa, two managers, me and Jacob. Seems like QA managers are either clueless in general (regarding QA) or clueless regarding what our developers do. But anyway, everyone was happy with what our QA from Africa does. She’s really efficient and smart.

Then we had a 1-1 with Jacob. Discussed my “sarcasm” - I explained to him in details why I wanted to discuss with him the sarcasm that maybe not everyone appreciated. Seems to be ok. Then discussed the “mainline” stuff I’ve been doing lately, and the sidekicks (I have 2.5 of them). Jacob said that because of Facebook’s demands/requirements/audit I should rather focus on Scala version upgrade in one project (I was doing it a month ago, but then put it on the back burner). Ok, sure, will do. Sent a question to Mahesh regarding whether he knows wtf is going on with my library project, and will freeze it for a while. Mahesh’ advice did not work.

Also Jacob told me that I probably have the highest engineer salary in the company. Nice, and I told him that I agreed to go seriously down, because I really liked the interview. So, East Coast, how do people live there?

Just went around the two PRs, but can’t seriously work on them now, have to drive to Toulouse right away.

June 2023

June 01, 2023

Ok, I had convinced Karen and Jorge to review my PRs, and merged them to master (and between themselves).

And then Chris raised his voice again - why t.f. do we need these two methods if we can have one. Then he wanted to see the discussion that I had with our partners. Etc. Then he started arguing with me. And convinced me. We can just do upsert. It's idempotent, by the way. So, we are both happy. I'm simplifying my code. Our partner is okay with all this too. I believe Chris is right.

What else? Nothing. Feeling pretty sleepy now. Hope to send a pr for review within about 10 minutes from now.

June 02, 2023

The pr with all that yesterday's fix was approved, I merged it, then another, so the functionality is there. Karen noticed that tokens are not required anymore; they should be, so I'll fix that. But the API is ready.

So I went ahead to make sure the "free-form json" is also there; it more or less is, but I intentionally blocked it on the web boundary. Will have to add a test case and enable it.

That's it for today.

Oh, also we had a hangout in the evening (2AM for Karen), talked about movies, mostly anime (where I'm totally clueless). Jacob has several adopted kids. I think such people are heroes.

June 05, 2023

Ok, one funny case is when a test suite crashes in the app it launches when there's a test to run (after the app launches). But I did not have time to investigate it, because Karen is coming up with improvement after improvement in the code that I want to deploy to let our partners play with it. What's that in the minds, I don't know. And fuck if I care.

The improvements are: "don't limit the rate in getter API, because it's not required", then "oh, no, let's limit it, we can control it via config file", "oh, change the config file, disable this rate limiting", etc. Wait, we are talking about calls for small chunks of data, for staging, for our partners to try.

But fuck if I care. I'm tired today.

June 07, 2023

Combined flying TLS->LHR->CLT with working. Like, deploying something to staging, sending a PR, etc.

In the middle of the day (about 3:30, just arrived to CLT) Renato asked me on slack, what of my tickets are done, what are almost done (I had something like eight of them). Only late in the evening (no internet at home) I made a list of all those cases, with comments, and wrote to him.

Meanwhile, our Texas partners are happy that the API is available on staging.

Karen seems like forgot what was where, and continued bugging me what I did with rate limiting in the method that was removed. Wrote to her a long explanation, and a bit of philosophy - that we need to have tests to validate the behavior, not guesses. But actually I did make a mistake about 3 weeks ago (it was bad merge to master) - a big chunk of test code just disappeared. I'll restore it, probably this weekend.

Ok, that's it for this long day.

June 08, 2023

Karen, strangely, was up until about 3AM Tokyo time. So we had a 1-1, discussed all those config files, eventually she approved my pr updating all those files, then I merged it, posted a pr to deploy to staging; everything worked; then I posted a pr to deploy to production, and it was approved too.

So, huyak-huyak. Production.

Updated all those jira cases, or, rather, just closed them.

But one thing. I can't login to docker dashboard for prod. Wrote on slack. Swathi said: use this url. Etc. No, it does not work. But by the time I tried all that shit, Swathi ended her work day. Cool. Hello, production! I'm done for the day too. We have pager duty people (my turn is Sunday).

When will I do something productive (which brings fun)? Oh, tomorrow. Mañana.

June 09, 2023

Ok, the jetlag is hitting gradually, so I'm pretty dumb.

Spent half a day trying to extract from Karen and Swathi how t.f. exactly should I get in to kubernetes dashboard. I'm being kicked out etc. Karen wrote that she just updated a document, two days ago, describing how to get added to the list of people that are allowed to go to prod. So I pulled that project, added myself, sent to Swathi for approval. Swathi writes back that this project is archived, and we don't do it this way. I got Karen involved, and basically got the two in a conversation, out of which I had phased out.

But my stuff was not working anyway. How come? Eventually I commented out a line in my script (I'm clueless, I just script what they tell me to call). Not I can get to the dashboard; so I could look into fresh logs (nothing bad there), and I feel ok.

Started the next thing. Eventually Copilot wrote me a one-liner. But later I discovered that I did not need even that (although Copilot's code looks good). But then I added some tests, and found that our ancestors treated JSON in a weird way. They fill optional fields that are None in the db (it's Mongo) as nulls. So I asked both teams, partners and our "titans", and closed the work tabs. Enough for today. Anyway, I have two or three tickets closed by Copilot. All I need to do is write a couple of test cases.

June 12, 2023

Ok, I'm on pager duty, Karen on the plane, and we went for a walk with my wife, and the pager paged.

It turned out that in one of the services (or jobs) a bunch of requests were failing, but all we know in the logs (we had a pretty good workbook with good links) that only request ids are saved. The rest is somewhere in Dynamo, etc. So I created a jira case, and wrote some code in that project, to log the profile id, the functionality requested, and the status - big deal. It went through well, except that Mahesh already created a Jira case covering the same problem - which case was happily ignored.

Anyway, another thing is that support reports users problems with timestamps - in one place they see local time, in another they see EST. Bad, bad service, but which one?

Did not do anything else, because, well, I'm busy, I'm on PD, you know. My favorite pastime. Солдат спит, служба идет.

June 13, 2023

Not on pd today, so I kind of worked. First, the "privacy service" deployment was approved, and I pushed it to staging, together with the job. Everything works so far. Tomorrow morning will check the logs, and request the deployment to prod. Just to be able to trace the failures. Then, also, spent time chatting with Chris regarding the PR where I just had tests prepared to validate the functionality I was going to push next. But Chris was fucking picky, "remove this value that is not used" (the value is in TestBase.scala, and whose fucking business is what it is for).

Actually there are two. One is to check whether we are on a local machine (that's in tests), and another checks if we are in a debugger (then all timeouts become huge, to give people time). Since the only use case is commented out...

Anyway, while I was commenting out these two, on his whim, I, by mistake, had also pushed the next step, the functionality. It's a one-liner, but you won't find it, because there are also more test cases.

Chris got pissed off, and told me to send this to review to someone else. Sure. I have friends (his last name is Friend, but he's not one). Sent it to Karen.

Karen also reminded me how t.f. can I run the server locally and avoid problems with Mongo. I have to go to Mongo dashboard and add my IP to the list of IPs allowed to talk to Mongo. On staging machine. Who fucking cares, right?

Then we had a “standup”, then we had a “planning”, where Renato was trying to convince me to do what MVP needs and stop there (as if I fucking care, my stocks bring me more than I earn at my job anyway) - sure! We’ll get back to it when our users start using it. Big deal.

Also, I got back to my OKR, where I was planning to upgrade one project to Scala 12 (or 13?) and start using ZIO.

We had a short, about 5 minutes, 1-1 with Karen, she is on a jet lag, so why should I waste her time; and since she answered my question anyway... Everything’s fine.

Also, Priyanka in Africa was happy that I reviewed here PR.

I won’t tell here, but IMHO female Indian programmers are best in this category. Absolutely. Next follow the Chinese ones, but for the Chinese gender does not matter. As to the Indians, males are idiots, females are extremely smart.

June 14, 2023

Ok, deployed the diagnostics update for privacy job to prod. Ok. Half an hour later saw those errors properly reported. Posted it on slack.

And then kaboom, Chris comments: “oh, these are just test data, and the errors are legit.” And Karen added: “yes, I have to find time to clean up the privacy database on prod”. И эти люди запрещают ковырять в носу. Chris pretended to be a big proponent of quality and reliability, and now what.

Well, whatever, I did my part. And I’m not on pager duty anymore. If I were, I’d have created a jira case.

Karen is still reviewing my pr, the one that Chris refused to review. She asked me to write a description of the changes. That’s about 6 lines of code in business logic, and about 200-300 lines of tests and test data. Yes, I added the description, it’s important, is not it. Still not approved.

A good news is that we now have Newrelic. Which is way way more reasonable than Kibana. I just took a look, found some errors reported, investigated, cool. If I were on PD, I’d have created a jira case.

Standup, and again I had to explain the status. The status is: our partners are totally happy, and the only thing left is to produce as output, on demand, exactly the same data they send us. We did not, before. We had an idempotent monad, so to say.

Also, played with converting one project to use scala 2.12 (no shit). Turned out that somewhere deep, in the depths of sbt and ivy2, sbt looks for jars in e.g. `~/.ivy2/localcom.typelevel` - neat, right? Asked people around, nobody seems to know. (sigh)

June 15, 2023

Karen did not approve my PR last night, said that left a comment. I could not find the comment.

So in the morning she wrote that she has a question, whether the functionality on GET has changed.

I had to admit that it did. Now I don't return `null` when the value is not defined in Mongo. Her idea was that for our users it can be a surprise. I think not, but OTOH, who knows those frontend people.

Anyway, Karen was happy, approved the PR, I merged it to master, deployed to prod, and decided to test it tomorrow. Too exhausted today.

Started working with converting one project from Scala 2.11 to Scala 2.13. Found it depends on one tiny ancient project of ours, which does not have a 2.13 version (as if it's important), so decided to tweak that one. A small project, but converting it to 2.13 is just a torture. So far no progress.

We have a DLL hell in Scala world. I'd love to switch to Rust, where no such thing is possible. But sbt is totally crazy. Also discovered that when using sbt of version above 1.4, you bump into tons of stupid new errors, hardly documented, it's all about "libraries compatibilities". Hey, Yokota? You are not there anymore... shit...

Oh, btw. During that "we are hiring now" meeting, a couple of engineers insisted on not going, during the interviews, too deep into Scala mysteries, and instead check if the candidate can investigate a server failure. Сельские механизаторы нах.

June 16, 2023

Working kind of randomly today.

The biggest thing is that I found that the "extra data" that I'm supposed to accept disappear at some early stages. See, I slapped together that functionality, but since nobody uses it and everybody wanted it asap, I decided I release it first, and then make sure it works. That's in the schedule. No, it does not. Funny.

Also, we had a meeting today, we are hiring. Pretty nice, right? Wrote to Stephen, who's interested. He was kicked out from my previous company, about two weeks ago. I'd call it systemic racism. An Indian CEO, you know.

That's it for today. But I'll investigate tomorrow, pretty curious, right? You get a JSON from a PUT request, and when you parse the request, kaboom, *some fields are missing*.

June 20, 2023

So, pretty progressive day today. Deployed my shit to prod, so far so good. Except that when I started looking at the logs, I found that, shit, we don't pool mongodb connections. Every call goes with a new connection. This is stupid. Wrote on slack channel, but it seems like it's nobody's business (including me - maybe later).

Had a 1-1 with Karen; agreed that next I'll try to do my best to upgrade our code (all of it) to Scala 2.12 and 2.13, whichever succeeds. So far it sucks. A teeny-weeny project that I'm playing with does not "compile" because "style checker" does not accept build.sbt - yes, the same file that was there for the last 6 years. And I "reformatted" it with IntelliJ, does not help. Most probably it's about space characters here and there. I tried to update style checker config - then sbt complains about missing feature checkers. Omg.

Otherwise, I spent half of the time at the conference; missed the beginning of Jiri Adamek's talk; he's cool anyway. But too much bs talks, in the style of "I'm a category theory scientist, here's my picture of how machine learning works, and it's a category. These people work (do they) in Topos Institute; nobody seems to be aware what a topos is. I feel like next I'll have to write "Topoi for Dummies"... maybe.

That's it for today. Having fun tweaking build.sbt.

June 21, 2023

What's up today? Wednesday. The day of meetings, meetings. Pretty tired. Especially when Ashok 🧑 posted a complaint that his attempt to update station returned HTTP 503. It turned out later that he used HTTP GET to post the json, and that the json was not json, but a piece of javascript (keys without quotes). I only returned to this shit when the meetings ended. And me and Karen



spent some efforts trying to explain to him that the documentation says that he has to use PUT, not GET.

Anyway, and also, since we are on Newrelic, I went ahead to check if my logging fix is ok, and could not. So Karen sent me a neat sample url for Newrelic, and everything went smooth.

Also, since the "epic" was about MVP, I kicked out two paranoid cases from that epic, it's definitely not MVP, it's 100% quality proof, which I love, but it's not MVP ("minimum viable product"). The MVP has

been around for almost a month now, and even Ashok tried it (not very successfully, so far).

Oh, what else: I've cleaned up as much as I can, and now I'm ready to do the scala version transition in the projects. Really, the geniuses from the past made it hard. 1) Nix that knows better what is where (and does not let anybody make any changes) 2) our local tool that handles a uniform solution, e.g. for "style" - that's how my "build.sbt" does not follow the "style" (even the 6-year-old version of that build.sbt, from github, does not).

June 22, 2023

Spent most of the day struggling with sbt and a "stylechecker" written by Eugene's friends. Eventually found how to block this checker from checking the style of plugins.sbt and build.sbt. The rest of the day I spent trying to figure out where sbt finds the scala version which it injects into jar names. "2.11" is never mentioned in the project. Here's what turned out: "If the Scala version is not specified, the version sbt was built against is used. It is recommended to explicitly specify the version of Scala."

Configuring and maintaining build scripts is really not my area; I'd rather do monads. But anyway.

Also, had to slap together a jira case regarding the thing I promised to do: pull all stations from our partners, stuff them into our system, then read them all and make sure that the data are identical. The dataset is not that huge.

That's it for today. Kind of a waste of time, solving sbt puzzles.

June 23, 2023

Mostly I was busy participating in the conference.

Otherwise, ok. Figured out why that small 2.12 project does not build on jenkins, updated a config xml (one of four that there are) for jenkins, sent for review. But it's Friday.

Then what: I have just one case to finish to transition from "MVP" status to "completed" - that is, have a test that takes the whole list of stations, pushes them to the db, and checks that all of them when retrieved in json format are exactly the same.

That was about 10 l.o.c., and it worked as expected. Mission accomplished. Sent a pr. Will close the case as soon as it's merged.

Also, Davis was asking about how better to transition from Int keys to Long, and he introduced a generic HasId trait, and I pointed him at "WithId[T]" trait, written by one of the great ancestors and maintained/used by me. Let's see if he agrees. See, it's all about scary generics in the code.

June 26, 2023

Merged the PR in which my big test checks that everything's ok in my import. Karen added a reasonable notice that the code deployed to import job is a little bit too old, and it makes sense just to synchronize it with the code that was tested, just in case. Sure. So I did.

Also, she found that my pr with an update for jenkins configs actually does not contain anything. What?! I updated a couple of xml files, ran the push script, and? And. It turned out that, since my push script first runs a build and tests, these xml files are successfully overwritten by other scripts. Nice. So I had to tweak some groovy files to reach the effect I needed. Oops. Then we had a discussion with Karen why it's needed.

Also, last week I closed the jira spike since it was about MVP. Renato reopened it, because, as he said, it's him who is supposed to close spikes, after talking with stakeholders. Or, nice, fine. Meaning, I guess, I'm not even responsible for spikes. Ok. Not that there's anything to talk about, because it's all done anyway. Just none of my business. This part I love.

So, that's it for today, basically. Will proceed tomorrow.

June 27, 2023

Was mostly closing the issues.

Ok, the build with that tricky old project eventually does run, so I sent it for review (not much to review, but we have rules, right).

Then, I sent a bunch of other pull requests, and gradually was merging them and/or deploying.

When an import job was deployed, I checked its logs in kubernetes. Turned out, there are systematic errors on one of the records it imports.

So I did import all those records from Texas, found the one that was failing, created a test case... but anyway. The error was that the number of listeners of a radio station was reported as 90.5 - one listener was half-human, I guess. Probably a GPT. Of course when we switch from pull to push, they'll get a proper error report, 400 something (added a test case). Informed the partners.

My main task for this "sprint" is to make some projects build with 2.12 - tons of weird stuff in our build system, but I'm working on it.

Another funny case is a piece of code (I just reviewed it casually), where the same instance, depending on some flag, is called either with map or with flatMap. Since intellij does not understand the build file, it does not know what's the type of that value. And of course no types are explicitly written in the code. Proud hackers. So I wrote to Chris, asking what this can mean.

This is a pretty short week/sprint, thanks to the 4th of July.

June 28, 2023

Spent the day doing almost senseless stuff. Like, you update a groovy script that generates jenkins configs xmls, updating the version of the build of a job. More, updating it for staging is ok, but if you update it for prod (what I wanted to do), jenkins fails in that build, saying that it can't find some unrelated dependencies. So, still investigating this shit. Also, got on Chris' nerves regarding a piece of code that, depending on whether some value is defined, calls a map method or a flatMap method of some value, whose type I could not figure out, because intellij can't "import the project", throwing an exception. Chris, after attempting to convince me that the code is okay (which I may agree, but there are no tests), eventually told me the types of a couple of variables. Meanwhile I fixed the build, by specifying the JDK version, and it stopped crashing, and I could add values types. The code sucks, so I decided to clean it up. It also turns out it uses ScalaZ, and it uses Cody Allen's library for god knows what. Also, friended with Cody on linkedin. I would not say we had good relations with Cody when we worked together at Salesforce; he's a Scala Cats guy, so we are from different camps. I might have unfriended him because of Jon Pretty affair. But Jon is in the past; he's probably got into religion now. Anyway, and there was a "standup", then a "team sync", all rather senseless. Then Qian gave us a talk about Oauth 2.0, where, it seems, all the terminology has changed compared to Oauth 1.0, if I'm not mistaken. Qian explained it all perfectly, in spite of his heavy accent. Then there was a "roles and responsibilities" meeting, with a big boss, who seems to plan to unite our team with our partners. This was the first time in America that I heard good words about myself. Never, since 1998, I heard even one nice word about me. I thought it's just American culture. And this is the first East-coast company that I work at. Now I'm more and more negative about California, Bay Area, West Coast. As Samantha from Google told me once, "it's jungle out there". I did not know it. Was I a masochist? Or just a naive sucker? Anyway, проехали. I worked on. Just a small code cleanup, after all that ancient shit like `val legacyTransactorTask: Task[HikariTransactor[Task]] = transactorFor(primaryConfig)` - that's after I extracted this variable and added the type.

And so it goes.

June 29, 2023

Spent most of the day trying to figure out why `t.f. import` is not working.

Eventually got totally stuck in the dumbest call stack in Mongo library. Obviously these JavaScript pros are clueless about Scala. I'd say, just idiots. Found the problem eventually - it could have been detected at compile time, and if not, in tests. But now. It only crashes when you deploy it and run. Thank god not in production.

Also, coordinated with Chris the things I was trying to fix (make readable) in his project (written by weird New-Yorkers between 2017 and 2020, when Scala was considered an obscure language, similar

to Haskell, that can only be used by the smartest idiots savants.

Imagine a mix of ScalaZ, Cats, and some very custom code written by Cody, a nice but weird guy, working now with Runar. Oh...

I think I made it readable, sent it to Chris, but actually, as Chris wrote to me, refactoring it may make zero sense, because the complexity will go away as soon as transition to the other db is finished.

Well, ok, do I mind? I just sent it to him for discussion, not insisting on anything. I got a feeling I'm his best buddy in the company now.

June 30, 2023

My pr that fixes the import job on staging still was not approved by anybody.

OTOH, I did clean up the oauth service tests (and already introduced caching for some data inside); sent it for review. No rush.

At 2:30 had a weird 1-1 with Jacob. Had a short casual talk about how we spend the holidays, and then I asked him about Stephen - no answer. I asked later again, no answer. He just changed the topic. Maybe the company is not hiring, after all.

OTOH, he sent me a document which either I or him are supposed to fill in later, and gave me a hint about a position of an architect. Not the kind of hint I'm supposed to respond. But I can't tell him I don't care much. Probably for my salary it will be more convenient for them to upgrade the title. Do I care? I don't. It's a nice source of income, but stocks bring me more. And also I really don't fucking care. But I'll be glad to work with this nice company. As I already planned, I was going to retire, get half of the previous salary (that's what I do), and enjoy not only what I do, but also the people around. So all this suits me.

That's it. Since my projects problems are in the limbo, I'll be able to do the stuff I was planning to do, but did not have time. Neat.

July 2023

July 05, 2023

So my fix for the weird import job was eventually approved, and I merged it. But then I had to deploy it (and run), which is done by updating a groovy file, generating configs for Jenkins, and, well, merging it all. So spent some time (actually working on build.sh which would kind of automate this shit).

Also we had a sprint closing meeting, and when asked about whether import will be deployed today, I was saying “I don’t know, it depends on the speed of reviews”, and when asked whether I’m going to add endpoints for updating markets I said “I don’t know, it depends on what we agree upon with Mike, maybe there’s nothing to do at all”. So we agreed to postpone it to the next coming sprint - and everyone was doing the same.

Karen was sick, so we did not have a 1-1, where I wanted to discuss with her, maybe we should not have a cache, if the data are requested once a day (which I don’t 100% believe). But she’s sick.

What else. Oh, id length. Amazon, during the long weekend, increased the default length of user id for Alexa. We store it (wrapped in some stupid stuff) in Postgres, and the column width is not enough to fit it all.

Ok, you can tell me what we should do. I know what we should do. I also know what we should not do: trim the id. Had it in Orkut, when users started mixing their accounts with other users, because “encrypted” ids started overlapping.

Here my question about whether it’s 1 to 1, a mono, etc, was not answered, because, well... you know programmers. They never learned the basics.

As a result, the sprint postmortem was replaced with a rush “war room”, when they tried to find a solution that they would understand. So I had one more hour free for work.

Also, I was exhausted struggling with our jfrog, and asked on slack whether anybody knows who is jfrog’s admin, and I hope it’s not me. Guess what. Jacob was, and he now delegated it to me. I was volunteered.

That’s it for today.

July 06, 2023

Ok, so, I did deploy that import stuff to prod, or at least I think I did (will check tomorrow, it runs once in two hours). There are two updates to be done (and merged) before the thing is actually deployed. What happens if you do just one, I’m clueless (and not that I care).

Had a 1-1 with Karen; discussed this, and discussed the need for caching for 6 minutes the data that is requested (for each account) once a day. No sense to cache. Also discussed the next step, import of

other data, which is supposed to be a part of the data we already import. Agreed to postpone talking to Michael until RJ decides whether it's really important.

So I was spending time upgrading oauth service to 2.12 - it's over 2 weeks already, or 3. The success was that it started building. No need to go ahead and update jfrog (where I still don't have rights).

But what's funny, when sbt started the compilation, it found a bunch of compilation errors. It turned out that one of our libraries, its build for 2.12, just does not have the server functionality. 2.11 does, 2.12 does not. And more, neither compiles. I pulled that library project, tried to build for 2.12 - fails, tried to build it as is, for 2.11 - it fails, does not build. Some shit is bad in mvn repo somewhere out there. We depend on external shit.

Ok, I'll check it. tomorrow. Karen tried to convince me to use this or that specific build, but well, if the code is not there, then what? Let me just fix it all. The code was last time updated about 4 years ago.

That's it. Funny, eh.

July 07, 2023

An interesting day today. Yesterday Karen convinced me not to deploy to prod because this build was not properly checked on staging. Ok, do I mind? Sure. So I rolled back that shit (it was already deployed - approved by Karen, by the way).

Then it turned out I was not the only one to roll it back. Chris did it too, around 10pm. Kind of strange, right?

In the morning, having learned of that, I tried to figure out what's up. I already know that all Chris' bitterness comes not from hatred to others, but from hatred to himself. So I try to be as friendly with him as possible.

One record, when imported from our partners, does not get in. Just not in the database.

And so I started investigating the code.

In short, we run, in parallel, import from five geographic regions. It turns out that many records come from all the regions, under the same ids. Then Edan, the guy that left 3 years ago, merges all these records. Yes, leaving the region in the record to whatever comes first. And merging some fields that are collections. And after that he checks that the rest of the data is the same.

So the key is not unique. The record is unique.

And so on.

In the beginning I was totally amazed how this shit could have been used for so long. But then I thought, wtf, I'll just have to find the people with whom I could discuss it. It's not our team, it's something company-wide, and more, most of the decision makers are now gone.

Left some comments in the code, and thank you guys for that.

So by the end of the day I just managed to slap together a bypass that would make import work.

And why is it? It's because this project is approaching EOL. It'll be gone in half a year, why bother.

Ok, at least I added more checks to the tests, so that this would not get through, at least.

As a result, I spent all day like crazy, running around, pinging Jacob, Karen, Chris, Renato, and why? Just because I could not believe such stupid shit is possible.

July 10, 2023

Surprise, we had an all-day meeting. Starting at 9AM. Teams reporting this and that. Our partners said that our API is almost ready for them to use, 6 access points, 3 left to implement.

Well, except that it has been ready for about a month now, and they know it, and there's just one access point, and whether anything is left, is questionable. Maybe. Will talk to Jacob tomorrow.

Also, Priyanka, a QA engineer from Africa, found a bug in my, well, API; she's a great QA. Fixed it by the end of the day. The trouble is that we receive a JSON with "true"/"false" for one of the fields (only one), but we are supposed to return true/false on request. So there.

Also Karen approved my new caching code (uses Google Cache, with the code that Fry rewrote 100% when he took it over from me and Bob).

Also I'm working on this funny Scala 2.11 to Scala 2.12 transition. Crazy, crazy.

That's it for today.

July 11, 2023

Half of the day there was a meeting, so I was watching it and doing nothing.

Also, merged a couple of approved prs. Ran the fixed import job, but it's too late now to check the results. Tomorrow. Oh, ok, actually I just ran an integration test, and the record is there.

Anyway, spent time trying to just Make Sbt Build for Scala 2.12.

Oh fuck, how much effort it takes. Cross-version builds, etc. But at least that part is there now.

The problem is, in that weird project that I want to adapt to 2.12, the build fails because it reports some file from maven central is broken. It does not tell which file. And all jars downloaded from mvn are correct on my machine. So, a mystery.

So, it sucks.

Also, had a 1-1 with Karen, discussed how we can influence our partners to deal with weird data they are giving us, and how will they be able to do push, overriding their own data. Agreed that I will write a doc first.

That's it for today. I'm almost out of patience.

July 12, 2023

Still getting my small PR reviewed by the two Chinese, Karen and Qian. Qian found that there are two sample json files that my tests stopped using, so I should remove them. And he also found an unused variable in the code. Cute!

Our QA, Priyanka, reported a pretty disgusting bug, where a 500 is returned with the message “non-200 code returned from...” and then an IP address. I tried that IP address, it was also returning 500. Well. That's not how software should work.

There was a “planning meeting”, of the big dept of which our dept is a part. Ok, they talked, we listened (playing sudoku, etc).

Also, wrote a doc regarding wtf is happening with those imports, push or pull, etc. Will discuss it with Karen and Jacob tomorrow, then we will see.

Basically, almost nothing else.

Went to work today. To the office. Which is 22 miles from my home, in Charlotte. Funny, right? I had no clue. Met Bobbi there, they needed my passport, to confirm that I'm legit. Stayed there in the office more than I expected, because her server was not working. So she showed me around. There are four radio stations on our floor. One was broadcasting. Shook hands with the guy that talks on another. And yet another. Saw the server room, with the hardware of the four radio stations. Cute!

That's it for today. Tomorrow I'll work.

July 13, 2023

Ok, it was a crazy day today.

1-1 with Jacob was again moved to tomorrow. NP at all, of course.

Managed to deploy import to prod, and even checked how it runs. I may need to update one more script, that actually does run it, at least for staging - where it was updated already. But it's too late for this week, so, Monday.

Besides, got through with deploying the service on staging, and it works ok, I checked. Notified Priyanka, who's busy anyway, she's flying to India now, from Senegal(?).

But the biggest fun was that this staging config update, I had to do it manually.

See, since I had one update for prod, and another for staging, I had two directories with these branches. That's my practice for years, have two or more copies of the project, to do development in parallel.

And now what. That other directory, the thing did not even compile. Haoyi's ammonite suddenly stopped finding some libraries. When running in a just pulled master. Meanwhile in the first directory, the thing started crashing with "unknown exceptions". And in Jenkins, it started crashing with another exception.

People were saying that it may be thanks to Nix glitches. Nix interferes in the development process in mysterious ways, and should be eliminated.

Later in the day, when I edited the 4 files manually - they were supposed to be updated by the ammonite script, and pushed the update into a branch, and sent it to Karen for review, she noticed that two lines of one file have magically disappeared. Thank you Karen! I looked into the PR, and found that also 206 files were silently deleted by the script. The script that has been working for about two years. Or three. Programmers, eh.

Somehow I managed to slap together another PR, with all sane updated, but still, what's up with the scripts, behaving differently in different tabs in terminal? I know, env returns different lists of env variables, but what's up exactly?

(sigh)

Tomorrow. Production issues have been resolved. And I'll have to talk to Jacob.

Also, Bhagav created a channel in slack, dedicated to the functionality that we are providing them. I am providing. And asked me if it is already available. My notifications from before July 4th were successfully forgotten. OTOH, my initial promise was: mid-July. So that was their plan. And also I sent again the url of the swagger page where one can experiment with that API. So we are good. Except that they'll encounter surprises - described in the doc I wrote yesterday and shared with Karen and Jacob. We will see.

July 14, 2023

Not much today.

Experimented with this transition to Scala 2.12; it turned out that our "playpal" library was totally reworked for 2.12, and as a result, I'll have to make tons of changes in oauth. Which is, I'd say, pretty scary. Of course we have enough tests, but...

I view it as some kind of Revenge of the Nerds. They created weird software, and left. Now what. And what exactly is there? Just a web server, a regular node for microservices. Why not keep it simple? Why not move it to ZIO? Have to talk to Jacob.

Or just port the “playpal” lib from 2.11 to 2.12, and make it work, and not bother about the shit they wrote.

I was supposed to talk with Jacob at 11, but he moved it to 3pm, then connected shortly and told me that they had spilled a gallon of muriatic acid in the pool, so he has to take care. Baking soda will help, I guess. One gallon is not much of a deal for my pool, but he probably has a much smaller one. So there.

Okay, will talk Monday.

July 17, 2023

I planned to talk with Jacob and discuss the options of just going ahead and upgrading playpal to use Scala 2.12 - instead of readapting it to a newer version of it, left to us by the nerds, as a revenge.

But it turned out Jacob will be away for a week (probably still taking care of the gallon of HCl in his pool). But he left his phone number, so I planned to talk to him over the phone - except that he’s in a later timezone, so it would make no sense to call him at his 8AM.

And I started working on bumping scala version for playpal.

And by the end of the day I figured that, fuck, it’s impossible. Typesafe has radically “upgraded” Play, and if you use the new version of that Play, we already have a version of Paypal for that, so I’ll have to make a lot of updates in the server which I was going to upgrade after playpal. Fck, fck, fck. No other options except just updating the server.

So that’s what I will do. No shortcuts.

OTOH, got a lot of experience with sbt. Oh, a lot of experience. Learned something, yes.

July 18, 2023

No success today.

Ok, I figured that updating over a dozen projects to adapt to the new version of scala play would be too much, and we should only upgrade the adapter. The fact that the new scala play allows us to introduce all kinds of configurations is not important, because we already have the code with all the configurations, so what next? Just make an adapter that would work with the new version.

Created a jira case for making this happen, in the next spring.

And I was working on it today.

Also, notified our partners of the questions we have and the problems they will have. Hope they will understand it.

I did want to talk with Jacob, but why bother the guy who is on vacation? We have discussed all this with Karen.

July 19, 2023

Tried to spend some time creating a proper adapter out of the existing pieces. But it turned out I'm on pager duty, so I'm not supposed to do anything useful.

Instead, some client-side guys reported that pagination does not work; they give a limit and get the whole sequence from "collections" service.

The service does not log anything; it also does not compile for me in intellij; also 4 tests fail when I run it in sbt. They don't fail on other people's local machines. Also, there seems to be no tests for pagination, and no way to run the stuff locally, since it uses some other, tests, database. So there.

I'll just create a bug in Jira, and that would be it for today.

Too late anyway.

July 20, 2023

Not my lucky day.

Eventually I managed to run that server on my machine, but... but the server crashes pretty frequently, with obscure exceptions from aws. Well, at least I could add "println" to see that it receives the requests. Ok, it receives GET, but not POST. Maybe it does receive POST, but I don't see any trace of that, and curl does not get any response.

Ok, and of course I tried to do POST on staging, via curl or via swagger; it returns ok there (but not locally, the same request) - but it does not append the data I send.

So, fck, it's a waste of time, but... but I have to do it.

Also, I tried to reproduce the behavior on production, and kaboom, I can't log in (as a regular user), and more, it says I have a wrong password, and when I try to "reset the password", it sends me an email with a wrong url, and when I ask my colleagues where t.f. should I file frontend bugs, nobody knows.

OTOH, at 12:30 had a pretty good meeting with our partners (with the presence of Karen and Renato). It turned out, Bhargav is a pretty bright guy, and he did implement all that push absolutely correctly, so all my caveats are perfectly covered. He also managed to explain to me why the data look so weird. It's not a duplication, they pass the correlation, and it's on our part that the implementation sucked.

I'm less and less impressed by those nerds that slapped together the code, not understanding the meaning of data (I bet nobody knows what "cume" is), missing tests, not providing solutions for proper testing and debugging. Хуяк-хуяк и в продакшен.

Ok, what can we do. As much as we can, right?

July 21, 2023

And today no such aws failures ever happened; the server did not crash nor hung.

So I've repaired my scripts, and, running them, have successfully found the cause of the problem. You have to pass the API version in a header (Accept:...). Sure it takes time to trace.

As a result, I updated my script, documented the findings, closed the jira case, wrote to our customer users, sent the PR with these docs updates for review. Karen is still reviewing it, asking to add this and that to the docs. Waste of time, but she's precious.

Discussed another case with another bunch of reporters; decided to think about it, what is error, what is ok.

Also, Renato told me that suddenly our partners managed to convince the management that what I do is so very important that it should be done asap. Well, ok, do I worry. Ok. But in their docs they want strange things. Like "order the output dataset". How can it be an endpoint?

Well, whatever. I'll investigate. But I'm on pager duty, so I can't work too much on real stuff. I should be on alert.

It's 6pm, it's time to close the work computer. The only reason to keep it open is that Copilot helps me with the book. Not much, but it does.

July 24, 2023

Ok, I'm on pager duty, but nothing is happening, so I decided to do two things. One is to write a proper fix to that shit with content-type, getting rid of acting on the version number in that content-type; the other is to gradually shut down the old import job.

The latter, I looked through the whole codebase and could not find any place where a flag is set for it to run (but it does). Asked Karen, she does not know either.

The former, I spent the whole day trying to make shit run, not return a null from the service. It does return null when I remove that dependency.

4 hours later (or 6) I got it. It's "mock". Mock accepts a different parameter value, and returns a null. In the method that can't return a null, you just look at it. You can check the return value for null. But it's mock, and mock does not even call that method.

So you see how I hate mocks. They break the contract. Silently.

Ok, anyway. Sent the PR. Karen wants me to add a test case that checks that stuff works if you don't specify limit. Of course there's a bunch of such test cases there, I'll just add a comment in the test suite.

Boring, but I agree with Karen. I like this kind of people.

And of course I did not have time for anything else.

July 25, 2023

Ok, last night I was too tired, and just fell asleep.

Funny, Karen asked for a new test case for that functionality fix for pagination; I added the case.

Then Chris discovered my pr. First he noticed that I have a profile id and session id, for a test account, for staging, and informed me that staging is publicly available. What?! WTF? Checked it from my home machine. Yes, it is. WTF. Will have to talk to people, starting with Jacob. Anyway, it's not a problem to remove it from the script (and keep it in my private file), but wtf.

It was not enough, then Chris started attacking my API fix. The defect remains since 2018, when it was last time touched (and the users complained from time to time). And he started asking like "Do we know why it is this way? Maybe if we fix it (return paginated response if page size is provided in the request) something will break? Who knows what is it for?"

I wrote to him some philosophy regarding software development, but then I tried to find Jo Price, who wrote it in 2018. There are about 20 Jo Prices on LinkedIn, so I don't even know whether it's a boy or a girl. Contacted William Ho (first, sent a connection message), and asked whether he knows Jo. Will ask some people internally too. But I don't fucking care about Chris' crazy opinions, and will have to talk to Karen, Jacob, Renato and who knows who else, to make this get through. Carefully, though.

In another development, I could not find how t.f. can I change the frequency of stations import in kubernetes. Asked Swathi, asked Karen. I had five questions regarding where it's specified, and how the job is launched. Swathi does not know, Karen kind of does not know, but she pointed at two files. Eventually I had to learn how kubernetes deals with cron jobs. Okey-dokey, I changed the frequency, Karen approved, I merged it. Still, I have to deploy something somewhere. Will check.

Also, I'll have to write a doc regarding future developments in our pull-to-push transitions. Ok, that's for July 26th (Cuba's "revolution day", storm of Moncada?)

July 26, 2023

So, ok, I did send a friending request on linkedin to William, but so far he did not respond. Ok, whatever.

Meanwhile, I did find a solution for controlling cron job on kube, and successfully made it run once a day. IRL, Raghav etc, our partners, seriously want us to just stop it for good. So I also have a pr for kicking out this part of code. But Karen wants to keep it. The Chinese don't throw out anything, it's a known cultural feature. We have probably about 100 build tasks in Jenkins that did not run for several years. Just try to remove it. Objections!

We have a “test automation” meeting, where Priyanka demoed her big pile of tests. I got a feeling she hired a big team in Africa or in India. She does not even answer my repeated questions why tabs are 8 characters in her scripts. But anyway, if she can do it, cool. It works, really.

Then had a 1-1 with Jacob. First we talked about lake Michigan and huge sand dunes, but then there was this “perf review”. In my view, the purpose of this shit is not to promote anybody, because we are not growing anyway. So he gathered, together with a lot of good stuff, a lot of bad stuff. Like, in one cell it says that I'm the most qualified Scala expert at the company, and in another it's that I by mistake told Qian that his code can crash in certain conditions, and it cannot. How the fuck does he know, right? Do I have to write a test for that? Well, whatever. I know the purpose of it, it's all about salary and promotion.

I don't fucking look for salary raise or promotion, I'm fine with what I get, and my main income is on stock market anyway, which, btw, grows like crazy these days (not all stock, of course, but patience, patience).

What's good in this review is that Jacob failed to properly hide the sources of bullshit. So I know now who's my friend and who's my enemy (as if I care about enemies, I don't have that sense, I just don't see enemies until they assault).

What's interesting, I asked Karen about this fix in “collections service”, but she ignored it. For an obvious reason of never getting into a conflict of opinions. Let idiots keep their idiotship. Well, ok, it's not my company, right? I have some of its stocks, but it's really pathetic. It will grow some time... maybe.

Anyway, no mood for working today. In that table, with review, there's a column for my comments. I copied the table, so I wrote and edited my comments there. Will reread them, and maybe tomorrow, or Friday, will copy-paste them for Jacob to see.

Jacob looked like he's sorry, almost guilty. He's a nice guy. I imagine his face when he decides to fire me (eventually, when the company does a layoff... I was never laid off, by the way... and it's probably too late for me). I will try to tell him politely that I don't fucking care about opinions of people like Chris or even Qian (whom I actually respect). But explaining to him that I don't fucking care about career (never was) may be challenging.

Anyway, all this gives me tons of free time, so we went to the optometrist, then to Lowe's, then to the pharmacy (where they mixed me with my wife, and said they never saw me before... but that's a different story).

July 27, 2023

In the morning I saw that my “run import every day at 6am on prod” was not approved, but there was a comment. The comment was about my comment in the yml file about the meaning of the cron params, and my comment contained a mistake. Qian found that defect.

Eventually I had fixed it, it was approved, and now our job does just that, runs once a day.

Also, since I was contacting Victor about getting in touch with Jo who wrote that weird code, it turned out that Victor remembers all this. Chinese memory. So, the problem is, we have “old clients” who send wrong requests, and that’s it. Since we have over 200 types of client devices... some of them are probably 10 years old... in short, better not touch it. As a result, I restored my PR with instructions, and that’s it. Merged it, no need to deploy anything. Or maybe. Maybe I will, so that swagger docs show my instructions.

Then there was a “standup” (on the couch), where I bravely reported this shit; then Qian was giving a talk about how we are going to do privacy management properly (you know, there’s a fucking state of California, who don’t want us to keep identities of kids under 18 - although all we have is just an email... it’s called CCPR). Well, whatever. I liked the way Qian organized his doc - as a spike in Jira. Will do the same.

Then we had a 1-1 with Jacob, to whom I had already, in the morning, filled the column with “actions”; some of them are just questions - but I swear to God I will never use the word “question” at work, it triggers some weird reaction. So I wrote something like “I need information regarding communicating with client teams”, “I need information regarding projects/directions priorities”. All that he gave me. One of the item (there was a stupid complaint about me deploying to staging an app that was failing on start - this, btw, triggered my adding an AppTest which launches an app locally and does some checks), and another complaint when Qian disagreed with me about an NPE possibly thrown in some piece of code (I’m paranoid about NPE), so I wrote something like “I will never make mistakes again”. This last one gave Jacob a relaxed feeling - that I’m not crazy, that I just laugh at all this.

Actually, in the beginning I asked about the purpose of this doc, and he said it’s about what to tell about the management why I’m being paid so high, and to explain how I don’t need a raise or a promotion. I assured him that I’m not looking for a raise or a promotion, I’m fine the way I am.

In the end we also agreed that I will document the architecture, because it’s still undocumented. Nice, right?

Anyway, after that I needed a break, so I fell asleep; then I figured I don't need a computer to think. I was thinking about how t.f. I can design a proper adapter to use to switch previously written code to play framework 2.6 (which does not use Netty, but uses AkkaHTTP (which is pretty good per se)). I could think while doing some carpenter work, so that's what I did.

July 28, 2023

Woke up at half past three, having figured out one problem and the required solution.

Our codebase, for a bunch of servers, is still in Scala 2.11 - and why? Because it uses Play framework, and Play had changed *a_lot* between Scala 2.11 and Scala 2.12 - you need to rewrite tons of your app code to use the new framework. Assholes, right? So, what did Jo Price (ubiquitous) did? He rewrote *some* of the adapter code to Scala 2.12, but not all. And so tons of server code needs a lot of rework. Now, with dozens of servers, how much time will it work? It's been about 5 years since they started trying. I guess we need several more years to be done.

No Way. We need a proper adapter, so that the production code hardly even notices that something changed in the library. It's not the production code business to react to each version change by rewriting.

And I guess I almost know how.

So I fell asleep, and then most of the day I was just thinking. And experimenting, because updating the existing adapter project codebase, without a ready use case, makes not much sense.

Eventually I figured. I take one project (which I was in the process of upgrading to Scala 2.12), and keep all the adapter changes there. That's what I'm doing now.

Besides, I decided to write a table: project/scala version/play version/planned upgrades. I started collecting this information, and filling the table, when suddenly another google spreadsheet popped up. The one that has exactly the same data, and existed since god knows how long. Mahesh was the last to update it. Asked Mahesh. He said, he just updated it.

So, interesting secrets, right?

Anyway; and another funny thing.

Friday are usually "tidbit days", meaning, you don't work on production code, but spend time fixing this and that, refactoring, improving, cleaning up. And Karen asked me whether I spend my time doing this. I said, yes, and I did some yesterday. And she said: cool, create a jira case for what you did, add it to the current sprint, and report done. We'll get more points.

This is what I appreciate! I was a trade union organizer in the USSR, and I know how to earn points. Well, same on stocks (market was nice today). So, ok, I guess things will go smoother from now on. I can provide them with jira points effortlessly, while doing good stuff.

July 31, 2023

Not much fun today, but at least some progress with those upgrades.

Also, slapped together a one-pager “design doc”, and hope that would be enough. Our partners meanwhile have upgraded their doc a lot. Still asking for an endpoint to reorder records in a database table.

August 2023

August 01, 2023

- eventually the “upgraded” project I’m working on compiles with Scala 2.12 - but not tests (some tests pass, though)
- still working on making sure the remaining tests compile

Raghav asked to add bulk push import - could have guessed earlier, right? Actually, we can use the same api... maybe

Also, Karen suggested that I add points for the additional things (bulk import), so I did, 10 points here, 11 there, with details breakdown.

Still the idea of “reorder the table” API is a mystery. Karen got some thoughts about what they mean, but that’s doubtful.

Well, whatever.

We have a new contractor, Howard. Pretty silent. Nobody asked him to tell about himself, and nobody introduced themselves either.

No, it’s not boring. It’s a fight against Scala compiler and the libraries (some mysterious stuff from Bill).

August 02, 2023

Had three meetings today. And worked on the two pieces that I’m working on. One is - a ‘design doc’ regarding the import shit I’m going to do. Before that discussed it with Raghav, we agreed on an obvious thing: new endpoints will be accepting bulk data, the old endpoint, I’ll try to make it flexible, but if now, I’ll create one more. No changes in API! Ever!

Also, someone from the client side complained about HTTP 500 being received when he sends a station name instead of station id. He figured that on one hand, it’s his mistake, but I told him 500 is not the right code; and I tried another trick, and it was returning 404 instead of 400.

So I created a Jira case, for the second half of August. No rush, right?

Then there was a “planning meeting”, after which I was totally confused. Is it so that I’m supposed to do what I just want to do? (That’s what I do.) I’ll have to create a couple of Jira case (ok, not a couple, but five) for the future dev, but so far this upgrade to Scala 2.12 is what somehow is viewed as important. Ok, it is. It dragged for over 5 years.

The next meeting was about sprint postmortem. And strangely they asked me to give details about this Scala Play shit that made our services stay so behind. And I explained to them what t.f. was going on, and how I want to eliminate this kind of blockers, and do it pretty soon.

Another complaint of mine was about random meetings. Switching the context is kind of hard. So Jacob suggested to have meetings back-to-back, and leave developers tons of continuous time to do the real work.

After all that I was kind of exhausted, but after the pool I eventually got back to normal, and managed to have all that code (the transitioning one) compile, and most of the tests to pass. Except for three tests, which I'll fix tomorrow.

BTW, the PR is about 50 files. Before sending it for review, I'll ask the people, do they prefer me to split it into 4 chunks, or not. Have a vote.

Anyway, I happily closed my work computer after all that. Mission not exactly accomplished, but remember G.W.Bush? (A nice guy, by the way, was not he?)

August 03, 2023

Basically, just worked all day. Managing the play framework configuration problems. Kind of figured out some stuff, so what's left is just so-called "guice", created by two crazy guys, Bob and Kevin. I don't know who this stupid shit became popular everywhere. Scala people, at Typesafe, don't they know better?

In short, fixing this crazy "dependency injection" is left for tomorrow. Other problems are resolved in the code. Well, but next is launching the app, creating an integration test for the app, and making sure this shit works. And then minimizing changes in the production code.

I looked into other projects using Scala 2.12 and Play 2.6 - seems like the authors were half-clueless regarding all those configs, modes, etc.

August 04, 2023

Friday, communication at work is pretty rare, everybody seems to be busy or just away. I spent the whole day trying to make those three failing tests work. The problem was with "dependency injection", thanks to Crazy Bob and (maybe not so crazy) Kevin - their Guice, I never wanted to deal with it. Today I got into it, and, basically, decided to stay away.

Yes, you can add annotations to constructors. Then you'll need to add annotations to the data types used there. Etc. When I approached about 60 classes that need annotations (that includes `Set[String]`, `DateTime`, etc, I decided that it's enough; rolled back all that shit.

Now the question is, why does Swagger generate Router class with these annotations? Also, why it has no bad impact on the tests in another project.

The last question is probably crucial. Maybe Davis just does not test that part of the code.

So there.

Also, there was a “happy hour” meeting at 2:30; I joined, hung out there for about 10 minutes, and disconnected. Boring.

August 07, 2023

Basically, spent all day working on that code.

As I planned, kicked out all that shit that depends on Bob and Kevin’s Guice. Cool, so we are good. That test does not crash, except for one test case, and others fail just because, well, it’s a prototype. I also have to figure out what exactly was Jo testing here. It’s all like “here’s a mock that returns this instance”, call that mock, check the instance, then check that mock was called.

Guys, mocks only exist because the parts are not flexibly connected. It’s all about DI, and DI won’t be done correctly unless you learn lambda calculus, learn it well.

Anyway, whatever.

Had a 1-1 with Jacob, explained to him wtf I’m doing (like instead of two years to spend on upgrades make it about 3-4 months); also Jacob said that my performance has improved after we reviewed it. Sure! I don’t ask any questions anymore. Just no questions.

Also, changed the pull frequency for stations from once a day to once a week. Deployed. Also discovered that one of the “import jobs” is run not as a cron job in kubernetes, but as a job on jenkins. WTF? Who knows, it’s been a long time. But good to know.

Then there was a tech talk, and a pretty good one. I asked the guy to post a link to his slides, he said he did, but he did not (yet). Whatever, I have his name. And some snapshots of his talk.

In the evening I took a break, did some French, went swimming, and now I’m sitting on my porch enjoying the thunderstorm. Now I’m done with that, and will continue the refactorings a little bit.

August 08, 2023

Had three meetings today: standup, planning, and 1-1 with Karen.

At those meeting I’ve been asking from time to time whether they want me to do some production code, but no, I should do Scala upgrade.

Я время от времени спрашиваю, не хотите ли, чтоб я что в продакшен понаписал - но нет, занимайся апгрейдом версии скалы. А и то, я тут самый такой скальщик; и этот апгрейд в сервисах, от 2.11 до 2.12, пытаются сделать с 2017-го года начиная. Ну и вот, а на нас уже насчет секьюрити наезжают, что у нас древние версии библиотек. Так что работаем.

С Карен обсудили такую телегу. У меня в тестах есть один тестовый объект, инстанс класса Session. Если я тест гоняю в интеллиджее, то один из параметров, дата создания - DateTime, а если гоняю в терминале в sbt, то это должно быть long. В смысле, иначе не компилируется.

Вот я ей и показал этот смешной фокус. Ну и когда показывал, когда показывал сорсы по линкам в интеллиджее, то я и понял, что у нас есть два таких класса (на самом деле девять одноименных) - с тем же путем, в смысле, the same package. Но в разных библиотеках. Нет, не third party,

In short, in that test fixture, I just use reflection to pass the appropriate value type to that constructor. Stupid? Well... we'll have to fix our libraries. The rest of the day I was experimenting. Took me some time when I went in the wrong direction, using MockWS, which basically serves nothing: you tell it what to return, and then you check what it returned. No reason.

So I'll be launching the app, the server, and just talk to that server. Working on it now.

When you see a 7-year-old comment about "these lines are here temporarily, until we upgrade from Scala 2.11 to Scala 2.12..." you see, something should be fixed.

That's it for today.

August 09, 2023

Spent the whole day fucking with the tests.

Figured one thing: a bunch of tests are not tests, but just mock tests. They mock some production code to return requested data, and call that mock from the tests. No way.

So I started slapping together pieces of code, and *almost* compiled one test suite. Basically, I'm kicking out guice and mockito. Mockito is still present in some of the code, but not much.

Crazy, now I have 64 files in that PR. That's a WIP, of course.

Priyanka gave a presentation of here integration tests, and showed what to click to run them in your branch (is not it cool?); and kaboom, wrote a nice one-pager doc about it an hour later.

Mahesh talked about Linkerd, but that's totally not my area.

That's all folks.

Just finishes working, it's half past eight already.

August 10, 2023

Spent the whole day introducing sanity in the tests. Refactoring fixtures, separating mock-dependend tests and play dependent tests (in separate fixtures).

By the end of the day one test suite started working. Funny, Mockito just was not working for me there. And what's the point, on every test case populate the mock with all the required reactions. Why not just program the mocks, if you have good decomposition into traits (interfaces). That part is already ok, our predecessors did it right.

So that's it for today.

Also, Samantha posed on slack 9 candidates resumes, and I looked through them - good, medium, bad. One guy is a TA in LA, right after graduating. Weird, right? I would have skipped him, but he has a nice list of projects he worked in.

OTOH, none of these people has a github link on his resume. Which is... cheap.

August 11, 2023

Spent all day struggling with play framework and its configuration.

Not that I've finished it, but there's a progress in configuring the app. Pretty boring, all this; people in typesafe are kind of, a little bit, retard. Or they like the Java DI solutions, and Bob-Kevin shit, Guice. So I have to do my own DI (and there's none actually, it's all a bs, when, instead of instantiating something they want to do "Inject", all the way to the bottom of the dependency graph. Oh, whatever. I can mock and I can DI, without this bs.

Also, early in the morning there was an emergency; Karen was on pager, so she had to wake up before her 6am. What happened. There's a radio station, KFAC or something. It's pronounced s "kay-fack". But Alexa pronounces it as "kfack". That's bad. So Safa updated it, adding a class that contains pronunciations (some), and sending it to Alexa to properly say it. That's it. As a result (of probably something else) some shit started pummeling our servers, and our servers just could not process all they are getting, and we had a stream of exceptions like "server did not respond". There was a lot of chat on this (and who knows, this could have been a DDOS that coincided with this small update, which, btw, was pretty much test-covered).

Life.

No, it was fixed before noon.

That's it for today.

August 14, 2023

Not much new, still working on that "upgrade". Yes, got rid of Mockito as much as possible, but the remaining parts, I just found a bug in this fucking Mockito. Tried different versions, eventually found the one that works.

But then this Mockito/Play, for some reason it wanted a couple of “implicit” values, one being a “web server” - whereby you only specify mappings, request->response.

Yes, I hated Scala implicits since 2008. Today I spent a couple of hours before I figure out that this implicit is injected in a couple of different places, during the call chain, and the one I (or rather previous people) defined, is being ignored, and replaced by another default.

So stupid. As a result, I just created a method that is parameterized explicitly. Now it works. I mean, it passes the request to the “server”; next is to populate that mock server with the right data (and what’s the point? FIIK, but the old tests should still work. So there.

So there.

At least I’m getting more and more in control of this old shit, written about 7 years ago, and a couple of generations of engineers that come and go tried to upgrade it. Well, at least I don’t depend on Nix here. Another fucking crazy environment for hackers that come and go.

August 15, 2023

Ok, spent a day trying to figure out wtf is with routing in the mocked tests.

Eventually, after 8pm, came to a conclusion that I can just mock one last thing that was not mocked, and that will be it. But will do it tomorrow.

Two other test suites will need more attention.

Fck.

Nothing else there, except that I skipped two (non-obligatory) meetings.

I guess normal people don’t use Scala Play, neither version.

August 16, 2023

Spent the whole day restoring proper mocks. Seems like I’ll also have to mock Guice Injections, or else things won’t build. That’s tomorrow.

Renato is back from the Philippines. Returned sick, with “flew virus A”, his wife and daughter are sick too. We had a planning meeting. I’m still amazed that they still allow me to go fix all this shit I’m working on. Well, I know that previous people spent a couple of years trying to upgrade this shit. Moral: Never Use Scala Play framework, written by the idiots at Typesafe who managed to get rid of Paul Phillips etc. Typesafe turned into shit factory.

Anyway, I’m in the process. No big deal to bypass Bob and Kevin’s injectors.

We did have three meetings today.

Oh, also, I signed up to upgrade OAuth service to use longs for profile ids. See how badly it was all designed, right?

August 17, 2023

Ok, spent all day struggling with Scala Play and Mockito and Guice.

Getting closer to reality. At least the tests start talking to localhost, which, strangely, is not answering, but at this level no big deal, really, will investigate and fix tomorrow. At least dependencies are resolved. Setting up proper port numbers must not be a big deal.

There's another PR, a bug reported by Priyanka, our African QA. She's great. Found that the same server does not accept long profile ids.

August 18, 2023

Spent some time investigating which port is being taken in the tests, and why. But then got tired of that shit, and took a break, and took another ticket. Filed by our African QA - long profile ids not accepted by the service.

Started investigating. Ha, she even wrote that our db table declares the field as INT4. Is not she cool?

So I ALTERed TABLE on staging - and then updated a bunch of tests to work with a long sample id.

But then, strangely, the tests were failing - until I figured that I also had to update the DB schema in the project, because the tests create test db every launch. Omg.

Anyway, it was a pleasure to deal with a sane code.

That's it for today. Will send the PR for review Monday.

Also, reviewed Qian's code; his code is really beautiful; a little bit verbose, but it's okay. Pretty well detailed.

August 21, 2023

Ok, spent most of the day at meetings and trying to figure out how t.f. I can log in to prod db in the oauth service. Eventually, by the end of the day, found out that the use id should be not amp_team, but amp-team. I would have discovered it later if I did not successfully log in to a bunch of other databases, where it is with the underscore.

Anyway, I'm thankful to Karen and Swathi for their patience. (Do you notice, I really prefer working with females? They are humans.)

Standup was okay, as usual. Then I interviewed a brilliant guy from LA. Seems like he's a Bengali; it's probably the first time I talked to a Bengali. With a Jewish-looking last name. Talent! Learned a new word, *concolic execution*. I gave him this funny problem:

```
scala> 105.3256 - 105.32
val res0: Double = 0.00560000000000116
```

He solved it without a second of hesitation. He wrote, in Scala, a solution that does it correctly. Within about 20 seconds.

So, I was pretty happy, expecting him to join.

Then I went to Workday, to write my response, and also saw that there's another document I have to get familiar with. And (to me) it looks like a hint that I'm on a PIP. WTF. No, I'll survive. I won't be looking for the next job, even now that I'm losing money every day (made one pretty stupid mistake, but I'm hopeful).

So, in an hour, I had a 1-1 with Jacob. By that time I decided that I won't talk to him about PIP. If it's official, he will tell me. If not, who fucking cares. He was nervous about this "upgrade" that I've been working on for god knows how long. Yes, I'm working on it. Discovering all kinds of crazy, disorganized things. He also wants me to document the architecture. Yes, I have to. Let me start doing it tomorrow.

Well, whatever.

Then I closed the case with that "long id" (having fixed the database table), merged my PR with a bunch of tests for that, and closed my work computer.

That's it for today, folks.

August 22, 2023

No fun. Spent the whole day trying to run the app in the tests. Basically, the services inside that app. Did not work (yet). For some stupid reason replaced one WS client variant with another (both are parts of play, and I just got distracted). Of course changing client does not help the server to start.

In other projects I was just launching the app by myself - maybe that's what I should be doing here, I don't know. Investigating weird behavior of crazy other people's code takes several times more time than just writing your own, simple code.

Well, whatever.

Actually, most of the day were meetings. And an interview.

A guy from North Carolina, first education in psychology, then he had a head trauma, could not continue, so he picked up another profession, took courses in a college; got totally hooked on programming, went to a startup, the startup was acquired by IBM, so he spent some time in Watson Hills, where

a French team imposed Scala, and he read the Red Book, FP in Scala (pretty hard reading), meanwhile he also played with Raspberry Pi, implemented some radio stations streams system; published an article about this; his code is pretty good (saw it on github), he also showed me his code on intellij; asking him to write something stupid would be just an insult, and we chatted about this and that - e.g. how did he learn unit tests? What's his understanding of monads. In short, a big talent.

That's it for today. I love talented people, a lot.

August 23, 2023

Half of the time spent on meetings. SNYK guy was talking about all kinds of threats from libraries pulled from external sources... Like, imagine, you refer, from your web page, a javascript deployed elsewhere. I would not. And the story how TransUnion was hacked. And how an .md file can be used to hack. And the like.

Then Priyanka was giving another presentation of her amazing collection of tests (these things make our product really reliable); the management showed total ignorance, asking us about whether we have unit tests, and are we ready to write integration tests. Well, whatever.

Meanwhile I was struggling with the server I'm trying to "upgrade" - trying to figure out how come the service (represented as an app) does not start; when it starts, how come it creates dozens of postgres connections (should create none), etc.

Am I wasting time? I don't know; at times I feel like I should just retire. OTOH, this project I'm working on, is the most complicated. Comparing to what Davis, supposedly, upgraded, I see that Davis does not test any HTTP communication. Zilch. Nada. Etc.

Well, whatever. No, I'm not exhausted. I'm just humbled.

August 24, 2023

It was pretty pathetic and hopeless. Yes, I do instantiate the application, including the test version (with all the functionality behind the web mocked properly); but I cannot launch it! With Scala Play, there's no main class or method, so you have to figure out. The internets don't help, it's either an ad or some ancient stuff about how was Scala Play 10 years ago - and some pathetic losers sharing their negative experience. Negative experience is good, we need to know that this shit is not good. But well, I can't switch to ZIO HTTP. I have to do what I have to do. Investigated this and that shit; Karen suggested to compare with another service (last couple of years maintained by Davis, who has been sick for the last week); and I did find some suddenly required config files. Does not help. sbt does not find "main class". So I eventually bothered Davis. And he suggested to copy-paste as much as possible from the project that he did. Ok. About 6pm it worked. Had to add another plugin. Did

anybody document it? No way. It's a hearsay, your honor. So, around 6pm I had the app launching, listening on port 5555; good. The tests don't pass yet, but it's not a fucking problem, because I'm in control now. While at 11AM, at the meeting, I just told that I do have some progress, but no good news. Now I have good news. Also, had a meeting regarding the two candidates we interviewed today. I'm the most enthusiastic about both. The Bengali one, and the North Carolina's psychologist. It's not my business to hire, but I liked both. Also, James D., as I suggested to him (ok, I advertised our company a lot), filed an application. So I added some good words about him. It would be really a great pleasure to work with him. We have known each other for over 10 years.

August 25, 2023

Yet another stupid day.

Now that I know how to launch an app in Play framework, I thought, well, my test will launch the "fakeApplication", the one with mocks. No way. The application obviously does not start. In principle, the app should consist of a server, but in this stupid framework "server" is a member of the app. What else is there, I've no clue. Components. And the server does not start.

So I started looking at other projects having the same scala version. One after another, none of them was launching the server. It was just macking receiving the requests. Could as well call the method from the test.

Eventually found an app where the server launches! (Port 19002 is opened, according to lsof). But where is it? Play framework does not call methods. It sends messages. Via Akka. So Thread.dumpStack does not help. Somebody called someone, that's it.

And then suddenly I found, in "scalatestplay" a class that launches the server. Creates it and launches. Takes the data from the application. Why does not it do it in my project? Because I don't inherit that trait. If you do, it has an `override final def run(...)` that is called when `run` in the main test class is called, and it is called in `BeforeAndAfterAll` trait.

Ok. big deal, I'll inherit this trait.

That's where the hell rose. Because I was using (obviously) a wrong version of Scala Play. The one in that project did have the sever and launched it. The one I have in mine - does not. The difference is: one is version 2.6.3, and another is 2.6.5. Mine is 2.6.5.

So I spent the rest of the day rolling back all those stupid dependencies to the same versions as M-
aresh and Qian use. These two guys write not only nicely-looking code, but also nice-looking build
scripts. I appreciate it a lot.

But I ended the day when the things was almost starting. Except that some tool could not find a
method in class Object.

That's when I decided to delay it till Monday. I already had some shit with Scala compiler (it was crashing, over a week ago), with tests not compiling in intellij (so I had to run it in sbt only). Now what, now some troubles with Java Native.

I bet when you read all this, you will decide never to use Scala. I'd suggest never to use Play. Ok, and I'm becoming more pessimistic about ScalaTest. I know, Bill (the author of ScalaTest) has problems at home, his father is very sick, but could not it be left alone for a while? Or is it that his two crazy women, twins, pro-Palestinian, are destroying ScalaTest? Who knows.

So, the only reason I do all this is because, well, I don't have anything else to do. I do make some hints to the management, asking, maybe we have something more important? But it seems like no, nothing there. So I proceed.

August 28, 2023

Well, not much happened today. Except that by the end of the day tests stopped crashing. Basically, some of them are still failing, but at least I found the cause of crash. The version of session service for Scala 2.12 has a different fucking API. So when we mock it, we mock the old version, while the code talks to a new one. So mock fails.

Jacob wants me to document the architecture, and I still promise and promise, but I work all day on that fucking project. Well, I'm getting close. Things are getting smoother. But there are 73 files in that PR. I'll split it when I'm ready. No big deal.

That, and also we had a funny brainstorming session regarding how can we provide our client-side partners with a collection of test accounts, so that their tests will not overlap on the same account, causing confusion. The brainstorm was mostly silent (yes, it's not the Bay Area, where every idiot is screaming loudly their nonsense; I was the only one Bay-Arian, and so I was coming up with suggestions, and the reaction was - silence. Eventually I was offered to write it all down in the Jira case that our partners created - and that's what I did.

August 29, 2023

Well, all I'm focusing on is this project, to make all the tests properly configured and pass. Some pass, some fail. Some are still weird. Had a 1-1 with Karen, but I was waiting for a furniture delivery, so it was just about several minutes, exchanged the opinions, and that's it.

Jorge is leaving. That's sad.

I finished all that shit around 7:40 pm, because it was getting too dark, and I had to go swimming.

That's it for today. Seems like the scope is bigger than I estimated. The service, I'd say, is just disorganized. Feature after feature in the same "oauth" service. Something wrong. And yesterday Qian

already said that nobody knows what exactly this service does. Ok, we generate tokens (of more than one kind), we generate sessions.

Well, whatever. First, have to make it work.

August 30, 2023

One thing I don't understand, how they are so patient. I've spent more than a month fucking with this project. Still not done. Today I found one funny thing: Mockito's mocks are static - while my test (as a choice of the appropriate Play Test framework) is parallel. So one test mocks something, another does - the same - and then verify says that I called a method thrice, and I was supposed to do it once.

First, whose fucking business is how many times a library called this method inside that library. Second, what exactly do we check here? Nothing. So I kicked out all verify in one test suite, and only three test cases fail now. And I spent most of the day trying to figure out wtf. What I found. My client code sends an http request. I spent time mostly tracing it to make sure that the data I provide in the body is *actually* the data sent over the channel on that port I give (19001).

Ok, it turned out yes.

But it also turned out that "by default" I used the "request parser" that ignores the body contents. So that's on the server side. That's the edge I actually had to mock... probably.

Anyway, as soon as I made the parser behave as what is normal in production code, the clock struck 7pm, and it was time for me to go swim and then have dinner. So here I am now.

Another problem was mocking profile service. I had a manual mock, but then Safa added some code... well, whatever. Now I have two "mocked profile services".

Does not matter, I'm on the right way.

Will write another blog, then do French, and then try to finish two remaining test cases. Now that I know.

Renato scheduled the next sprint, and, since my stuff is obviously not finishing these two weeks, and since I'll be on pager duty... short, I'm free to be done when it's done.

Also, both guys that we interviewed recently are going to be hired. Good. Probably they skipped James - he's really too good for our company. As was Prasad, who applied half a year ago.

August 31, 2023

Ok, as I found yesterday, Mockito does not seem to care about tests running in parallel, and keeps mock the cache not in thread locals, but in static, global. So it gets confused when we call the same

method in parallel from distinct test cases. I fixed it by kicking out the `verify` calls, since they make no sense. Today another surprise. When I run a test suite, it works ok. I run another, it works ok. Then I run them all in `sbt`, and they crash. What happens? They run in parallel, and the servers are fighting for the same port. Not a problem, right? I actually find a piece of code, 5-year-old, in another project, that just assigns a different port for a specific test suite. So I announce the port for the current suite right away at the top. I won't compile if I don't.

So ok, except for one test, proxy something, all tests pass now.

But when I decided to push it again, my build script started complaining about "scala style". First a couple of "style configuration" files are miraculously missing. Actually, none of these files exist in any project that I have in my `projects` folder. That is, about 20 or 30 of them. What is it? FIIK. Had to google it; it turned out such files can be auto-generated. So I did.

Now my `build.sbt` and `plugins.sbt` are found non-kosher. Fuck, A file with about half a dozen plugins listed not good enough, stylistically. What kind of Jane Austen should we hire to write `plugin.sbt`? So I had to google it and find how I can ignore a file. It's called "excludeFilters", but the form is the same as `.gitignore` and others.

So by the end of the day I felt pretty happy that I have just one test suite that sucks.

But not actually. Because when I woke up after my nap after lunch, it turned out that the team was looking for me, because? Because the server I deployed in June or July is suddenly causing 100% errors in another server which talks to it. WTF, right? They hurried to roll back that deployment (Chris, as always, demonstrated his eagerness to fix other people's mistakes.)

But WTF? I wrote to them I'll deploy the latest version (which has been running on staging for the last month too, and tested by Priyanka's suites), but, as Karen noted, it was too late on the East Coast (it was 5:30 PM); right, we don't want to bother Davis, who is on pager duty, and who did not, I guess, exactly get rid of his heavy covid case that was torturing him for the last two weeks. So we agreed that I slap together a pr for redeployment, but will do it Tuesday. Because Monday is some special day after which one cannot wear white socks (see "Serial Mom" movie, where Patty Hearst did).

And besides, Karen explained the problem. Today she deployed that bad version 216 - only today; while there was already version 218 available. And that's how all this started happening today. Nobody knows whether (and why) the newer version was or was not deployed earlier. But well, whatever. Chris, I believe, is eager to become a guardian of the builds, but nobody will let him do it. Oh, whatever. At least we are ok with Karen. And with Davis. And (maybe) with Jacob.

That's all, folks!

Oh, and what was the problem? The problem was that our partners send us json files where one of the boolean fields has values in quotes: "true" or "false". Other fields are perfectly ok. When I was

introducing a free-form data storage, we got those booleans in quotes, and started giving them away on demand, in quotes. The fix was to replace booleans in quotes specifically for this field.

Oh, whatever.

September 2023

September 01, 2023

Not much today. I wrote to Bhargav about this stupid bad boolean data that is sent in quotes; he immediately created a Jira case, assigning it to himself, to fix it. Communication, fucking communication. Normal people will improve things when you just talk to them. Ok... As a result, when he fixed it, I'll have to kill the weekly pull job, which misbehaves anyway.

Besides, I have just one test suite left to fix. Before going ahead and investigating, I simplified the test code (test code readability, that's something that almost nobody cares about). Now cases look like this:

```
val response: Future[(something)] = request.post(or something)  
  
from(response) expect OK withJson(expectedJson)
```

That replaces boring repetitions.

These thinks I was writing while waiting for my daughter in a restaurant in Hickory. And then got a slack post from Karen telling everybody that today's work day ends at 1pm, and see you Tuesday. Cool. I'm still interested in investigating at least one failure (getting 500 - meaning the request did not get properly passed to the right service).

That's it for today.

Oh, also reviewed three resumes. One was especially ridiculous: the guy claims to be a "senior engineer", versatile in Java, Spring, Python, Scala, JavaScript (do you expect HTML to come next?) - and in his list of positions he was always "leading" something, and "helping" beginners - right from his first position after graduation. Well, ok, that was kind of my life story: I started "leading" about half a year after starting my first job - but that was because I was a rare case, an experienced programmer right after graduation. I was единственный такой хлопчик в Енакиевому. Anyway, it turned out we already hired two engineers, and that will be enough for now. Pity we missed James. Really pity. James would beat all of them, and probably all of our engineers (including me of course) if joined.

September 03, 2023

I know it's Sunday night. Just tried to satisfy my curiosity: what's wrong with running this proxy server if the code a) was not changed, and b) looks pretty neat (ok, I'd repair the structure, but later).

Wound up with a bigger puzzle. It fails in run mode, and passes in debug mode in IntelliJ.

If you think it's about delays and servers not being started properly before tests start - no! I did add a delay. No. The error reported is something like "and instance of this class cannot be cast to that class".

Puzzle,,eh? I even know where it happens, which line.

See, every time I report shit like this, somewhere a young programmer decides to stay as far from Scala as he or she can.

September 05, 2023

Was spending time with Mockito, like crazy. The problem is, if you need to mock a method that has an additional list of implicit parameters, this does not work every time. Plus, what's the implicit value? Play for Scala 2.12 uses Akka HTTP, and all is done via Akka actors, and ExecutionContext is something from Akka, not the standard 'global', which is available in the test. So one needs to match the parameters, using any [ExecutionContext] as a matcher, and it does not work every time. I was debugging through that crazy Mockito code, until I figured that I don't need all that.

And now I know why Mockito. People use it if they can't refactor their code into manageable chunks. A trait with 50 methods, a method with 20 parameters, and the like. These people think they are programmers. They are hackers. Their brains don't serve them properly.

So what I did. 1. Extracted a trait I need. Just one method in that trait is enough for me. 2. The class that implements does not need to get an implicit ExecutionContext in every call. It actually gets it in a constructor. So it was a cargo cult, passing the same context twice, in the constructor and in every method. Kicked it out from those calls, from those signatures. All unit tests immediately started passing.

Now I have integration tests left. There's some problem of dependencies (mockito again); maybe I should just kick it out of integration tests. Will do it tomorrow.

Tomorrow also I'm on Pager Duty. Chris left me his duty report. One of the services leaks memory. That's what I'll spend time hunting and fixing. I love this kind of shit.

Of course we had some meetings. The shortest ones are with Karen, we have a perfect rapport. Several minutes, done.

September 06, 2023

Ok, I'm on pager duty today. Almost nothing happened, except for one server that is eating up memory, and I'll have to investigate it. That's one case. Another is, have to add something to something in the database. That's in my plans... for tomorrow.

Today I also promised Karen to deploy to prod the version of the OAuth server that has some tests updated (to check for long ids to work). IRL, no production code was touched, because it works, but I did add good test cases. So, just to make sure we have everything in sync, I had to deploy.

I almost did. Will do it in the morning. At least I merged it.

Also, all my tests pass. A kind of a victory. Next I'll have to make the app run (and work).

Had a 1-1 with Jacob, discussed future development (generally, just kicking out Scala Play and using something more advanced and maintained, e.g. ZIO HTTP). It's important to separate business logic from the library tricks. We do have microservices, but each service is a mixture of Scala Play, Cats, Jsons, and the business logic per se. Why? Jo P. to blame. It was all in 2016-2019. People were happy being themselves.

Oh, whatever. I'm at least glad to fix it a little bit.

The last trouble was the “style checker”. Eugene B. probably forgot all that shit. I spent a bunch of time making build.sbt be accepted by the “style checker”. Then I had about 40 files not accepted. So I had to read some docs, and eventually found a “flag” in sbt which makes this style checker fix the source.

Now I have 97 files in the pull request.

Ok, whatever. But the build in Jenkins still fails. Why? Because scoverage does not find its logs in scala-2.11/scoverage. Why 2.11? Because its a parameter somewhere. So I updated a gradle file, replacing find scala-{{scalaVersion}}/scoverage... with find scala-*/*scoverage. Yes, it's a shell script template. When I built it and pushed, I had 71 files in the pull request. But it was all the same, so Karen approved it, and I deployed it.

Will continue tomorrow.

Oh, whatever.

September 07, 2023

So, today I spent some time working.

First, check out one obscure database, logging in to which was some kind of puzzle. Found there a table that contains “providers”. About 500 records; I had to add a value to a list of values for mexican records - there are 38 of these. Now the question is, what about providers in the database with which our end (back end) deals regularly? I believe I did create the “providers” table. And the data are kind of overlapping.

Oh, whatever.

Then I also worked on the two versions of that project; now I have already 101 files in the PR. Somehow the build still fails, and it's all about the script that does not find the coverage data. I'll fix it, but... well.

Also, started working on that “memory leak” in a pretty weird app. Looking.

Also, we had a “hangout”, where we discussed this and that. When I mentioned Norway, it turned out Davis just returned from there, brought his Covid from there. Still, he loved Bergen. I have to watch the movie “Menu”, it’s somehow related.

Etc. A bunch of people in our team would love to be musicians. Karen would love to be a designer. Well, sure.

Ok, so we did have fun. And I continue having fun.

See ya!

September 08, 2023

It was an exhausting day. Ok, I was fixing both the branch where I now have 101 files, trying to synchronize it with the one for the older version - it does require some attention. Meanwhile I was bombed with not very reasonable user support cases.

One Natasha was complaining about her playlist suddenly losing a bunch of songs, and she only has 20 songs, and she does not remember which ones disappeared. I spent some time investigating how I can check out her playlist. Turned out that there's a url for it, and there's an API by which I can download all that information (104 songs listed in total). But what I found: you see some Mexican song on a page, and what sounds is something in English, and the lyrics are improperly rendered. So I contacted the web guys, and wrote to our customer support asking what we can do with Natasha.

Another case is when someone closed his account, and opened another, via Google, and he does not know why his account is still free, while he says he paid. I don't know.

And yet another was totally messed up, so Karen wrote to that support person that we need more info; I also joined and said we need more info. But I only wrote to those support guys, and Karen was sure I was not participating.

Oh fuck, and that definitely is not the kind of job I'm created for.

In short, my sugar jumped up, and I'm still trying to beat it down.

By the end of the day I sent one small PR, with about 10 files, all in tests (and one README).

That's it. Tired. Tired of communicating with people.

Ok, and Chris also found time to tell me that I should be reading emails from support - exactly at the time when I was in the dentist's office, of which I informed them yesterday.

But. Do I fucking care? No. I'm immutable. Supposed to be immutable.

September 11, 2023

Ok, so nobody approved my small PR, and I decided to work on that memory leak.¡

We also had a “standup”, where two new guys were introduced, and we basically said good bye to Jorge, a Bolivian Scala ZIO guru. One new guy is from the North of North Carolina. A pretty curious enthusiast, hope he’ll make it.

Then I switched back again to the memory leak - started writing a “report memory” method - just wrote the name, as suddenly copilot wrote the rest of the code. I only refactored it a bit, to keep values lazy.

Ran the app (ok, for that I had to update the version of the load-testing infrastructure, never met it before, and it was not compatible with Scala 2.13, which was in that service. No big deal, upgraded, fixed some code (naming, etc), and I observed the memory leak, when I launch the server and pummel it with the load tester. Perfect. Now I have to find the culprit. But since it is WSDL and XML, these two usual suspects will be investigated.

Updated the Jira case, to properly inform the people.

Qian even reviewed that code (88 files, because of “copyright year change”. No big deal.

That’s it; worked a little bit in the airport, because I’m on my way.

September 12, 2023

Just a formality. Managed to get approved and merge one small step for humanity (upgraded oauth service to Scala 2.12), and still discussing that small shit with memory leak tracking>

Could not make it to Jorge’s talk regarding what should we do with ZIO in our codebase. (I’m not 100% sure anymore).

Had a 1-1 with Karen. I may have a false feeling that we are the best buddies. Maybe I’m naive. But anyway, talking to her is a pleasure. But I even could not extract from her her perception of her level of knowledge of Japanese. So-so, is a usual answer to everything. Still, we cooperate.

What next? Next is, I could not go on with the progress, because it’s fucking France, and everything is either slow or not available. I’m talking about the internet connection and about my time. Ok, whatever. Tomorrow.

September 13, 2023

In the morning I was feeling so good, started to do this memory leak tracking.

Tweaked the logging a little bit, and found that there are two sorts of leak:

- when requests are coming, there's at least several megs gobbled per request
- when the server is idle, every 40 seconds it eats between half a meg and a meg.

Ok, now what. Now i need a profiler. So I'm trying to figure out Java Flight Recorder, following hardly readable Oracle documentation. By "hardly readable" I mean they have gray font on white. Why? It's probably about Larry, who else. Have to ask Andrew maybe. Or Dima. Or... no, not Sandeep. We kind of hate each other.

Anyway, and we had a standup, then a planning meeting; I said it will take between 3 and 5 "points" to find and fix the leak. I could not say "4 points", because we don't have that coin. Only 1,2,3,5.

I actually fell asleep during the planning meeting. Thank god my cam and mic were turned off.

Later I woke up and worked with JFR. We will see. I also asked Jacob if we can purchase "ultimate" IntelliJ. He'll investigate.

There's also JProfiler (still alive), I worked with him for a bunch of years. But I would prefer IntelliJ, and I'm curious about JFR.

Our noobs are really cool, especially this North Carolina man (there are no guys or dudes in Carolinas). That's how I liked him during the interview: he's an enthusiast. I love enthusiasts. At times they break things (sigh, I was wrong when I fired JEF for breaking our Forth...).

Anyway, I'm still under jet lag, so not exactly behaving reasonably. If you could know how hard it was to do my daily dose of French.

Ok, an update. I decided to to runtime.gc() every 10 seconds.

Problem GONE.

So it's all about JVM parameters regarding gc behavior. No big deal to fix; and even I don't fix it, I can just run this gc thread every 10 seconds. Big deal, really.

How come I, a functional Scala programmer, am supposed to do these devops things, I wonder. But ain't it funny?

September 14, 2023

Now I have to kind of schedule my work. Because, well, in the morning I like to work, but then I stay till almost 23 working.- it is stupid.

Anyway, I did check that during the night no memory was leaked, I rest my case, and started looking for the right jvm params.

Found them, put them in build.sbt, started playing with it again.

At 17 we had a meeting, and Renato asked me first, and I told my story, and so he was happy, Jacob was happy, and I almost felt like some people became jealous. Wow. I grow my feelers.

Anyway, I slapped together my PR, but there were interesting comments from Qian, regarding “why do we need debug output in production code”. Right. Should I start quoting Dijkstra? I could. And I’ll find that quote, about a plane where all the measuring devices are removed after the plane is built.

Anyway, the problem is different now. It works ok locally, but in Jenkins it complains either about “stylecheck”, in the build file (I barely touched it, and more, the plugin does not allow me to say that a build file should be ignore), or about “coverage report” being in the wrong folder (that I already had in another project, and I’m also still clueless).

So by the end of the day I just split that 88-file PR into two; one with “copyright years being updated” (88 files), and another with 6 files that contain some meaningful stuff. I could have left just one update, the jvm gc params, but it won’t “check the style” anyway.

Well, tomorrow.

Jeremy, our new enthusiast from NC, is shining; he updated a bunch of docs, and now suggests to fix our builds by using sbt 1.4.7 everywhere (won’t work: Scala 2.11 requires sbt 0.14 the latest).

Oh, whatever. I like this kind of enthusiasm. I was like this... триста лет тому назад.

Anyway, feeling good.

September 15, 2023

Ok, managed to merge the updates for build.sbt (no shit, just one file, and is “not validated” by the “style check”. And I can’t disable checking it (then the checker complains).

But still, Qian does not like the idea of having a class called Dbg, with memory logging. He said I should keep it in git stash. He generally does not look like an idiot, but probably, partially, he is. He does not control his jealousy.

Still, could not do anything with configs for prod for that project. Basically, I did not even have to do anything with the code. Just update jvm params in those config files. And you know what? The “build” for config files fails. It fails on ammonite being unable to find coursier. Asked (carefully) on slack whether anybody had that problem. Howard from Caracas says: same here, I had to take Ammonite v.1.6.8 - the newer versions fail.

Fucking nice, our software have outlived Haoyi’s software. (Haoyi is the son of that famous Singapore president.) Haoyi, of course, does not keep that ancient version. So it’s nowhere to be found. My options: 1) manually generate the files - no big deal, been there; 2) convert those “Scala scripts” into regular Scala programs and not depend on Haoyi’s life path (he is in Databricks now, I believe - or gone; it was me who advised him to go there, not to Salesforce). I guess I’ll do both.

So, I was totally exhausted by the end of the day, by hardly having anything achieved. Oh, btw, I'm on pager duty. The best time for it is the weekend. No customer support calls. Why? No clue.

September 18, 2023

So, Chris wrote to me that I was wrong regarding OOM. I asked for a proof. And he showed to me that, well, it's not about OOM, it's about Out of Direct Memory. The thing grabs memory in JNI, and uses that memory, it's "faster". And it seems like (by now) that on every exception caught there's about 2k bytes left unreturned. Since there are tons of exceptions in mysql driver...

In short, I have to try to properly reproduce this shit and then find a version of the driver that works. Or just have a fix; must not be hard. But first, need to reproduce.

That's it for today. Ok, at 22 I have a meeting with Jacob. Interesting life, right? Second shift.

September 19, 2023

So, I spent my day working in buses, in a church, and whatever.

Discovered that a) memory is requested, but never returned. Tried, eventually, the latest version of http4s (Polish guys are supposed to be the coolest, no?)

The code grabs memory, and never, never, never returns.

At 10pm had a 1-1 with Jacob, and he told me that, weird, all this memory loss started happening two weeks ago. What the fucking fuck, that code was not changed since October 2022. Now what.

Ok, by about 11pm I managed to launch the September 2022 version. And? No fucking leaks? Now what. It's the same code!

Puzzle.

Also, I was on pager duty, and there were some cases, but I was anywhere except for the place I could react. Walking between two cities, sitting in a bus with no connection because my French phone was dead... Omfg. Tired of all this.

But we have nice people here. Davis said "ok, I've restarted that service". (Of course, the question why does it fucking fail all the time... oh, whatever.)

September 20, 2023

Still trying to introduce some time of schedule, not to work from 10 to 11, as I have a tendency to do ("because they crave").

Anyway, to log the behavior of all those “buffer allocators” I pulled out two more classes, their source code, to my repo, and added some logging.

Now what. No leaks of Direct Memory. But there leaks of jvm heap. That’s easier of course, just need a profiler. Experimented a bit. Found a class that, when represented as a source code in my directory, works perfectly (allocates and releases that JNI memory), but when used in its binary version, it does not. Well, and of course there is a mismatch between the source and the binary. That’s typical. That’s everywhere.

So, ok, as a result, I can now focus on JVM heap.

Besides, during the standup my connection was breaking all the time. Ok, I switched to my android (turning off the modem). Still breaking. Then what. Then I stopped my tests. Oops, it works now. So the tests were gobbling the bandwidth! (And the computer power, of course). Cool.

The standup was at my 17h; then a talk that Qian was giving, about some features he implemented. Features are okay, but the style, though pretty scalaish, is way too bland for me. Not enough abstractions. Juggling tons of values, constants, parameters. Not my style. But well, I’m not his mother.

Then there was a 1-1 with Karen, lasted about 5 minutes, during which I was entertaining her about my findings in memory leaks.

Then there was a meeting dedicated to integration testing. Our African QA was talking about what she does, what she can do, what she cannot do, and three managers (all females, because it’s QA) were asking dumbest questions possible, and she politely answered.

Then a long break during which I did my private stuff, and also worked with the documents (seriously, documenting the “architecture”, which was last time “documented” (partially) five years ago. Because I need a profiler.

Then at 21:30 a 1-1 with Jacob, and I entertained him too regarding the memory management in Jetty; and asked about the Ultimate IntelliJ license. Since it’s work in progress, with the upper management, he suggested to get a trial version. Omg, how did I forget it. Will do, but tomorrow. It was a little bit too late. Also told him how I enjoy doing what I’m doing. I really do. It’s like Rubik cube for me. Rubik kubernetes.

Good night.

September 21, 2023

This morning was pathetic. I tried to get into the vpn, no way. Took pictures of those stupid cisco messages, posted it on devops slack help board - silence. Restarted cisco Any Connect, etc. Tried to use another network (I have three: fiber, my french phone, my american phone). Did not help. So I did what I could: documentation, etc.

But after about 16 heures everything started working. Just judging by the time zone, it was most probably our devops or whatever woke up and fixed their network, and that was it. No worries next time, right?

As I wrote already, I've discovered that there is no direct memory leak in this service if I copy a couple of source files from that netty library into my project and use them instead. So, wtf? But instead, there's a heap memory leak.

Chris, as always, was getting on my nerves, so I a) ran it again, with the old code - reproduced both direct memory leak and heap leak. Showed him the logs. He said, ha, heap does not leak in production. I believe I know why: production has bad gc parameters, and for some reason it helps.

But anyway. I experimented with pulling the freshest versions of that memory management, then the freshest versions of the code that uses it, etc. In short, I've learned how to get the dependency tree.

But no, the freshest (this Spring's) code still leaks.

Chris is good in pointing out errors, but he is too slow in understanding what he is told. But anyway, I'll have to talk to him slowly, in simple words. He's okay, in a sense. A useful idiot, or something.

Anyway, so I continued looking around. And found, on SO, an advice: set "directMemorySomething=0". Which will turn on the leak checker (I know that when I was running my shit, leak checker was turned off, and I was clueless how to turn it on).

And now what. No leaks. No direct memory leaks, no heap leaks. Instead, at times leak checker throws an exception because some client code attempts to reuse a variable for a new buffer while it was pointing at an old, non-released buffer. And the checker catches it.

Bingo. Will slap together a PR tomorrow, and send it for review.

Our nubes are shining. They learn our shitty docs like they are egyptologists reading the discussion between a Man and his Ba. One of them already decided to pick up the "common library" I started, and start populating it. That would be amazing. How come we interviewed and hired two great engineers? I don't know, but it's cool. Let's see if we can also hire James, seems like he filed his resume to Jacob.

Oh, and Renato wants me to review the jira cases and update the epic, write the tickets so we could finalize what our partners want us to do but we never had time (our precious time). Ok, tomorrow. Mañana.

September 22, 2023

Having fun today.

Ok, slapped together a PR with the title "tracking memory leak"; wrote a long description what's inside. And because it's just a code update, and the solution against memory leak consists of a JVM parameter,

in principle, I could instead send a pr for deployment config. But see, I want the freshest code to be deployed. The only difference in the freshest one is that I added some memory tracking - under “debug” flag, so it won’t run if debug is off. So Qian, theoretically, would not object (although, fuck, you have tons of metrics reports, what’s wrong with one more).

And thern the hell was, err, let out. Questions! (Fuck, rtfm, right? I mean, the PR description). Chris was objecting more and more. But naive Carolinian Jeremy just approved it after having read the description etc. He’s a fast thinker, btw. So while Chris was calling it “cargo cult” (I’m sure it’s pretty insulting, and I’ll remind this phrase to him more than once), I just merged it.

Jacob wrote that it was not exactly polite, and it’s better to resolve the issues. Oh, well, of course. How can I explain to that idiot Chris that this flag means “use leak checker”? I mean, I wrote it, but how does he know I’m right? He could google it, right? But he already has an opinion.

Besides, the style checker suddenly complained about the code. So I’ll have to go figure out what’s wrong there. It does not complain when I build locally. Will check. Actually, I can deploy the older version, with the new flag.

So I now have two prs: deployment and an update for that merge. But now everybody’s scared, afraid of Chris, and just don’t want to get involved.

So I wrote to Karen asking what would be her approach, should I set up a meeting and explain, s-l-o-w-l-y, how the thing works? Prepare slides, whatever? I can.

And I went for a walk, as I did when I worked in Salesforce and dealt with local assholes.

OTOH, why not give a presentation? I can.

I already wrote about the origin of those memory leaks; that can be fixed later on. Let’s first intercept and prevent these leaks. But well... we will see.

As a result, I did not write the documents about the 2 or 4 or even 6 spikes for the updates requested by our pardners. The ones I promised Renato. So I wrote to Renato, and he asked me to make sure we’ll have it by Wednesday. I guess I’ll have to write that stuff (no big deal) Sunday, because Saturday I’ll be busy.

I really don’t like conflicts. If that idiot was not that adamant, I could explain to him. But he behaves like a troll. At work. Tricky.

September 25, 2023

Okay, it was a funny monday today.

Turned on my computer, tried to get into VPN. Ok, it sends me the code to my phone, but does not show the input window to enter the code. That’s Cisco Any Connect. Repeated it about half a dozen

times, stopped Cisco, restarted Cisco. Did not work. Then I rebooted. Did not work. But, otoh, half of my tabs in Chrome just disappeared. Did not I feel some relief? Well, I kind of did.

It turned out actually that everybody was having this problem with Cisco. Some blamed their network, some just complained.

Meanwhile, the project I've been working on stopped compiling. I mean, it was showing Java errors (it had some generated Java code inside).

I started investigating. No clue. I had the same project, but master version, in another directory (I have just two directories for parallel development, in Salesforce I had four). Same shit. Same shit. Started investigating deeper.

What I found: I have a different JDK now? Something called "Semaro", from IBM. I know IBM has been brain-dead for about 30 years. Not it pushes, magically, via Apple, its buggy JDK. Enforces.

So I spent some time looking for what exactly build of Java 8 should I use, and where to get it. Seems like OpenJDK's 8 is kind of gone. And we are still on Java 8.

Discussed it with Jeremy; he suggested using sdkman.- ok, I installed sdkman. Did it help? Not exactly.

Eventually I just found a release of openjdk 8, installed it manually. The only problem, IntelliJ does not recognize it, and insists on having that Semaro, which I had already deleted.

But gradually IntelliJ got used to the one I use.

Chris still discusses the one-liner fix. Now he is objecting to the contents of README file, where i listed JVM GC recommendation from Oracle, together with the link to Oracle. Ok, I already mentioned "cargo cult". I guess I should be using this term more frequently, to scare the shit out of this idiot.

But since it started working on my machine - that including the mysql server that was dead all Sunday - I was able to grab proper logs for the doc. Wrote a doc, showed it to Karen, and asked her, what would she suggest next, show it to the team, or? She suggested to get the graphs from Grafana. Omfg, how can I get graphs from Grafana if my stuff is not deployed? So I just grabbed the graph that shows memory leak. Had more discussions with her, regarding what is missing in a three-page explanation why we need a one-liner jvm parameter.

Playing an idiot... I really don't like it, especially if one idiot is playing a Wise Man caring about performance.

Tomorrow I'll have a 1-1 with Karen, and I will suggest, if Chris is still resisting, to close this case with "won't fix". Let Chris fix it. He's jealous, Qian is jealous. That's our profession, everyone is jealous. As if I'm fucking their wives in the background.

September 26, 2023

So, posted again my PR for review, and Chris did not hesitate to put new demands. “Remove these sentences from README.md” (do they break the build? Are they incorrect? - no response). “Remove this class that logs the memory usage to logger.debug” - why? Do we lose anything? It’s all lazy, and has 0 impact on performance.

And I asked him if he has any objections regarding the quality and correctness of the code and its performance. Then he eventually said that he refuses to review this, and I should ask someone else.

So during the 1-1 with Karen I discussed all this, and she agreed to review - she had some RFC (requests for comments) - I love adding comments to code, so I did. We basically laughed talking about all this; now she’s discussing all this with Chris, in github. I hardly participate. Just add some clarifications. Karen understands the code, and Chris obviously does not. I probably should ask Jacob too, how can I get rid of this incompetent reviewer. (Sure, using the word “incompetent” is forbidden. Last time I heard it in America, it was in the movie, “Home Alone”.)

Anyway, whatever. Let them discuss. It’s not my money that is being wasted during these discussions. Karen is a manager.

Also, I returned to that Oauth code transition. Gradually. Also, I started writing some new “architecture” docs; every time stopping and starting to think, no, this is a bad big picture (the previous two versions were bad too - including the docs from 5 years ago).

Anyway, I relax and enjoy. And anyway, the next quarter is coming, and there will be new stuff for me to do.

September 27, 2023

After a bunch of nonsensical microscopic “improvements in README”, I got that shit approved, and merged it.

Having explained to everyone individually my further plans of testing it in Staging, how, etc.

But what do you think. Jenkins still resists. First, it does not find Gatling library on mvn - the one that my sbt finds easily when running locally. Second, it insists on having already built this shit, so this version number is taken (ok, by “client” part). I’ll have to set the “override” flag somewhere in sbt configs. Later, later. I’ve just bumped the version manually now. I really want to go to sleep seeing that this shit is ready for deployment.

We had today:

- standup meeting
- Mahesh was talking about how he conquered mongodb that now deprecated map/reduce (obviously, after monads having lost to programming idiots, monoids are next)

- planning meeting for the next sprint of our marathon (I was asked whether I'll continue investigating/fixing/deploying this memory leak - that's it for the next two weeks)
- sprint post-mortem, where we were all thanking team members for this and that. I thanked the two bright new guys for their enthusiasm and docs updates, and the whole team for being nice and non-conflicting (yeah, right); and everybody was thanking someone or more; I was thanked for bravely investigating and fixing this memory leak (this old plumber can fix plumbing!), but in the end I also congratulated the whole team on successfully celebrating Thanksgiving two months ahead of schedule. This kind of sarcasm was not understood by the masses.

Oh, and in a separate thread, our partners got a lot of sanity. They removed all this “sort the table in the database”; instead they only left a question in their doc: ”can AMP backend (that’s us) return records in a given order. (There’s a good reason for this: paged output.) We sure can. Discussed it with Bhargav, and the issue is closed. He is also almost ready to deploy to prod the code that does not wrap that specific boolean in quotes. Mission (almost) accomplished.

Also, at that planning meeting one girl complained that it’s a lot of work to manually get a list of stations, per country, and fill in a spreadsheet. I told her I can help: step 1) download it from partners; step 2) use jq to get the necessary data and produce a csv; step 3) import csv into Google spreadsheet (or into Excel, whichever she prefers). Everybody’s smiling. Everybody’s happy. Jai guru deva om...

That’s it for today.

BTW, did you notice how John Lennon and Steve Jobs are similar? Two charismatic assholes pretending to be peaceful wizards. (Fuck, who’s asking... time to change my id.)

September 28, 2023

Crazy day. First, it was declared that we have a 7.5-hour planning meeting. They more or less honestly reported the status, the growth, etc; well... not impressive, but manageable. Then a weird guy was reporting how we bravely produce hundreds (or thousands) jira points, and never catch up with the planned estimates of jira points, which means that nobody’s cheating. Reminded me a Columbian guy, he owned a farm in the mountains (of Columbia), and he visited it like once a month, and he made sure that he gave the peasants enough tasks so that they would not be able to finish them all, and, consequently, won’t dare to complain.

Then suddenly that meeting ended. It was actually split into smaller parts, team-wise, so I was only invited to three, of which I attended one and a half.

I deployed my shit to staging, but it was not working. Why? Because linkerd was crashing. Good that Mahesh was informed, and removed linkerd from configuration. Then my stuff started working. And so far it does not leak the memory. I’ll have to run my load test against stg, but somehow the existing Gatling test does not want to talk to an external server ([https](https://), credentials, etc, and everything’s in

xml...) Whatever, I'll write my own load script, big deal, curling it in a loop. Have to find logs in new relic, too.

So I went to the movies, to relax, by 20:30. But the movie was pretty stupid ("something weird in Venice", or something). Kitsch. So, after about 70 minutes, I went home. Opened my work computer.

There's a panic now, coming from, well, Columbia! Wilmar writes that the service he's working on suddenly stopped accepting data from our partners. What data? The same fucking data as the official import job was importing for years. That service, written obviously by "other people", duplicated the job (which was going to be removed soon). And their code wants that specific boolean in quotes, and if it's not... You know. I wrote to Wilmar, wrote to Karen, wrote to poor Raghav, whose PM Jaime already created a jira case regarding "roll it back".

WTF, right?

I also had to add the fact that nix just stopped working, or, rather, showing up on my machine; Jeremy says it's not needed, just need to install a proper ammonite. I did it according to his new docs, but it does not work. We agreed to get together tomorrow and follow his instructions and see what's going on. Should I also say that IntelliJ again does not see OpenJDK@8? Well, whatever. At times it does. SBT is ok, in this relation.

What's interesting, that crazy Chris is now in good phase (I believe he's a bit deranged, MDP or something), and is pretty helpful.

September 29, 2023

So, what's up? Since last night I had deployed that shit to staging, I had to pummel it with load test. Except that the loadtest was not exactly ready to do it. Just did not work. And also, what's the url for the service, exactly? Omfg.

Then there was a "slack standup meeting" (my favorite: just send a message in a couple of lines describing what you are doing).

Then, since I was stuck with testing, I worked on another project (the one with 102 files updated), to send chunks with updates. Ok, I did. Except that I was tired from config service refusing to return metrics.conf, and just copy-pasted it into "resources" - for local testing, of course.

And then it was Chris again. "Why do you do it", etc. I wrote to him in private, asking to take a break from picking on me for a while, and just ignore my pull requests. Suddenly he changed his tone, from his usual trolling to a normal friendly one.

I saw those friendly people back in Russia. Once there were three of them, and when they started bullying a girl on the escalator, I told them not to; they stopped, and got friendly; then, in the street, they waited for me, and assaulted, etc. I know this trick.

But anyway, he seems to have got the meaning of my message, so... do I care? No, I don't fucking care.

Also, he approved my PR, that's all I needed from him.

Then we had a “planning meeting” in the team. Thank god nothing where I was participating was mentioned. It was a “classical scrum” - a big project was vaguely, briefly explained (on most occasions somebody was already on it), and then we were supposed to vote whether it is S,M,L,XL, XXL, etc. Some are humongous. Everybody gave XXXXL to such projects. But does this look like serious work? No, of course. You are supposed to guess. Why not give it to chatGPT to guess. The bot can produce nonsense real fast.

Again, do I care? I just smile. Probably that's how Chinese engineers view their work in America.

Ok, and at last I got free from meetings, and could start thinking.

Eventually I decided just to write a script with curl inside. It turned out the script already exists, and is pretty much usable. So I only added a line to README.md, to make sure I won't forget what I've been doing, and pummeled Staging for a short while.

Then I looked around in New Relic (last time I used it 6 years ago), found the logs, saw the exceptions that memory leak checker throws - good! Then I looked in another log (why is it not in the same log?), saw fresh reports of “0 direct memory taken” - perfect. Mission accomplished.

So I copied chunks of these logs into the doc I wrote, send the link to Karen, and was going to send it to Jacob, but saw that Jacob was already asking me how I am doing. I'm doing just fine; it's a happy week end. Sent him too the link and shared my plans. It turned out that our SRE reboot that machine every 12 hours. Manually (FIIK why).

Sent a PR with this logging disabled (we don't want to see “memory abused = 0” logged every 5 seconds. We are good now.

And closed the computer. Oh, before I also merged another PR, with 8 files out of 102 committed.

October 2023

October 02, 2023

Ok, we are a little bit slow, right?

In the morning I checked out my work slack. There was some panic. The service I'm in the process of fixing (the memory thing) was being manually fixed by SRE twice a day, and now it seems to be crashing every two hours. So, could I just deploy it to prod, or at least to staging? Well, no, it's not safe!

So my PR at least was approved. Guess by whom. By the guy who was blocking this shit. Ok.

What's strange, he was also adding comments like "will it work?" And my answer was "I hope so".

Strangely, it merged ok (unlike Friday), and well whatever.

But I wanted to see the logs, which I saw Friday, but not today.

By the end of the day I managed to get Karen's attention.

And, well, WTF. I was too tired. Stuff works. And nobody fucking cares.

Will start deploying to production. Tomorrow. To a canary. While it's still crushing on production For the last at least month. And for the last five years in general. And they are afraid I'd break anything.

Does it make sense? Not to me.

October 03, 2023

So they wanted me to deploy to canary, as the staging version behaved ok.

I slapped together some config for canary, then Karen pointed at the defect ("0 pods for this canary"), and I fixed it and deployed it.

Then I was looking for ways to find the logs in this canary - logging in to production was kind of questionable, and why?

It turned out eventually that we have two productions, and one has a console, but another does not. Where was it deployed? To both.

Karen told me all this, and by the end of the day I saw my graph in Prometheus. But I need an overnight graph, to see that it behaves.

That was the end of the day.

Meanwhile of course there were meetings, meetings, more meetings. At least with 1-1 that I had two, with Jacob and Karen, we did it fast, and I told them, ok, now we have some time to get some rest between meetings. And everybody was happy.

I don't know, they company pays us for bragging? How about doing some real work?

Well, whatever.

October 04, 2023

So, this morning I went to prometheus (I don't care about the logs, it's not essential), and saw that the line is just flat for Canary deployment.

Fine! Update the config file to deploy to prod, and wanted to run a build and push - but oops. It does not build. 400 from the server. Bad credentials in a message. (why 400 then?)

What happened? Wrote to our team's slack channel. And then there was a flow of advices. Like "did you log in to vpn?" (did not I?) "did you log in to aws?" (did not I?)

Eventually it turned out that while I was speaking they changed the ids of aws clouds and their names. How come I did not know, right?

Then there was a flow of question to me: - why do I do this, why do I do that?

Guys, I run scripts I wrote half a year ago, after getting information from Safa and Swathi and Karen, now what?

Then Serkan patiently told me what to do, how to update the files; I wrote a PR to update the docs, posted it, and asked Serkan to review. He's a real, true devop guy, and he gave me tons of meaningful advices. That's what I'll be doing any time next time.

And suddenly Santosh wanted to chat on Teams. I would not say I know why, but we did spend about half an hour, and, after many attempts, my files, like `~/aws/config` etc were properly updated. Yes, the updates were mostly updating. But at least I know what EKS means.

In the end Santosh recommended to write down the commands we ran, in a notepad file (he himself uses VI), so that I don't forget. That's an interesting approach. An Asian approach. Europeans just document stuff.

Then I spent some time updating my deployment config (set the number of instances to 0, and the like). By the end of the day I was ready to deploy to prod (the graph of which was showing how all nodes gobble the memory, but not the canary one which I deployed last night).

And I almost deployed, but strangely, on jenkins, that deployment page was empty, no previous history. So I asked the people. It turned out, ha, when they created a new cloud, the history just disappeared (does it make sense IRL? not for me)>

So, hold on, I'm deploying to prod now.

October 05, 2023

Pretty short.

Ok, I got up at 2:30 at night, with strange paranoia; went downstairs, opened my work comp, and was reading Karen's complaints that HEAD requests return 400 on services on the new server. Eventually reacted rather stupidly that it's probably because of the new server configuration.

Then I opened Grafana, and looked at the graphs of memory gobbling. Last several hours: almost 0 (not exactly a zero, but it can stay like this for a couple of months. Then I looked at the graph for the last couple of weeks. Dramatic. SREs had to kill the process every day at about 4PM Eastern. Then I look at the graph for the last 8 weeks. And I saw why they started doing it. Before they did, the wave was going up-up-up, twice higher than the level they started killing it - and then magically it was going down. Probably there was some "direct memory garbage collector" that was turning on when the memory was reaching the limit. That's how it was.

Anyway, since there's hardly any gobbling now... I had updated my doc, and felt rather happy about all that. And went back to sleep.

In the morning checked what's going on: nothing.

At 11AM EST, when I just got out of my class, I wrote a "status report" via slack on my phone, a one-liner: "all done with 11107, problem solved".

Then, when I got home (it's about 6 minutes of fast walk), I went to the meeting, and there was no meeting. Everyone was reporting on slack. Cool!

Later I had a chat with Jeremy, who was given a task to upgrade one service that uses Play framework, from Scala 2.11 to Scala 2.12 - Jeremy was over-optimistic, so I showed him my branch, my efforts, and suggested to cooperate over a common library, and adapter library, that would make our code with business logic independent from those Lightbend idiots. He'll look into it. At least he knows now that it's impossible to convert a Scala 2.11 Play app to a Scala 2.12 Play app in a week. Or in two weeks. That's why we have some projects that have been in such state for the last 5 years.

So, maybe together with Jeremy, who is really talented, we could do something doable. We will see.

Then I had dinner, and now I'll get back to my work computer.

October 06, 2023

It was quiet today. So I worked on that oath project upgrade (with 102 files in pull request). Trying to make the code structure not depend on the version of play framework. WIP. Now three tests are failing. Anyway, whatever.

Meanwhile, also added some pertinent info to docs (how to deploy to canary - it was a secret knowledge until today), and got an amazing amount of good comments from Chris. Seems like he's just jealous, and when he feels like other people are posting something that he knows better (why t.f. was he keeping it secret for years?), he immediately gets involved.

Anyway, thanks to him we have some good docs.

At our 20h we had a hangout. Since Jacob was not here (some personal things, he has lots of them, with three adopted kids), we just hung out and talked about places, cats, dogs, Halloween, everything.

Seems like we have mostly cat people, except Renato, who has a rather crazy dog. Recently pissed on by a scunk. Recently... no, the story goes like this.

Renato was in Starbucks (why do people go to SBUX, I wonder) getting coffee; his dog was tied to a metal chair outside. Then the dog decided to run away (maybe saw another scunk). So he started running. With the metal chair behind him. Renato ran after him. For about a mile and a half. The chair was eventually destroyed, so Renato had to pay for it. But also the chair was banging anything that the dog was passing by. Like cars. Renato did not count the cars, but he was hiding at home, together with his dog, for about a week, after that.

As to coffee, Jeremy has a cat named Caffeine. A nice cat (I saw the picture); the only thing the cat does not like is when Jeremy puts on him a leather biker jacket. Then Caffeine just freezes. Jeremy also is taking care of two more cats, both belonging to his ex, who has gone to Vermont to visit his relatives, so while he, the ex, is away, Jeremy takes care of all the cats. One of them is a kitten who does not understand cats boundaries, so there are problems.

Safa showed his beautiful cat with a complicated name, something like "Maushuallia".

Karen does not have any pets. Other people were just silent and with cameras off. I also had my camera off in the beginning, since I'm in a robe, after shower. But eventually I just moved the camera up, so that they would not be watching me like this.

Basically we spent a nice hour chatting.

That was it; I merged and deployed the docs, but at this time I can hardly do any job now, too late. Mañana maybe.

October 09, 2023

Not much. 5.5 hours of meetings (the remaining half an hour is still ahead). But anyway, reviewed some code, and also update that "Oauth Scala 2.12" code. It's still not ready yet. Reposted my small pr that is going on step by step.

That's it for today.

Jeremy writes that he was “boosted”, and now he is sick. Karen is just sick. Jacob showed up for one of the meetings.

October 10, 2023

I decided to dedicate today to writing docs, so that’s what I was doing. The format of that .md is not exactly the standard .md, so I had to read some tutorials. Well, whatever. The picture of the architecture of our system is just changing while I’m reading it.

For a break, I decided to reproduce the PR I was holding for 2 months (108 files), and create a new one. Funny, some business code changed here and there, so... so I hope I see better now how to separate it from the underlying layers.

Then there was a presentation of machine learning people. When someone says “excetera”, I get suspicious right away. Now I saw that my suspicions were founded.

“It takes two days for us to process the data”. What?! First they read it into Postgres. Then the “ingest” the data into SQL Server (already idiots, right?). This “ingestion” takes 12 hours. (idiots, right?). Then they “process” the data. Why not directly from Postgres? “because our application only works with SQL Server” (meaning, all programmers are gone).

Someone asks them why do they do the same calculations twice, in two different jobs. The answer: “We have two different MongoDBs, and one job works with one db, and another with another”.

Karen asked questions, and they always had the dumbest answers possible. Oh, sure they use Spark. I asked them, why then SQL Server, why not let Spark work directly with Postgres or use Spark Streams and be (more or less) live. (They actually call their application “live data processing”). They said they tried, but it did not work for them. Neat!

Then, pretty late in the evening, we had a technical team meeting. We need to upgrade JDK (8->11 or -> 17 or -> 21), Scala 2.11 -> 2.13, and Mongo something -> 7.0 (which requires JDK 11). It was a pretty reasonable discussion, and I’m amazed how good the team is actually.

Then I worked a little bit, and now it’s almost 11pm.

October 11, 2023

Spent the whole day doing this.

Took the whole contents of that old pr, with assumingly 108 files (actually kind of less), and copied them over to a new branch.

Ran tests in the new branch. One test fails (a server returns 500). The same code in the old directory passes all tests.

Started investigating that 500. So far found that Play framework's logger throws a cast exception while trying to “log” an incoming http request.

Did anybody teach Lightbend programmers to write code? Casting. In Scala.

What's interesting, those exceptions in logger are not reported anywhere. They are just converted to HTTP 500.

There are more cases in Play framework. Do they have tests? Probably not. Or not enough.

Anyway, that's how I'm spending my time.

Also, wrote several jira cases for that scala 2.12 upgrade for that oauth module. Split the job into parts.

At 20:00 we had a meeting, reviewed the results of our last sprint. Seems like I'm a hero. Saved our devops from clicking the “перегрузка” button every 12 hours. But it's okay, our team is really cool. With two new guys we are even cooler.

I would say that it makes sense to hire more gays, but there are different kinds. I remember one celebrity from Google. Not a bad guy, but impossible to communicate. Well, whatever.

In 32 minutes I'll have another meeting. Let me finish my dybrs first.

October 12, 2023

It's still WIP while I'm writing this.

Jeremy is hyper-active in a bunch of projects; weird, and cool.

As to me, I'm still trying to trace wtf is going on. Probably the biggest issue is that the logger that throws the casting exception is different in the version for Scala 2.12 compared to the one for Scala 2.11.

As to why it runs here and does not run there: sbt clean does not exactly clean. You need to manually remove all target folders. I just pulled master, and it does not even compile. Why? Because of those generated files in target that are interfering.

Anyway, I'll fix it.

That's it for today. Nothing interesting. Just a mess.

Ok, an update.

What was happening actually. Our service code passes request logger to Play, as a “filter” (people obviously didn't know what “filter” is, and neither did Play framework people. This “filter” is called to “process” a request; in our specific case it's just logging the request. For logging, Play framework casts the request to the type it thinks it is.

In the version for Scala 2.11 Play framework, when we “add a filter”, add that filter, but its “mock version” ignores all filters. So the tests pass, nothing crashes.

But in the version for Scala 2.12 the guys in Lightbend decided to improve that “mock web service processor”. And they add the filter, and the filter does not accept a “mock request”, and crashes, and Play framework ignores all the diagnostics and just returns HTTP 500. With an enigmatic message.

Ok, what I do now? I’ll just fix that code, have our own version for all that, and that would be it.

Hopefully. it will work. Everything’s pretty slow.

Ok, and the last news. After I fixed one stupid piece of code in Play Mock code, all my stuff started working. They had a bug in their mocks. Neat. But I’m kind of celebrating.

October 13, 2023

Pretty boring today.

Karen found that I forgot to deploy to prod, in June, “market import job”. I actually did, but not the right way; and more, the build number was like 10 versions behind; how come? Who knows. Anyway, sent for approval.

Otherwise, working on cleaning the new PR; 101 files is kind of too much. Still playing with it.

In the evening it was “all hands”; it started with one guy playing a video how his 9-year-old son is playing some vr game. I turned it off in a minute or two. Let them have their fun, I’ll be having mine.

October 16, 2023

Did deploy to prod the “markets import”, which was 11 builds behind, strangely. Not that it’s important, but it’s better not to have surprises.

Then worked a little bit over docs, and otherwise over that “transition to Scala 2.12”.

In the evening there was something called “open code”, a company-wide gathering to show some code to each other, to talk about this and that. I tried to introduce some organization into this mess, but so far it was just that a bunch of JS(TS) programmers were showing their boring JS/Ts code to each other. Ok, anyway. I hope the ideas will settle in the organizers heads.

That’s it for today. I believe I’ll need to slap together some kind of “test queue” on my mac. The thing is, one test suite is too long, and if I’m working on the same project (different branches) in two (or more) different directories, and they can’t run in parallel (dealing with the same db config etc), I’ll need to queue them - and then safely do something very different. Will think about it.

October 17, 2023

Not much happening. Merged one pr with the tests updates (that's a preparatory job).

Maintaining the “95-file PR”, making it align with what's in master; but it seems like github does not get it, and I'll have to prepare a new pr, with the same files. THe thing is, github does not see proper latest master updates, and compares god knows what with god knows what.

Anyway, no big deal.

Also, started working on our play library (a reasonable adapter for that crazy library). We will see.

Karen asked for a link to that 95-file PR. Big deal, you know the project, just find all prs and/or branches. Sent her the link.

That's it for today. Ok, meetings, meetings. Discussions.

October 18, 2023

So, what's up? Working on a couple of PRs, one for Play, another for Oauth. The latter (a new version of PR) has just 46 files. Will have to decrease this number anyway. Split into chunks etc.

That's it for today. Hope to slap together one and send for review.

October 19, 2023

Jeremy sent me a pretty good review for my update for playpal; I did all this, kicked out some classes that yes, are not needed in the latest version of my server.

But overnight Jeremy also added suggestions to upgrade to using the latest releases of mockito and play and scalatest and scalatest for play. That did not compile, so I wrote to him.

And you know what? We got on Huddle, and he did not bother to tell me which imports to insert and which classes to rename. Of course I could do it, but my approach was to minimize the changes, knowing our reviewers. But Jeremy is a different kind of reviewers, and he does not mind to have it all done properly. Neat!

In short, by the end of the day I had published that library, and told him that he can use it (he needs it too).

Cooperation!

There was a meeting, later, regarding the need to change user registration processes, and Taka, the main developer, was pretty shining regarding the suggested solutions, so I hung up earler, since ok, we have a profi solving a problem, why bother to interfere.

October 20, 2023

Jeremy already uses my release of that paypal library. More, he suggested to upgrade the code to be multi-version, that is, to compile for the variety of scala/play versions. Pretty ambitions. I totally applaud.

As to myself, provoked by his idea of switching to the latest mockito and scalatest etc, I'm now upgrading the oauth project to be all that compatible. Makes sense, for the future.

And basically that's how I spent my day.

October 22, 2023

I had a dream

We are in a SCRUM meeting, and one manager, not from our team, interferes with a critique: "you people don't do a meeting correctly; you have to be more active and more comprehensible for the participants.

And I get up and say this: "Let me explain it all to you. Our team's product, AMP, is the heart of the whole software system of our company. Now imagine you are at a meeting of heat surgeons. Will they talk about jira points? No. They'll talk about tools, medicines, surgeries, complications. Most probably you'll understand zilch in what they say. But if you replace them with paramedics and illiterate healers, you can close the shop. Although you'll understand everything they would say. Our backend is written at the complexity level λ_2^{\rightarrow} ."

At this moment the dream ended, and another one started, with a black cat on a couch.

October 23, 2023

Since we were driving to the sea today, did not do much.

Mostly, attempts to run tests with the "latest versions" of PayPal, which leads to "abstract method errors", caused by mismatching libraries versions.

By the end of the day just wrote to Jeremy that my shit is failing with his shit. He wrote back that he is investigating artifactory problems with Qian.

Ok, enough for today.

October 24, 2023

Today was even worse.

Ok, I decided, let me try this with the previous dependencies. git stash sbt clean compile test

No, does not work. WTF? Seriously. git status - nothing new. WTF.

Ok, copied the directory to a backup, git clone ... - some tests fail. WTF, they did not, and they don't in github.

Did it again.

Also wrote to Jeremy, asking if he can comment on the tests failures that I observe (like 'abstract method exception'). He had tons of words, eventually suggested to switch to a slack channel.

Then we got meetings.

After the meetings, there was a flow of slack discussions. My explanations regarding wtf was going on (and "can you try to reproduce it?") did not work. More, it turned out, he merged to master his updates, which are supposed to work for all three versions of scala (2.11, 2.12, 2.13). How? The code in that Playpal library is different in different versions. Also, he either forgot, or did not bother, to include the adapters that I added to the library, although he did review them last week.

At 1PM EST we had a 1-1 with Karen; so we invited Jeremy and Davis. There Jeremy explained that the working version for 2.12 he just ignored (that's what he approved last week), and he seems to not remember anything we were talking about regarding this branch. He just made a "universal version". What kind of magic can it be? The rest of the hour I was asking him whether he tried this build with any server (Davis was nodding, Karen was just silent), and he said - no. His tests in that library passed, that's enough. The thing that this thing won't be working, e.g. for the reason that Play library has different structure and APIs for different versions, did not convince him. I'm not Chris, I won't go check his code. Will ask Jacob about it, later. Or Karen.

But the general idea that a library to be used in servers should be tested in those servers, and if something fails, it's the library that should have more tests, did not penetrate his hot head. "Here's my leg, it does not hurt."

I actually expected him to cooperate with me regarding the tests that were broken after I (I must admit) accepted his idea of taking the latest play, scalatest-play, mockito, but all I heard from him was something like "if your server does not work with the new library, fix the server". Actually what everyone will do is not to use the new library, that's it. Because it breaks the server.

Oh, whatever. Then, after that meeting, he wrote on slack that he has a migraine, and so will be AFK. (this means "away from keyboard").

Then Qian was giving a talk about kubernetes envoys (or is it http envoys?). I never understand what exactly Qian is talking about. He always talks fast, with a thick accent, and starts from the middle. I believe the same happens to everyone else. I probably could talk to him about how to give talks, but well, don't I have a migraine?

Actually, my work computer did. It discharged, and did not want to charge. So I closed it. Will get back to it right now, again.

October 25, 2023

Ok, I rolled back the versions of the stuff that my playpal depends on, and, since the branch is kind of free to go, just merged the updates there, and built the library (published to jfrog).

Now I'm trying to figure out how come the library is not available from sbt on my machine.

Of course spent most of the day trying to figure out wtf is going on in general, it was by the end of the day that I got to this specific problem.

Also I'm on pager duty starting today.

Also we are having yet another meeting in 13 minutes. We had a planning meeting, and when Jacob asked Renato how come I was not mentioned (for the next two weeks), Renato said that when I finish the case, he will put it into the sprint. Kind of nice, right? Like that turtle competing with the rabbit. I must be there to compete.

Jeremy, unfortunately, does not have any of his shit working. I mean, even publishing/deploying. Problems, problems. Oh, whatever. I'm on pager duty, I'm not supposed to work this coming week.

October 26, 2023

Well, did not do much. Was busy moving from one place to another.

Jeremy now notifies the team what will be delivered not later than today. Funny. No, it was not.

I still have problems with jfrog, but no time for that. There was a “pager duty handoff”, and Davis left me a couple of cases. I could not figure out wtf it is about, and Sked. He, and Karen, and Chris provided me with full information. In short, there's a collection of jobs, and the jobs should never run in parallel, but they do, because one job does not finish, when the next one tries to start.

Now imagine. They are cron jobs. Does it make it all more complicated? It should not, right?

But it's kubernetes, dynamodb, etc. I'll have to find a simple, and obvious, way to coordinate these dining philosophers. Is not it really funny?

Then, at 20:00, had a discussion regarding multiple anonymous users, and how to identify them properly, so that we would not look at them as unknown every day. We have... many of them. Unusually many. Ok, it was a good discussion. See, when the users deanonymize, and they have accounts on multiple devices, what can we do? I had that stupid trouble with my Toyota account in France. Had to create another email address.

October 27, 2023

Since I'm on a pager duty today, I'm not supposed to be doing any useful stuff, just be on alert.
Часовой есть лицо неприкосновенное.

Jeremy, meanwhile, failed to report the success he promised to achieve today. He reported that he's sick today. Also complained about a 100-file PR that is caused by 'style fixer'. Sure. Did not I tell them.

Anyway, since, as a pd guy, I got a couple of those jira cases, where I was supposed to take care of proper timeouts for some jobs.

So I learned some kubernetes. It turned out kubernetes is just fine. The complaint that Davis had was about jobs overlapping, but they don't. Then Chris wrote in that case (one of the two), that the problem is that we should make this job much shorter. I agree, a data-scraping job that lasts for an hour sounds fishy. Ok, I agreed, updated the description. But then Davis said that first we have to decrease the timeout and see if it fails again. Or, before that, I'm supposed, as a pd guy, collect the statistics of the time it takes those jobs to finish, and then... then what? Find an exact timeout value for them. Which I, frankly, consider a stupid idea. Saw it in Google, where every Friday afternoon pd was ringing, because boys and girls in Brazil were hectic setting up dates and parties, and our borg was not flexible enough to take care of this weekly usage boom, and also we could not (we save energy!) add more servers on Fridays.

Anyway, I'm still lazily discussing it with Davis, in the style of Denis.

Then a guy from some other team asked to add yet another Mexican provider. No biggie, it's just a sql statement. Log in to the write cloud, do psql, run one command. Next time I see it, I'll write a script for this shit.

Since I still have some working hours, I'll take a short break, and then'll get back to my work comp and do some stuff I really want to do.

Ok, and also sent a PR with docs changes. Not a big deal, just 10 files.

October 30, 2023

Today the time difference with EST is not 6 but 5 hours. I was not aware until like 15:45, in Intermarché, when my phone told me. Anyway, as I'm on PD, I spent the morning investigating one pager report (not enough disk space in one node where postgres is running). Not a big deal, but I had to ssh into that machine, and how? First, the docs say I had to decrypt the password file that I had successfully received from aws. To decrypt, have to call amp-kms. Wtf is this? It's in nix. What? I don't have nix... btw, how come? I did have it before, and I could run nix-shell, etc. Not anymore. That was my comp upgrade, I guess. So, should I install nix? I tried. But my os is younger than that nix, and nix does not get installed. Eventually... very eventually, like after a couple of ours experimenting and

even asking (whom should I ask, everybody's sleeping), I figured that I just have to call aws kms - after reading the code for that amp-kms that I found in one repo. Sounds familiar. Do we have aws kms anywhere? Sure we do. I wrote that script for encryption/decryption.

So, step 1 done, I decrypt the password. Now what? I'm still not allowed to ssh. In the docs there's a hint that if you can't, try to get this config file from... from where? They don't say.

So I just asked on slack, and started doing other things. Like that paypal, etc. Strangely, nothing compiles when I rolled back some bad changes. Bad because I counted on our Jenkins to deploy what it builds to jfrog and make the library available. It's there, but not available.

See, I'm a programmer, not a devops. I don't even believe in their security game. Imitation game - you know how fake this movie is. Same with the rest of security. When an app on your phone sends an sms to your phone and asks you to type in the code from that sms.

Oh, whatever. Eventually it turned out Davis had already fixed the disk problem, and I was too lazy.

What else did I do. Did some docs. Even deployed them by the end of the day. Tried to get an approval for my library build, but Davis (again) suggested I just rename the repo to keep the naming system. Kicked out that stupid PR of mine.

Also, one nube, Samin, I wrote in a review of his code that having 9 parameters of a case class, where all these parameters are of the same type is not a good idea - and he gave some strange excuses why it should be like this - do I actually seriously fucking care. I approved it, but wrote to him that (while actually his code is pretty good), this piece is brittle and may cause future problems. We will see, time permitting.

What else?

Oh, "standup". My standup started when we walked out of the store (Intermarché). I just walked along the streets, holding my phone at my ear, and participated in the meeting. Pretty neat, but you need much more attention than in usual meetings.

Whatever.

October 31, 2023

Such a mess.

Ok, there were a couple of pods running the same service, and I was getting "pager duty notifications" regarding their misbehavior. The solution was just to kill one, and k8s would resurrect it. But when the resurrection did not help, I was puzzled... until I figured that it was another pod with the same service, having problems. No big deal, though.

What else. Investigated some cron job solutions for two similar jobs that are still distinct, although the same code. Just different params. I'd love to have a sophisticated cron schedule in k8s config, but no clue how.

What else. My lib does not build because... because there's a lot of shit failing. But why? No clue. And the "playpal" part... oh shit, I don't know. Seems like what I do is not my area of expertise. So it's pretty pathetic, all of it.

November 2023

November 01, 2023

In the morning I checked my “pager duty” - no problem there. Looked into my “poweramp-lib” project, found why t.f. it was not building (some credentials were not in the right place... that’s what I thought). Then I saw a PR from Jeremy, he created a 358-file PR, for the build system, with the only purpose of specifying an env var to tell whether we are building stuff locally or in Jenkins. (You can always just call uname in a script or System.getProperty("os.name") in Java or Scala.

Anyway, whatever. I did not have a chance to finish my comment to those 358 files in the PR, because there was a 1-1 with Jacob. I expected it to be an hour later, but whatever.

There was Jacob and there was his boss Mike. They told me they have to lay me off or something like that, because of my social skills. They were absolutely happy with my technical capabilities, but socially I’m not a fit. I thanked them politely.

Went ahead and wrote a goodbye on slack, telling them where to find me. One of the guys immediately send a connection request on linkedin (and we connected).

For me, it was a pretty nice company. Who exactly made sure that I don’t work with them, I don’t know, what is Chris? Was it Karen? I don’t believe it’s Karen, we never argued about anything. But who else, was Chris enough?

Oh, whatever. I immediately slapped together my resume, and wrote to my recruiter Sean. My salary demands are pretty low now, I rather care about the atmosphere and the medical insurance. Sean immediately called me; I told him I can start December 1st; also assured him that I have no plans to move permanently to France. Just visit it regularly.

So, amazingly, I have a whole month for myself; we can go for long hikes and for long bike rides, travel somewhere.

Anyway, they did not send me any email yet regarding this. I’ll have to at least try to pull out of them some compensation for the vacations that I never took (I know it’s a stupid thing to do, I don’t recommend anybody to save their vacation days these days.)

November 02, 2023

Well, I’m free now. Checked out the positions Sean sent me. Discarded the one with a gambling house, even though Kit Langdon works there.

Agreed to go lower a little bit. Wrote to Sean; he called me; we agreed to talk over WhatsApp in the future (free for me).

Also talked with Gabriel (a cool guy in Uruguay) about the opportunities in EarCandy, described him the situation. Then he also asked me about my interests; so we agreed to exchange information.

Oh, and what about Victor? I wrote to Sean, mentioning Victor, but he did not mention back.

Well, whatever. It's obvious now that the company got rid of me just because the CEO told them to lower the budget. No hard feelings, but they could have lied less. Anyway, this is all in the past.

Did nothing today; in a sense, took a short break. School vacations end this weekend.

November 03, 2023

So, it's really a pleasure to take this long break, but I will need an insurance in December.

Anyway, still trying to get out of the company at least access to Workday: they gave me the id, but the password does not work. Wrote to them. So far I don't have any document stating that I'm not working there anymore.

We got together on twitter, there's a bunch of people suddenly out of work - including Debasish, for instance. Debasish calls it "funemployment". Oh, and Paul, and Gabriel, and... some other people too.

Sean, the recruiter, wrote to me that he got in touch with Victor, and filed his resume to the same company as mine. That would be a pretty cool idea, to hire a team. Victor could bring more people.

And I'm thinking about my colleagues from EarCandy, and, you see, while it was pretty nice to work with them, and I'd love working again with Raghav, and Safa, and Swathi, and Karen, and Davis, I don't think I had any buddies in that company. Why is it, I don't know. From every previous company, however evil, I was getting some buddies, with whom I'd love to talk, to work, to hang out. Not here. Isolation. I'm not 100% sure it's bad; just an observation.

November 07, 2023

Ok, since the HR told me to call user support, I called user support. A guy there told me that since I'm not an employee, he can't help me. (Am I not? I don't know.) I told him HR sent me to him. He said he will create a case (Jira case?) for HR to handle it.

So, I'm still in the air. I love this status. Many people don't. I do.

Also, filed my resume to a Swiss company. Interesting, right? We'll see.

THIS ENDS MY STORY OF WORKING AT EarCandy.